



# PyTorch Tutorial

## 05. Linear Regression with PyTorch

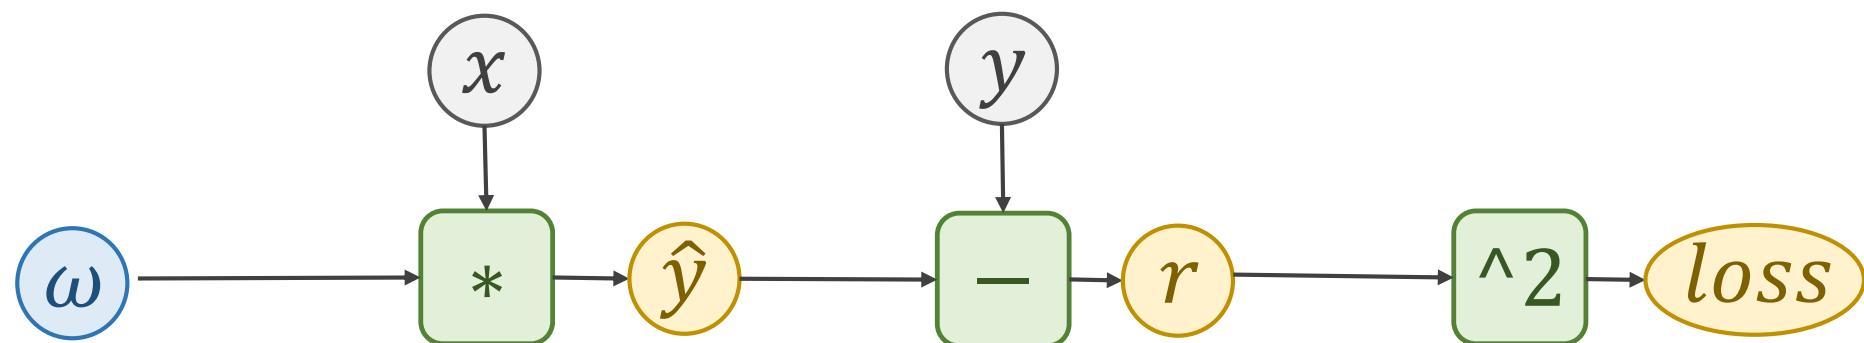
# Revision

Linear Model

$$\hat{y} = x * \omega$$

Loss Function

$$loss = (\hat{y} - y)^2 = (x \cdot \omega - y)^2$$



# Revision

```
print("predict (before training)", 4, forward(4).item())

for epoch in range(100):
    for x, y in zip(x_data, y_data):
        l = loss(x, y)
        l.backward()
        print('\tgrad:', x, y, w.grad.item())
        w.data = w.data - 0.01 * w.grad.data

        w.grad.data.zero_()

    print("progress:", epoch, l.item())

print("predict (after training)", 4, forward(4).item())
```

```
predict (before training) 4 4.0
grad: 1.0 2.0 -2.0
grad: 2.0 4.0 -7.840000152587891
grad: 3.0 6.0 -16.228801727294922
progress: 0 7.315943717956543
grad: 1.0 2.0 -1.478623867034912
grad: 2.0 4.0 -5.796205520629883
grad: 3.0 6.0 -11.998146057128906
progress: 1 3.9987640380859375
grad: 1.0 2.0 -1.0931644439697266
grad: 2.0 4.0 -4.285204887390137
grad: 3.0 6.0 -8.870372772216797
progress: 2 2.1856532096862793
grad: 1.0 2.0 -0.8081896305084229
grad: 2.0 4.0 -3.1681032180786133
grad: 3.0 6.0 -6.557973861694336
progress: 3 1.1946394443511963
grad: 1.0 2.0 -0.5975041389465332
grad: 2.0 4.0 -2.3422164916992188
grad: 3.0 6.0 -4.848389625549316
progress: 4 0.6529689431190491
grad: 1.0 2.0 -0.4417421817779541
grad: 2.0 4.0 -1.7316293716430664
grad: 3.0 6.0 -3.58447265625
progress: 5 0.35690122842788696
grad: 1.0 2.0 -0.3265852928161621
grad: 2.0 4.0 -1.2802143096923828
grad: 3.0 6.0 -2.650045394897461
```

# PyTorch Fashion

1

Prepare dataset

we shall talk about this later

2

Design **model** using Class

inherit from nn.Module

计算图

3

构造损失函数 · 构造  
Construct loss and optimizer  
using PyTorch API

4

Training cycle  
forward, backward, update

# Linear Regression – 1. Prepare dataset

In PyTorch, the computational graph is in mini-batch fashion, so  $X$  and  $Y$  are  $3 \times 1$  Tensors.

$$\begin{bmatrix} y_{pred}^{(1)} \\ y_{pred}^{(2)} \\ y_{pred}^{(3)} \end{bmatrix}_{1 \times 3} = \omega \cdot \begin{bmatrix} x^{(1)} \\ x^{(2)} \\ x^{(3)} \end{bmatrix}_{1 \times 3} + b$$

即把样本一维性都放进去  
必须是矩阵

```
import torch

x_data = torch.Tensor([[1.0], [2.0], [3.0]])
y_data = torch.Tensor([[2.0], [4.0], [6.0]])
```

# Revision: Gradient Descent Algorithm

## Derivative

$$\begin{aligned}\frac{\partial \text{cost}(\omega)}{\partial \omega} &= \frac{\partial}{\partial \omega} \frac{1}{N} \sum_{n=1}^N (x_n \cdot \omega - y_n)^2 \\ &= \frac{1}{N} \sum_{n=1}^N \frac{\partial}{\partial \omega} (x_n \cdot \omega - y_n)^2 \\ &= \frac{1}{N} \sum_{n=1}^N 2 \cdot (x_n \cdot \omega - y_n) \frac{\partial (x_n \cdot \omega - y_n)}{\partial \omega} \\ &= \frac{1}{N} \sum_{n=1}^N 2 \cdot x_n \cdot (x_n \cdot \omega - y_n)\end{aligned}$$

## Gradient

$$\frac{\partial \text{cost}}{\partial \omega}$$

## Update

$$\omega = \omega - \alpha \frac{\partial \text{cost}}{\partial \omega}$$

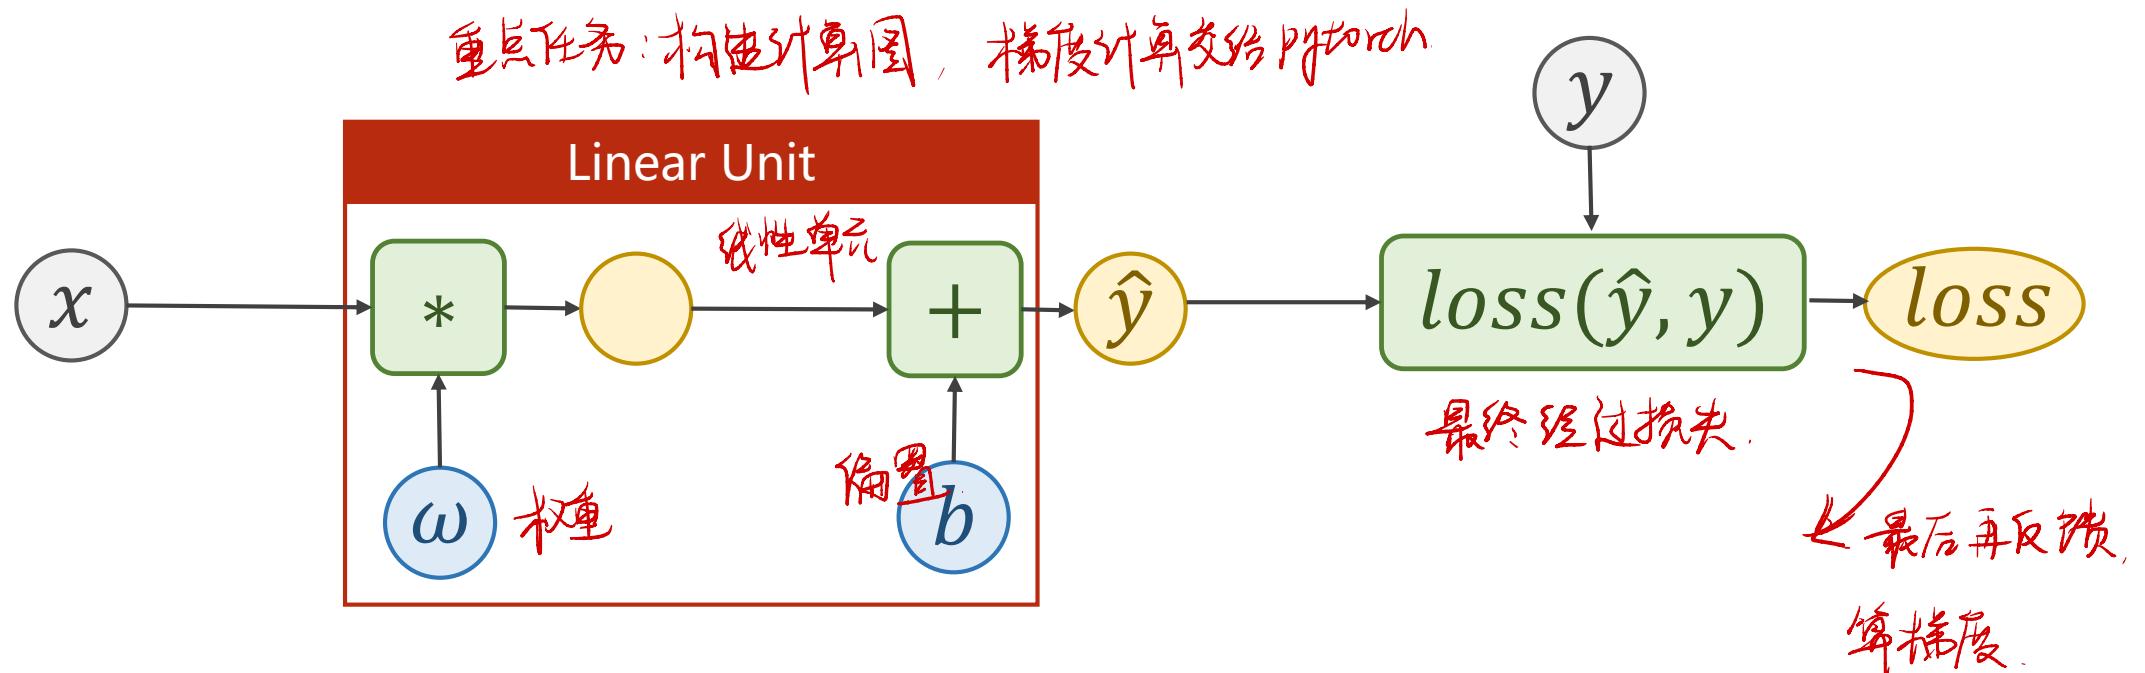
## Update

$$\omega = \omega - \alpha \frac{1}{N} \sum_{n=1}^N 2 \cdot x_n \cdot (x_n \cdot \omega - y_n)$$

# Linear Regression – 2. Design Model

Affine Model	$\hat{y} = x * \omega + b$
Loss Function	
$loss = (\hat{y} - y)^2 = (x \cdot \omega - y)^2$	

最终必须是一个标量，(非向量)



# Linear Regression – 2. Design Model

是一个类。 定义构造模型。

```
class LinearModel(torch.nn.Module):
    def __init__(self):
        super(LinearModel, self).__init__()
        self.linear = torch.nn.Linear(1, 1)

    def forward(self, x):
        y_pred = self.linear(x)
        return y_pred

model = LinearModel()
```

Our model class should be inherit from *nn.Module*, which is Base class for all neural network modules.

backward 在 torch 的 Function 中  
如求导等计算方式

# Linear Regression – 2. Design Model

```
class LinearModel(torch.nn.Module):
    def __init__(self):
        super(LinearModel, self).__init__()
        self.linear = torch.nn.Linear(1, 1)

    def forward(self, x):
        y_pred = self.linear(x)
        return y_pred

model = LinearModel()
```

Member methods *\_\_init\_\_()* and *forward()* have to be implemented.

# Linear Regression – 2. Design Model

```
class LinearModel(torch.nn.Module):
    def __init__(self):
        super(LinearModel, self).__init__()
        self.linear = torch.nn.Linear(1, 1)

    def forward(self, x):
        y_pred = self.linear(x)
        return y_pred

model = LinearModel()
```

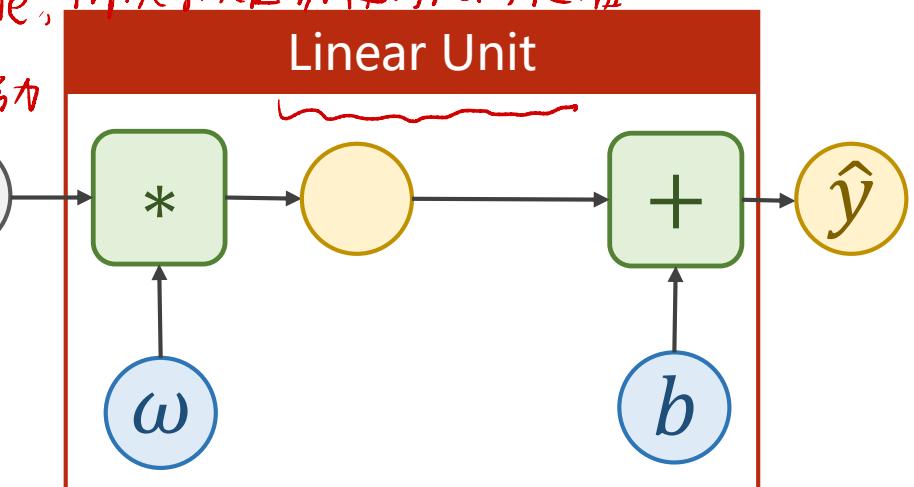
Just do it. : )

# Linear Regression – 2. Design Model

```
class LinearModel(torch.nn.Module):
    def __init__(self):
        super(LinearModel, self).__init__()
        self.linear = torch.nn.Linear(1, 1)
    def forward(self, x):
        y_pred = self.linear(x)
        return y_pred
model = LinearModel()
```

pytorch 中的一个类，继承自 module，所以可以自动进行反向传播  
Linear中包含权重、偏置，可以自动进行输入权重+偏置的计算  
↓ 是一个对象

Class *nn.Linear* contain two member **Tensors**: **weight** and **bias**.



# Linear Regression – 2. Design Model

`class torch.nn.Linear(in_features, out_features, bias=True)` [source]

Applies a linear transformation to the incoming data:  $y = Ax + b$

Linear 模型的运算

Parameters:

- in\_features – size of each input sample  
是否有偏置  $b$  部分
- out\_features – size of each output sample
- bias – If set to False, the layer will not learn an additive bias. Default: `True`

Shape:

$$[y_{n \times 2}] = [x_{n \times 3}] [w_{3 \times 2}] \quad , \text{但实际一般作} \downarrow \quad [y_{2 \times n}] = [w^T] [x_{3 \times n}] + [b]$$

行数为特征数量, 列数为样本数量

- Input:  $(N, *, in\_features)$  where \* means any number of additional dimensions
- Output:  $(N, *, out\_features)$  where all but the last dimension are the same shape as the input.

Variables:

- **weight** – the learnable weights of the module of shape  $(out\_features \times in\_features)$
- **bias** – the learnable bias of the module of shape  $(out\_features)$

# Linear Regression – 2. Design Model

`class torch.nn.Linear(in_features, out_features, bias=True)` [source]

Applies a linear transformation to the incoming data:  $y = Ax + b$

Parameters:

$$\text{Output} \begin{bmatrix} y_{pred}^{(1)} \\ y_{pred}^{(2)} \\ y_{pred}^{(3)} \end{bmatrix} = \omega \cdot \begin{bmatrix} x^{(1)} \\ x^{(2)} \\ x^{(3)} \end{bmatrix} + b \quad \text{Input}$$

tive bias. Default: `True`

Shape:

- Input:  $(N, *, in\_features)$  where \* means any number of additional dimensions
- Output:  $(N, *, out\_features)$  where all but the last dimension are the same shape as the input.

Variables:

- **weight** – the learnable weights of the module of shape  $(out\_features \times in\_features)$
- **bias** – the learnable bias of the module of shape  $(out\_features)$

# Linear Regression – 2. Design Model

```
class LinearModel(torch.nn.Module):  
    def __init__(self):  
        super(LinearModel, self).__init__()  
        self.linear = torch.nn.Linear(1, 1)  
  
    def forward(self, x):  
        y_pred = self.linear(x)  
        return y_pred  
  
model = LinearModel()
```

实例化模型，  
且可调用  
此对象

即对象中有 \_\_call\_\_ 方法

```
class FooBar:  
    def __init__(self):  
        pass  
  
    def __call__(self, *args, **kwargs):  
        pass
```

Class *nn.Linear* <sup>→ 可调用</sup> has implemented the magic method *\_\_call\_\_()*, which enable the instance of the class can be called just like a function. Normally the *forward()* will be called.

Pythonic!!!

即任何没有归类的参数都归给 \*args, 如 *def f(\*args, x)*, *f(a,b,x=1)*, *a,b* 作为元组传给 \*args

# Linear Regression – 2. Design Model

\*\*kwargs 传递旧参数如 x=1, y=2，并将其转化为词典

```
class LinearModel(torch.nn.Module):  
    def __init__(self):  
        super(LinearModel, self).__init__()  
        self.linear = torch.nn.Linear(1, 1)  
  
    def forward(self, x):  
        y_pred = self.linear(x)  
        return y_pred  
  
model = LinearModel()
```

Create a instance of class  
**LinearModel.**

# Linear Regression – 3. Construct Loss and Optimizer

```
需要写子即可求loss      → 继承自 module  
criterion = torch.nn.MSELoss(size_average=False)  → 是否平均值  
  
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
```

→ 只否将均变为一位). 一般默认即可

```
class torch.nn.MSELoss(size_average=True, reduce=True) [source]
```

Creates a criterion that measures the mean squared error between target  $y$ .

The loss can be described as:

$$\ell(x, y) = L = \{l_1, \dots, l_N\}^\top, \quad l_n = (x_n - y_n)^2,$$

where  $N$  is the batch size.

Also inherit from **nn.Module**.

# Linear Regression – 3. Construct Loss and Optimizer

构造损失函数  
Optimizer

```
criterion = torch.nn.MSELoss(size_average=False)
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
```

告诉优化器对哪些权重作优化  
→ 权重，会检查 module 中的所有成员，如果成员中有相应权重，  
就把包加到要训练的参数集中

学习率

class torch.optim.SGD(params, lr=<object object>, momentum=0, dampening=0, weight\_decay=0, nesterov=False) [source]

Implements stochastic gradient descent (optionally with momentum).

# Linear Regression – 3. Construct Loss and Optimizer

```
criterion = torch.nn.MSELoss(size_average=False)  
  
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
```



- Parameters:**
- **params** (*iterable*) – iterable of parameters to optimize or dicts defining parameter groups
  - **lr** (*float*) – learning rate

# Linear Regression – 4. Training Cycle

训练过程

```
for epoch in range(100):  
    y_pred = model(x_data) ← 算分  
    loss = criterion(y_pred, y_data)  
    print(epoch, loss)
```

```
optimizer.zero_grad()  
loss.backward()  
optimizer.step()
```

Forward: Predict

# Linear Regression – 4. Training Cycle

```
for epoch in range(100):
    y_pred = model(x_data)
    loss = criterion(y_pred, y_data)      ↗ 马损失
    print(epoch, loss)
    loss 是个 tensor, 但打印时
    optimizer.zero_grad()    会自动调用--str-- 函数, 所以没有问题.
    loss.backward()
    optimizer.step()
```

Forward: Loss

# Linear Regression – 4. Training Cycle

```
for epoch in range(100):
    y_pred = model(x_data)
    loss = criterion(y_pred, y_data)
    print(epoch, loss)

    optimizer.zero_grad() ←
    loss.backward()
    optimizer.step()
```

## NOTICE:

The grad computed by *.backward()* will be **accumulated**. So before backward, remember set the grad to **ZERO!!!**

# Linear Regression – 4. Training Cycle

```
for epoch in range(100):
    前向 < y_pred = model(x_data)
    loss = criterion(y_pred, y_data)
    print(epoch, loss)

    反向 < optimizer.zero_grad()          计算梯度 (反向传播)
    loss.backward()                      ←
    更新 → optimizer.step()            根据梯度和包含的权重
                                         ⏪ 进行一次更新
```

Backward: Autograd

# Linear Regression – 4. Training Cycle

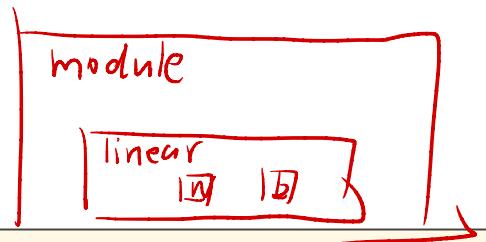
```
for epoch in range(100):  
    y_pred = model(x_data)  
    loss = criterion(y_pred, y_data)  
    print(epoch, loss)
```

```
optimizer.zero_grad()  
loss.backward()  
optimizer.step()
```

```
for x, y in zip(x_data, y_data):  
    .....  
    w.data = w.data - 0.01 * w.grad.data
```

Update

# Linear Regression – Test Model



```
# Output weight and bias
print('w = ', model.linear.weight.item())
print('b = ', model.linear.bias.item())
# Test Model
x_test = torch.Tensor([[4.0]])
y_test = model(x_test)
print('y_pred = ', y_test.data)
```

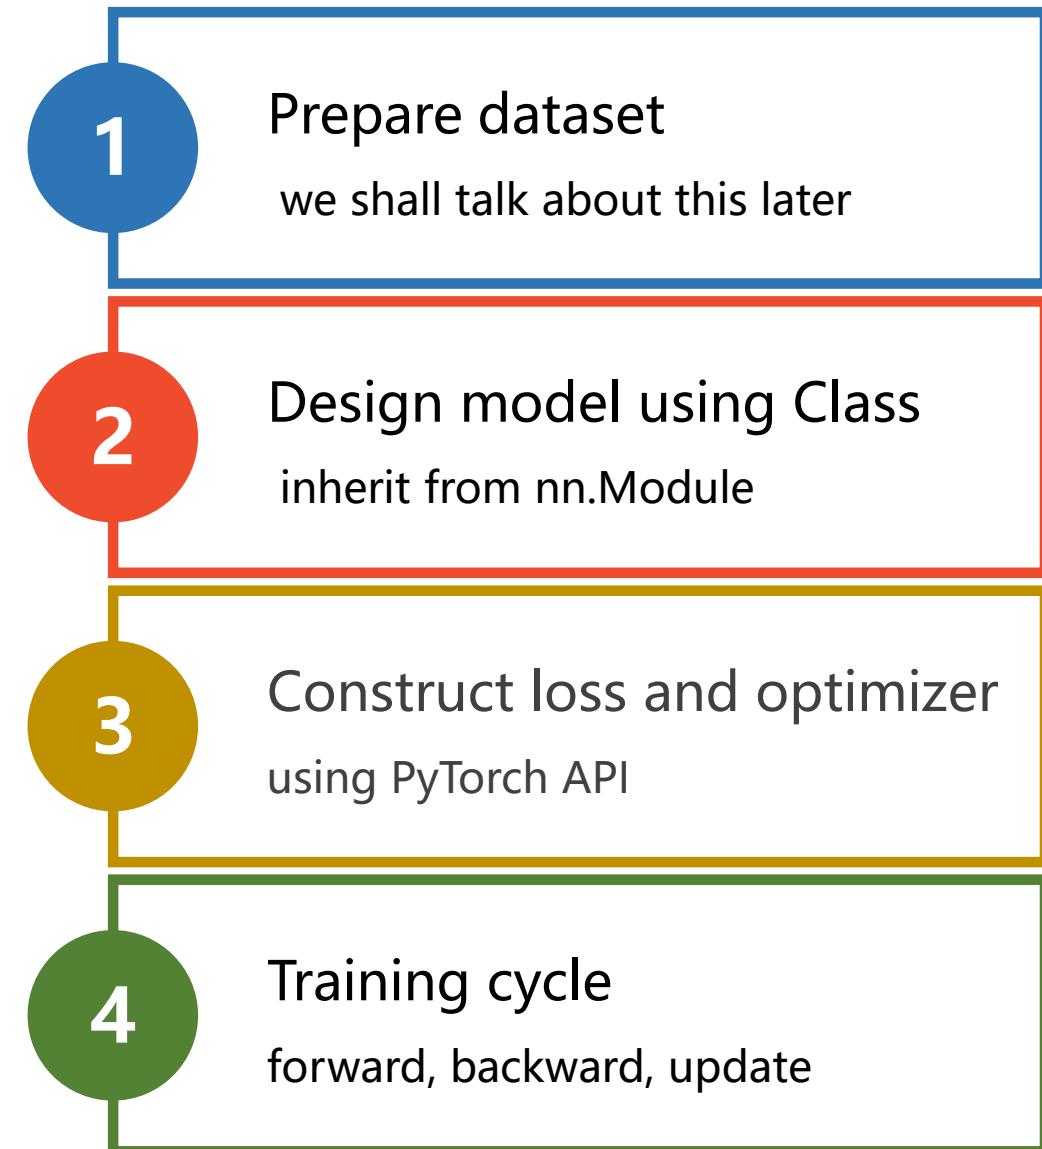
86 0.3036523759365082	986 3.594939812501252e-07
87 0.2992883026599884	987 3.5411068211033125e-07
88 0.29498720169067383	988 3.4917979974125046e-07
89 0.2907477021217346	989 3.4428359185767476e-07
90 0.28656935691833496	990 3.392528924450744e-07
91 0.28245046734809875	991 3.3442694302721065e-07
92 0.27839142084121704	992 3.294019847999152e-07
93 0.27439042925834656	993 3.247135396122758e-07
94 0.2704470157623291	994 3.199925231456291e-07
95 0.2665606141090393	995 3.1540417921860353e-07
96 0.262729674577713	996 3.1097857799977646e-07
97 0.25895369052886963	997 3.0668098816022393e-07
98 0.2552322745323181	998 3.020934400410624e-07
99 0.2515641450881958	999 2.977626536448952e-07
w = 1.666100263595581	w = 1.9996366500854492
b = 0.7590328454971313	b = 0.0008257834706455469
y_pred = tensor([[ 7.4234]])	y_pred = tensor([[ 7.9994]])

100 Iterations

1000 Iterations

# Linear Regression

```
import torch  
  
① x_data = torch.Tensor([[1.0], [2.0], [3.0]])  
y_data = torch.Tensor([[2.0], [4.0], [6.0]])  
  
② class LinearModel(torch.nn.Module):  
    def __init__(self):  
        super(LinearModel, self).__init__()  
        self.linear = torch.nn.Linear(1, 1)  
  
    def forward(self, x):  
        y_pred = self.linear(x)  
        return y_pred  
model = LinearModel()  
  
③ criterion = torch.nn.MSELoss(size_average=False)  
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)  
  
④ for epoch in range(1000):  
    y_pred = model(x_data)  
    loss = criterion(y_pred, y_data)  
    print(epoch, loss.item())  
  
    optimizer.zero_grad()  
    loss.backward()  
    optimizer.step()  
  
    print('w = ', model.linear.weight.item())  
    print('b = ', model.linear.bias.item())  
  
x_test = torch.Tensor([[4.0]])  
y_test = model(x_test) // 使用模型  
print('y_pred = ', y_test.data)
```



# Exercise 5-1: Try Different Optimizer in Linear Regression

- torch.optim.Adagrad
- torch.optim.Adam
- torch.optim.Adamax
- torch.optim.ASGD
- torch.optim.LBFGS
- torch.optim.RMSprop
- torch.optim.Rprop
- torch.optim.SGD

pytorch 的优化器.

# Exercise 5-2: Read more example from official tutorial

## Table of Contents

- Tensors
  - Warm-up: numpy
  - PyTorch: Tensors
- Autograd
  - PyTorch: Tensors and autograd
  - PyTorch: Defining new autograd functions
  - TensorFlow: Static Graphs
- nn module
  - PyTorch: nn
  - PyTorch: optim
  - PyTorch: Custom nn Modules
  - PyTorch: Control Flow + Weight Sharing
- Examples
  - Tensors
  - Autograd
  - nn module

[https://pytorch.org/tutorials/beginner/pytorch\\_with\\_examples.html](https://pytorch.org/tutorials/beginner/pytorch_with_examples.html)



# PyTorch Tutorial

## 05. Linear Regression with PyTorch