

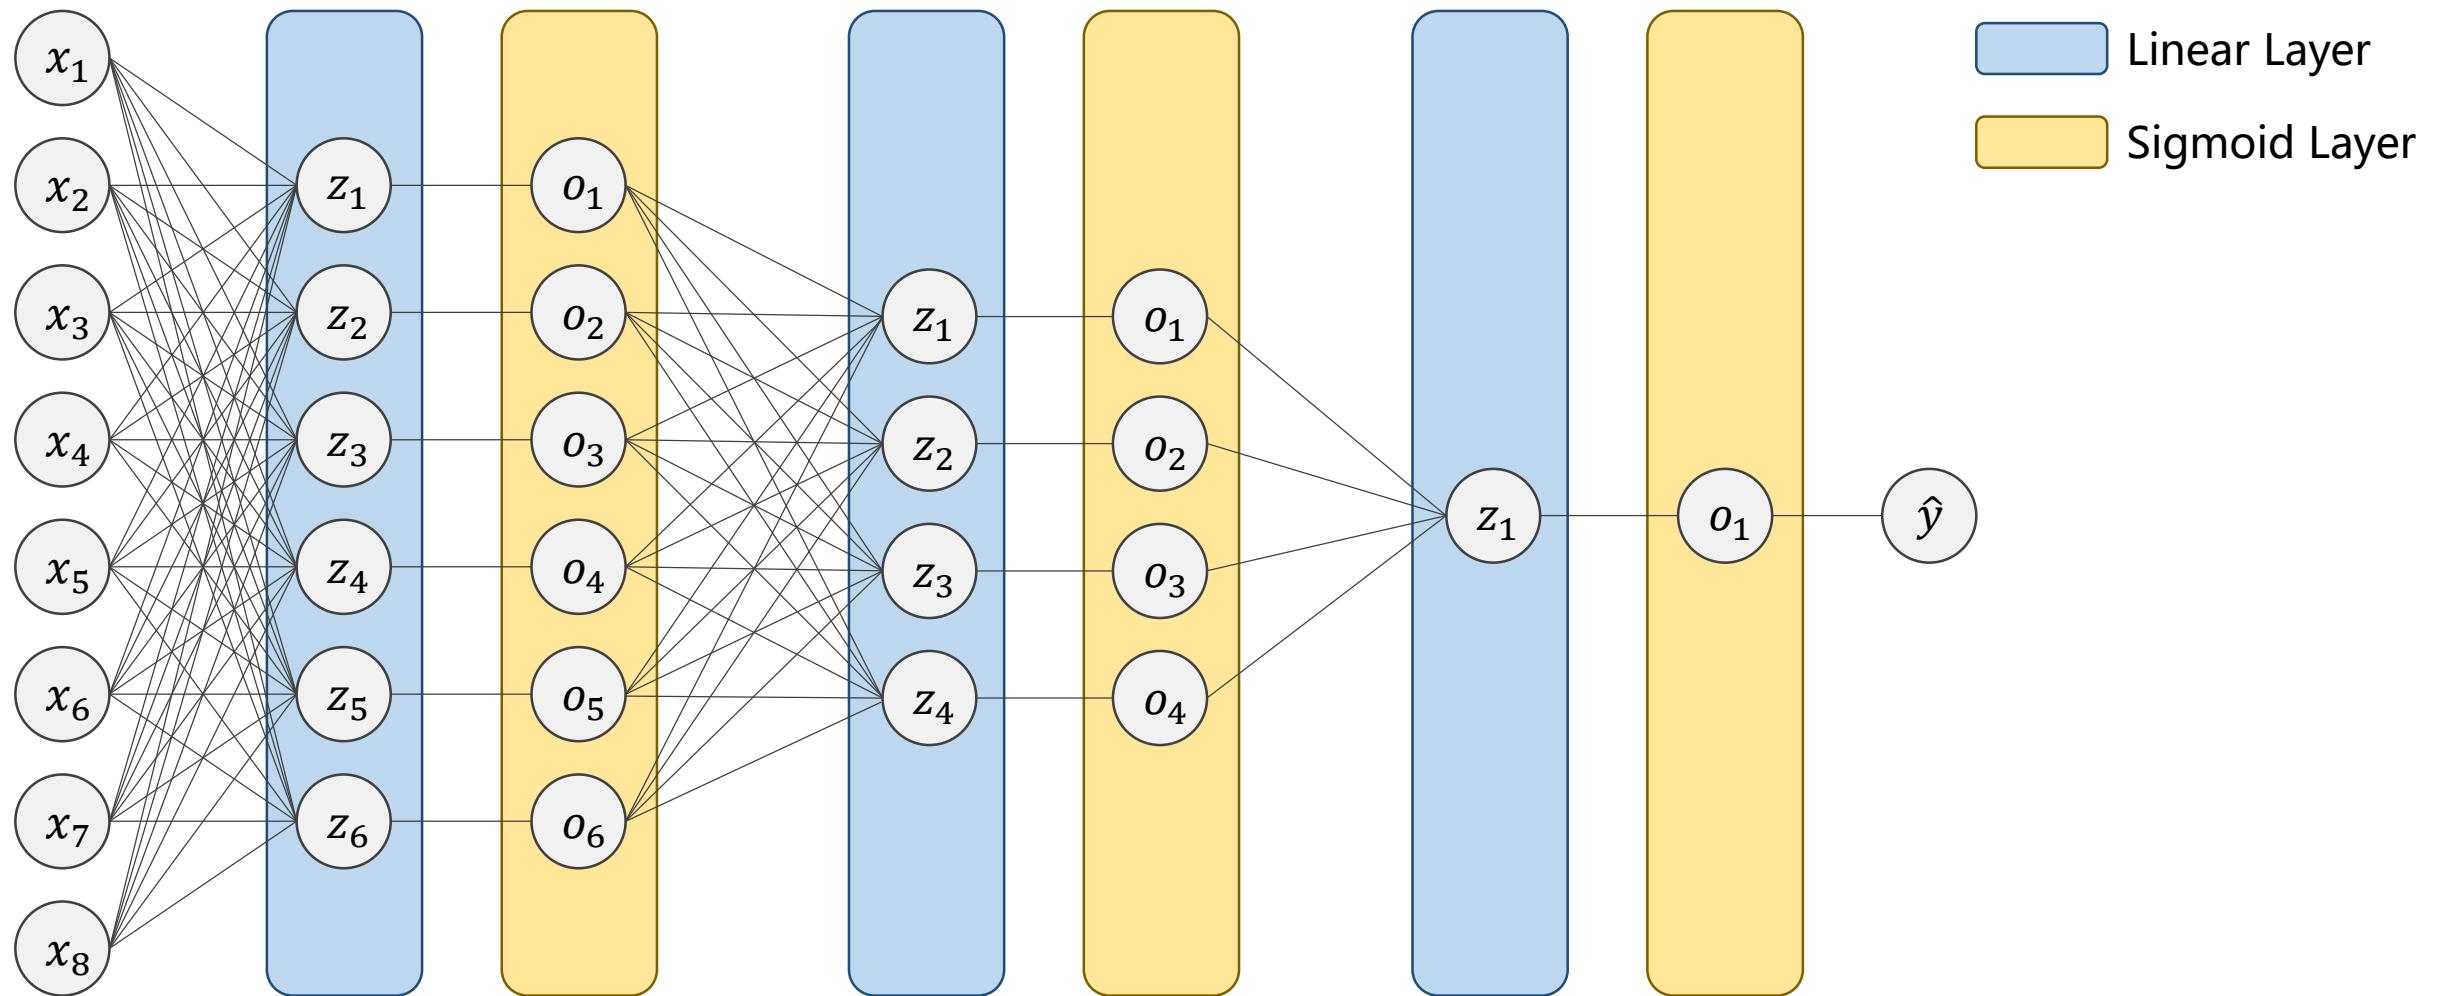


PyTorch Tutorial

09. Softmax Classifier

Softmax 分类器

Revision: Diabetes dataset



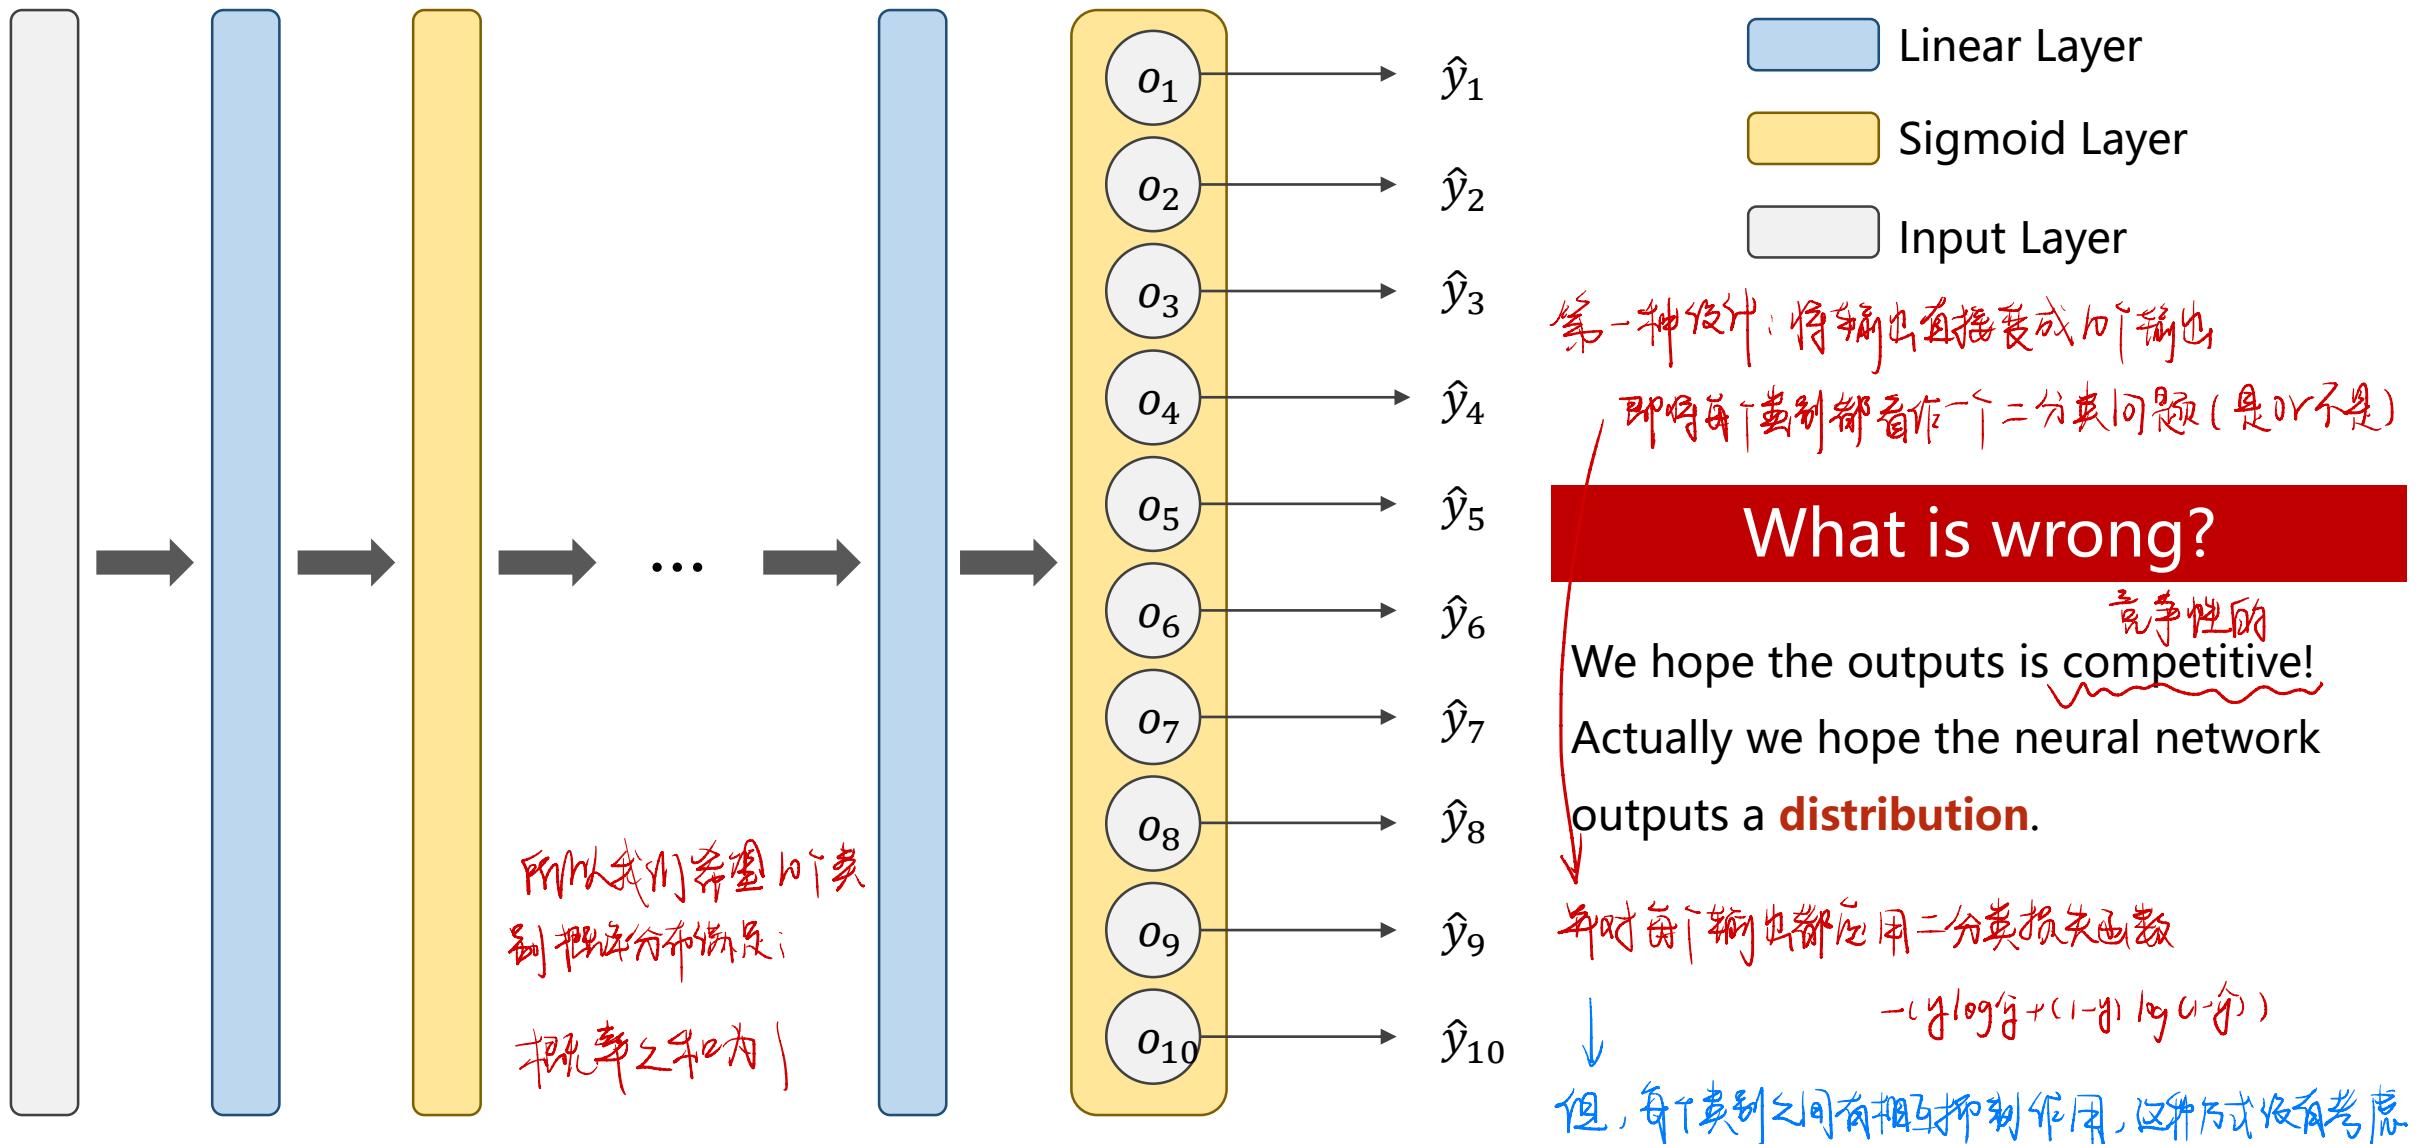
Revision: MNIST Dataset

5	0	4	1	9	2	1	3	1	4	3	5	3	6	1	7
2	8	6	9	4	0	9	1	1	2	4	3	2	7	3	8
6	9	0	5	6	0	7	4	1	8	7	9	3	9	8	5
9	3	3	0	7	4	9	8	0	9	4	1	4	4	6	0
4	5	6	7	1	0	0	1	7	1	6	3	0	2	1	1
8	0	2	6	7	8	3	9	0	4	6	7	4	6	8	0
7	8	3	1	5	7	1	7	1	1	6	3	0	2	9	3
1	1	0	4	9	2	0	0	2	0	2	7	1	8	6	4
1	6	3	4	3	9	\	3	3	8	5	4	7	7	4	2
8	5	8	6	7	3	4	6	1	9	9	6	0	3	7	2
8	2	9	4	4	6	4	9	7	0	9	2	7	5	1	5
9	1	0	3	1	3	5	9	1	7	6	2	8	2	1	5
0	7	4	9	7	8	3	2	1	1	8	3	6	1	0	3
1	0	0	1	1	2	7	3	0	4	6	5	2	6	4	7
1	8	9	9	3	0	7	1	0	2	0	3	5	4	6	5
8	6	3	7	5	8	0	9	1	0	3	1	2	2	3	3

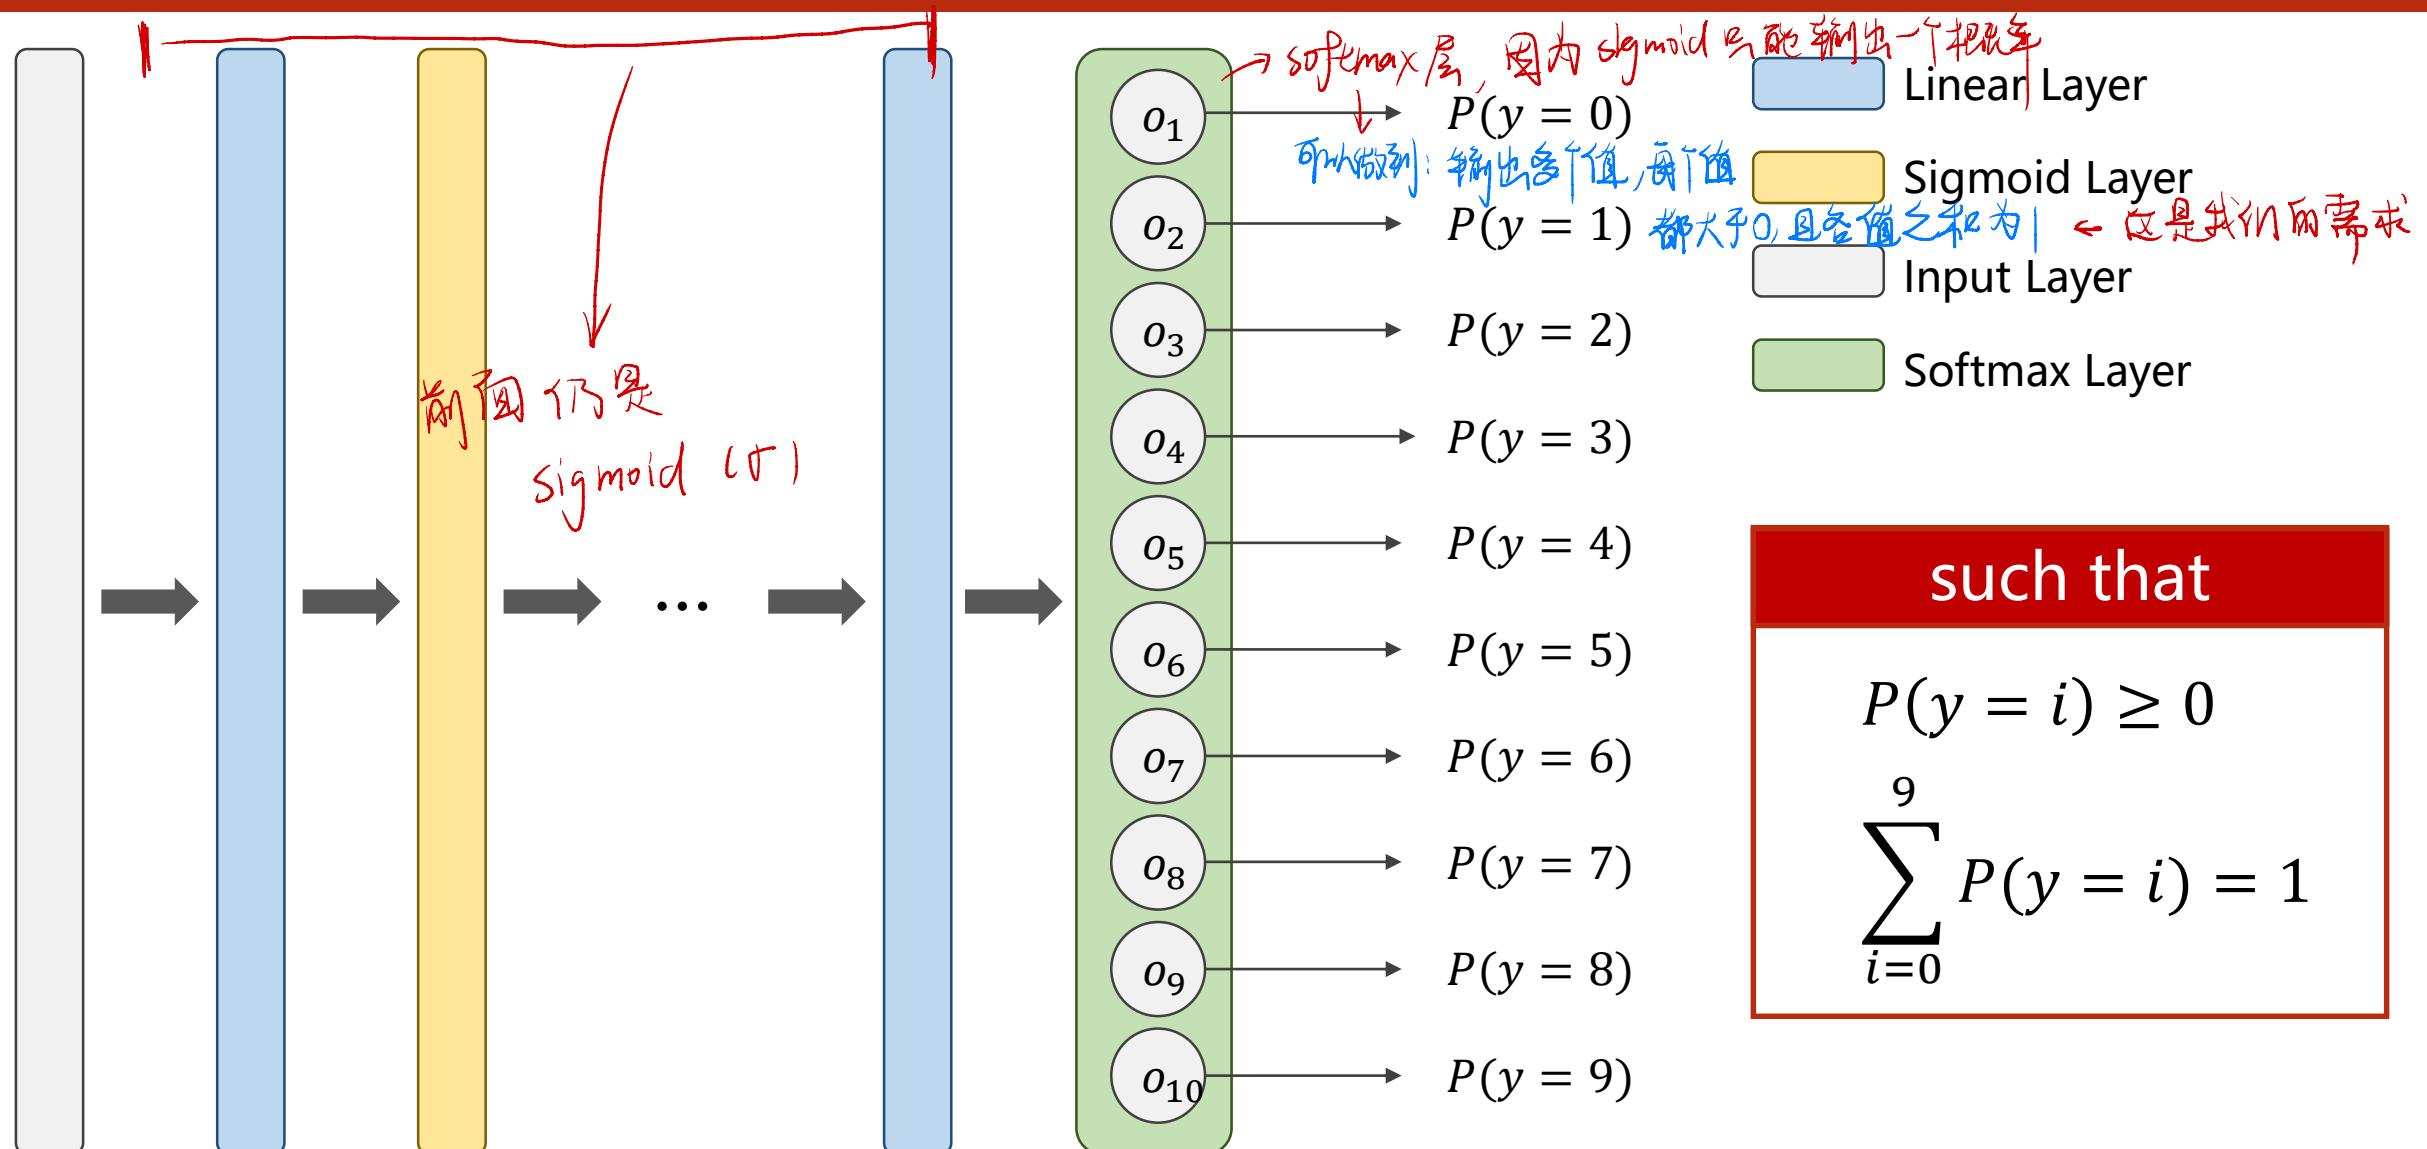
There are **10** labels in MNIST dataset.

How to design the neural network?

Design 10 outputs using Sigmoid?



Output a Distribution of prediction with Softmax



Softmax Layer

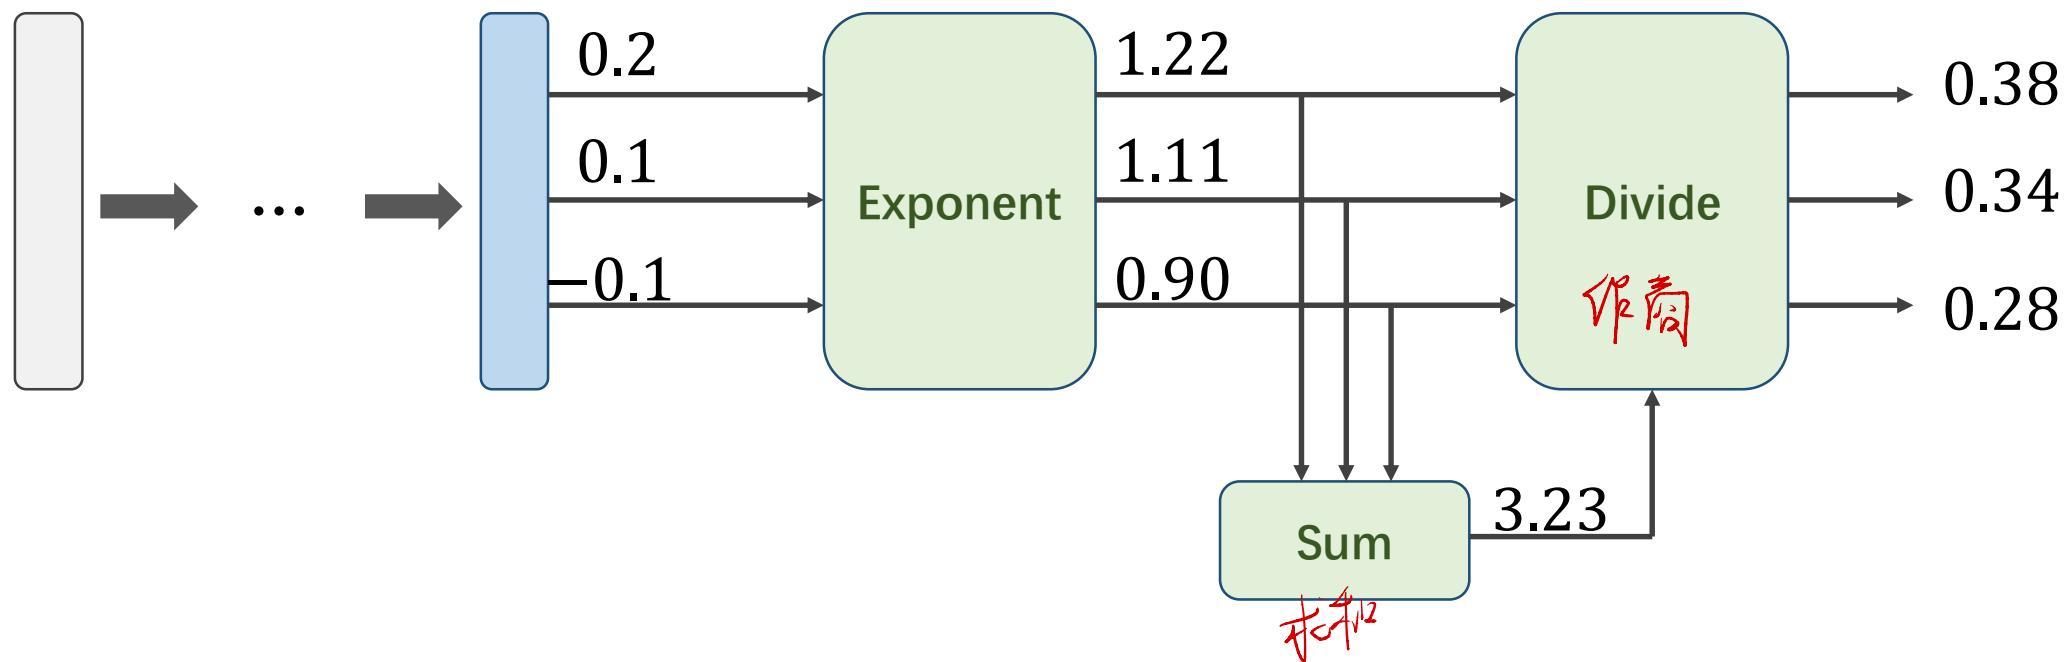
Suppose $Z^l \in \mathbb{R}^K$ is the output of the last linear layer, the Softmax function:

$$P(y = i) = \frac{e^{z_i}}{\sum_{j=0}^{K-1} e^{z_j}}, i \in \{0, \dots, K - 1\}$$

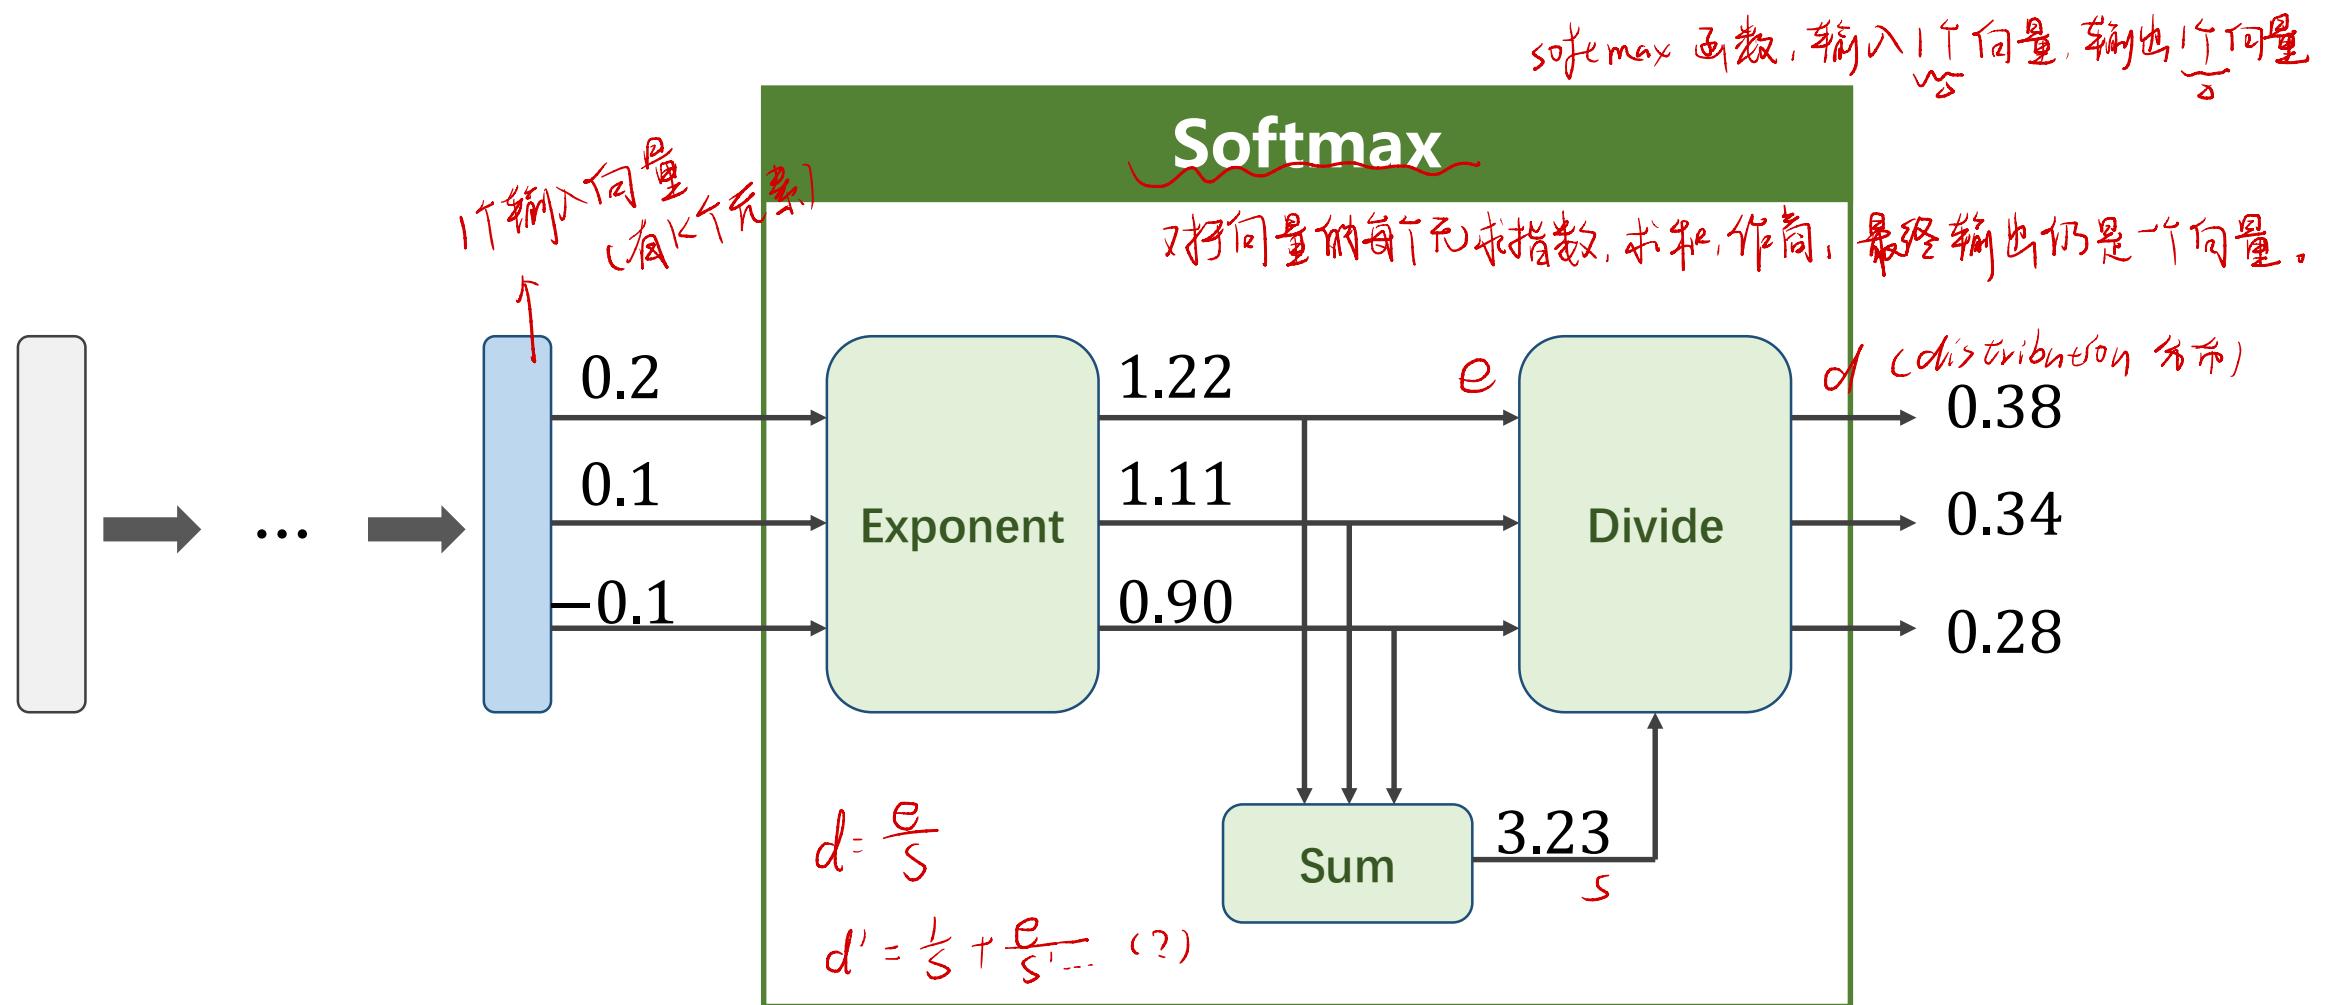
为什么用指数呢？因为指数可以保证大于0
softmax之前输出K个值 x_i
 x_i 是标量，
 $x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_K \end{bmatrix}$ 是向量

$z_i = x_i w + b$ ，
分子为求和，保证概率和为1

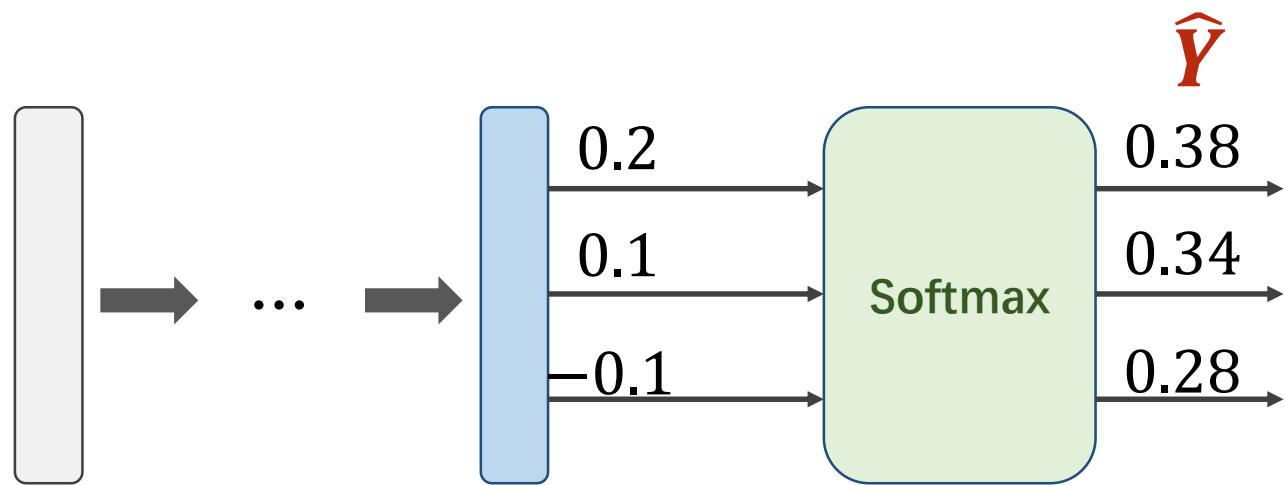
Softmax Layer - Example



Softmax Layer - Example



Loss function - Cross Entropy



我们看二分类的损失函数：
(交叉熵)

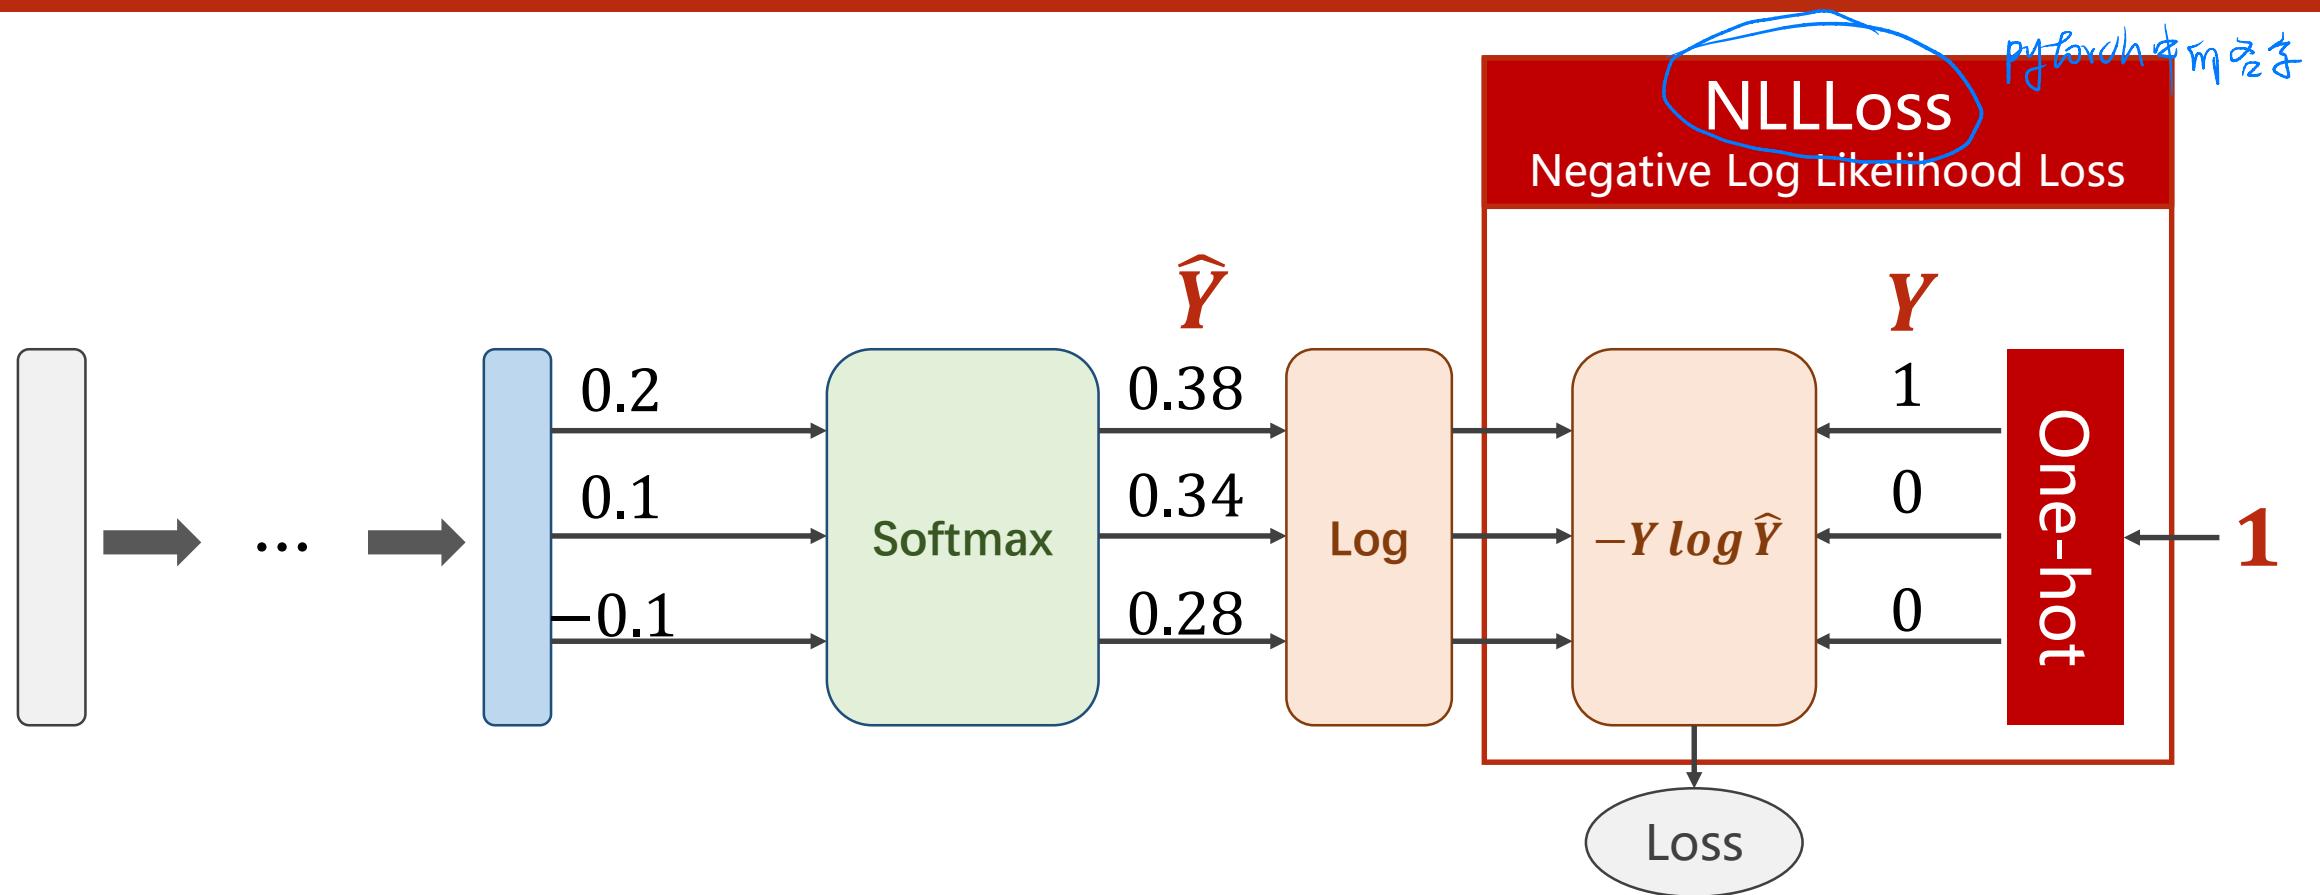
$-(y \log \hat{y} + (1-y) \log(1-\hat{y}))$ 在任何时候，两项中只能有一项非0

就是对于 实际分布 预测分布 两者相乘再相加，对于多分类仍相同

1	$\log \hat{y}$	只是2项变为N项，但同一时刻
0	$\log(1-\hat{y})$	只能有1项为0

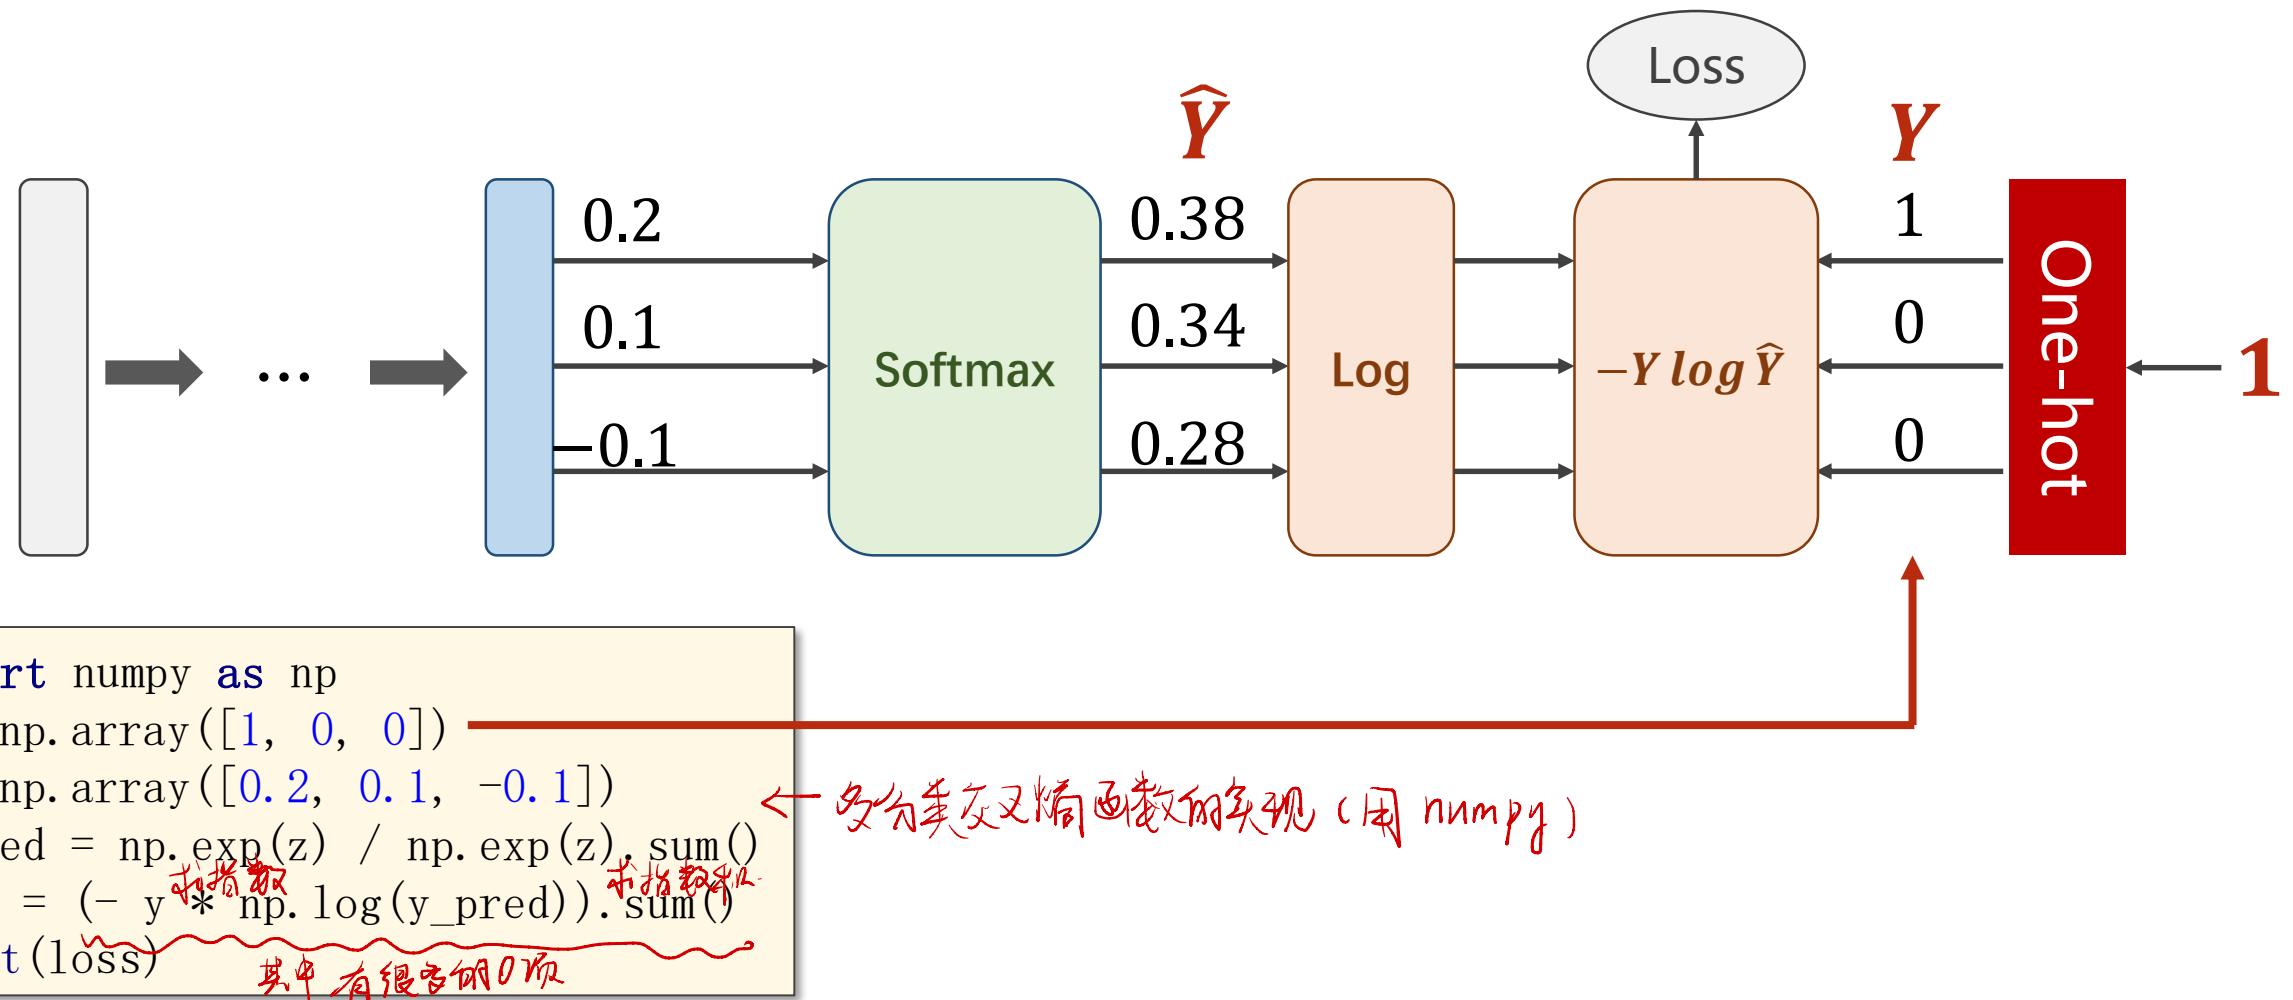
若 \hat{y} 为小于0.5的那一方，则 $1-\hat{y}$ 大于0.5
那么 \hat{y} 对应0, $1-\hat{y}$ 对应1
这和 \hat{y} 对应1, $1-\hat{y}$ 对应0是一样的。
 \hat{y}_1 和 \hat{y}_2 只是二分类中两个预测值，和多分类
中 \hat{y}_1 , \hat{y}_2 ... \hat{y}_n 地位相同

Loss function - Cross Entropy

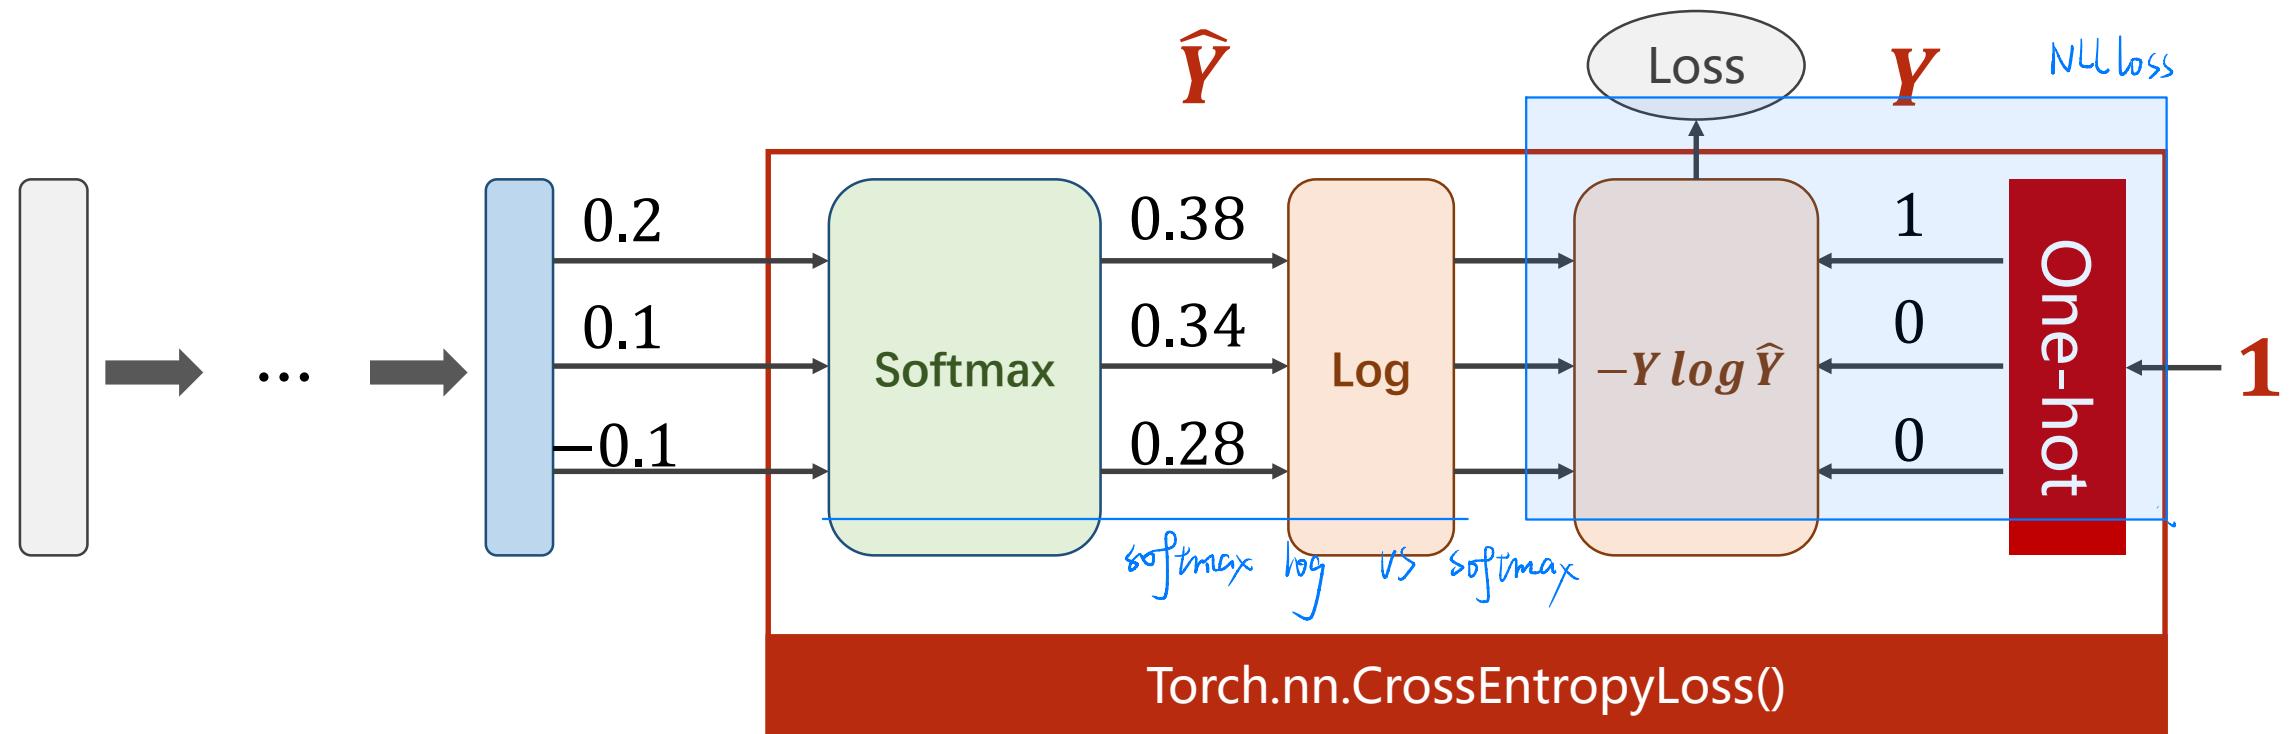


$$Loss(\hat{Y}, Y) = -Y \log \hat{Y}$$

Cross Entropy in Numpy



Cross Entropy in PyTorch



```
import torch  
y = torch.LongTensor([0]) // y是长整型张量  
z = torch.Tensor([[0.2, 0.1, -0.1]])  
criterion = torch.nn.CrossEntropyLoss()  
loss = criterion(z, y)  
print(loss)
```

//直接使用

CrossEntropyLoss 交叉熵损失

已经包括了 softmax 非线性变换.

*所以在使用交叉熵损失时, 神经网络最后一层不要作激活
(不作非线性变换)

Mini-Batch: $batch_size=3$

例：

```
import torch
criterion = torch.nn.CrossEntropyLoss()
Y = torch.LongTensor([2, 0, 1])
    第2类 0类 1类
Y_pred1 = torch.Tensor([[0.1, 0.2, 0.9], 样本1 - 2类
                      [1.1, 0.1, 0.2], 样本2 - 0类，比较响亮，损失小
                      [0.2, 2.1, 0.1]])
Y_pred2 = torch.Tensor([[0.8, 0.2, 0.3], 样本3 - 1类
                      [0.2, 0.3, 0.5], → 不响亮，损失大
                      [0.2, 0.2, 0.5]])
11 = criterion(Y_pred1, Y)
12 = criterion(Y_pred2, Y)
print("Batch Loss1 =", 11.data, "\nBatch Loss2=", 12.data)
```

Batch Loss1 = tensor(0.4966)

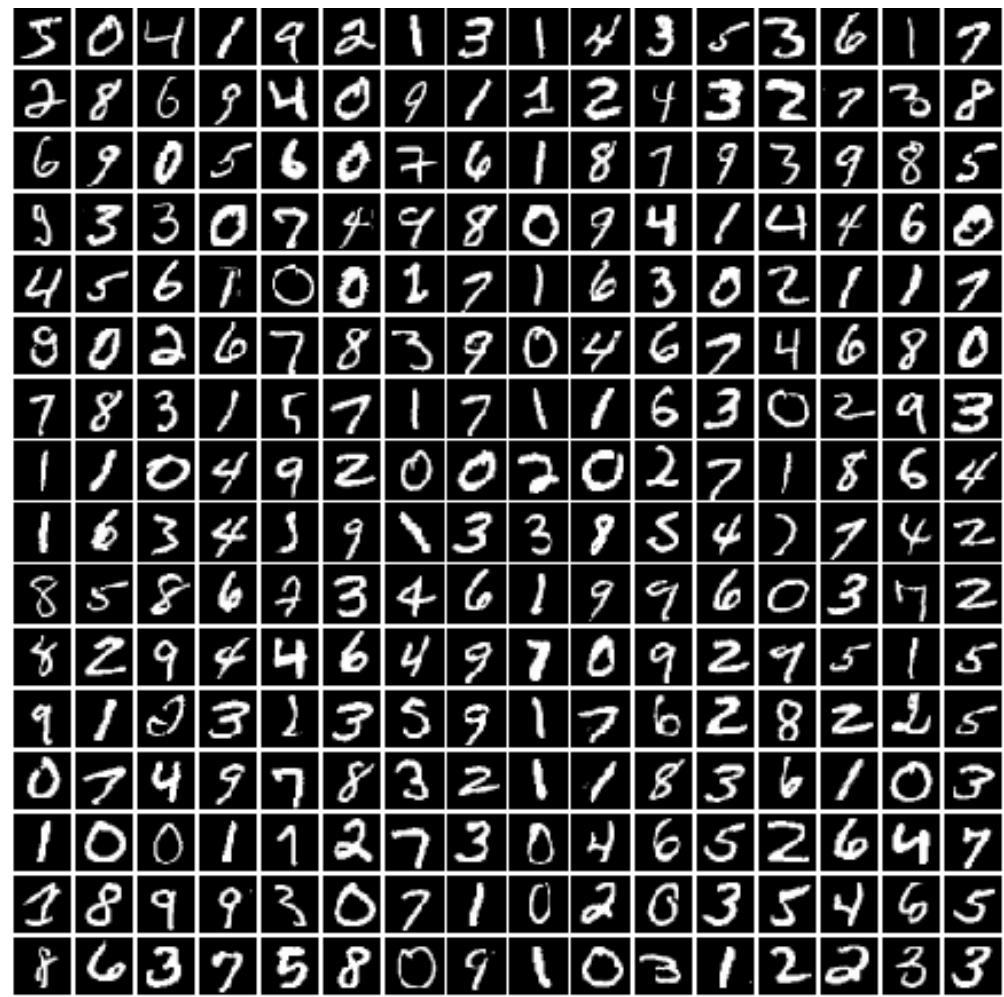
Batch Loss2 = tensor(1.2389)

Exercise 9-1: CrossEntropyLoss vs NLLLoss

- What are the differences?
- Reading the document:
 - <https://pytorch.org/docs/stable/nn.html#crossentropyloss>
 - <https://pytorch.org/docs/stable/nn.html#nllloss>
- Try to know why:
 - CrossEntropyLoss \Leftrightarrow LogSoftmax + NLLLoss

$\underbrace{\text{LogSoftmax}}$ 
相当于一个激活函数

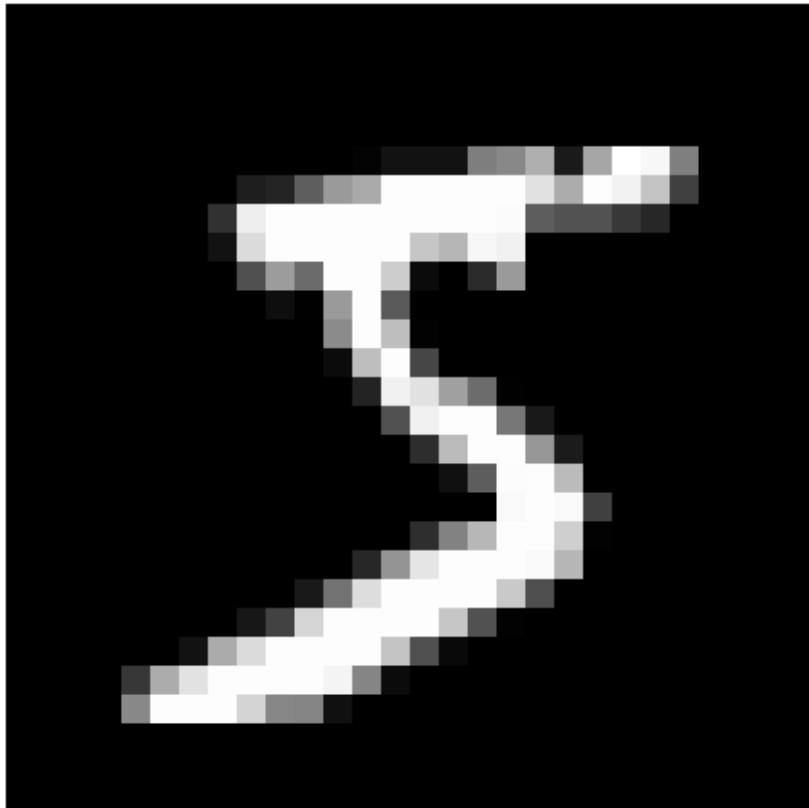
Back to MNIST Dataset



There are **10** labels in MNIST dataset.

How to design the neural network?

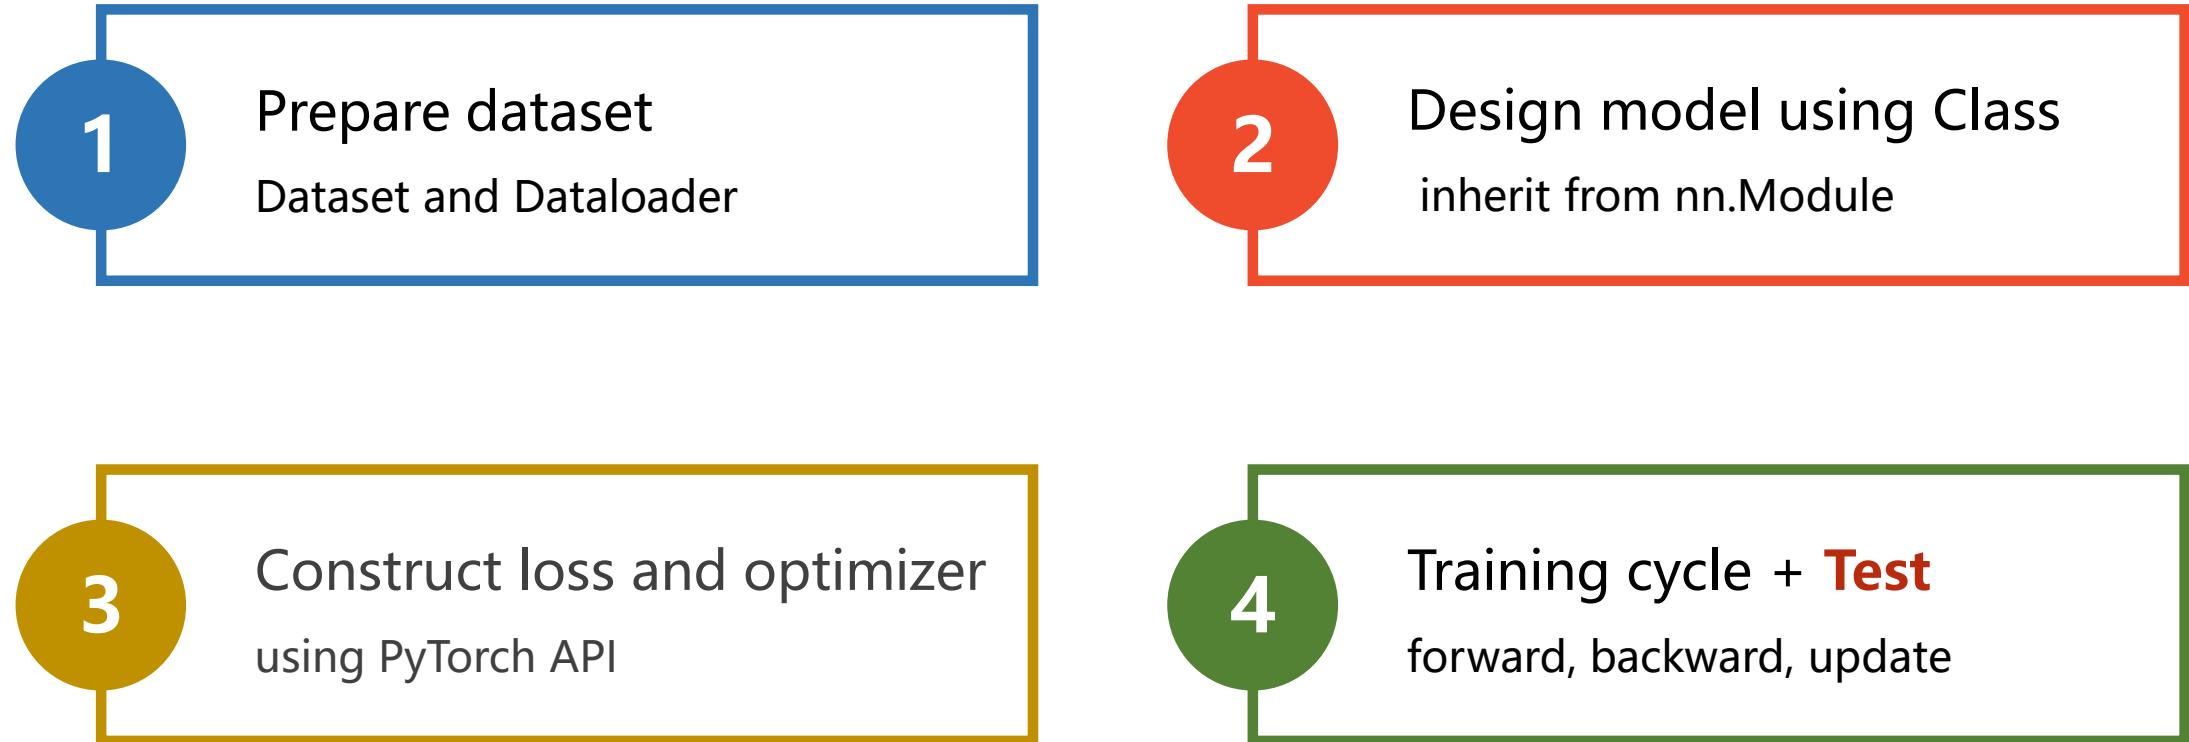
MNIST Dataset



$$28 * 28 = 784$$

图像用矩阵
保存

Implementation of classifier to MNIST dataset



Implementation – 0. Import Package

```
import torch
from torchvision import transforms
from torchvision import datasets
from torch.utils.data import DataLoader
import torch.nn.functional as F
import torch.optim as optim
```

For constructing DataLoader

Implementation – 0. Import Package

```
import torch
from torchvision import transforms
from torchvision import datasets
from torch.utils.data import DataLoader
import torch.nn.functional as F
import torch.optim as optim
```

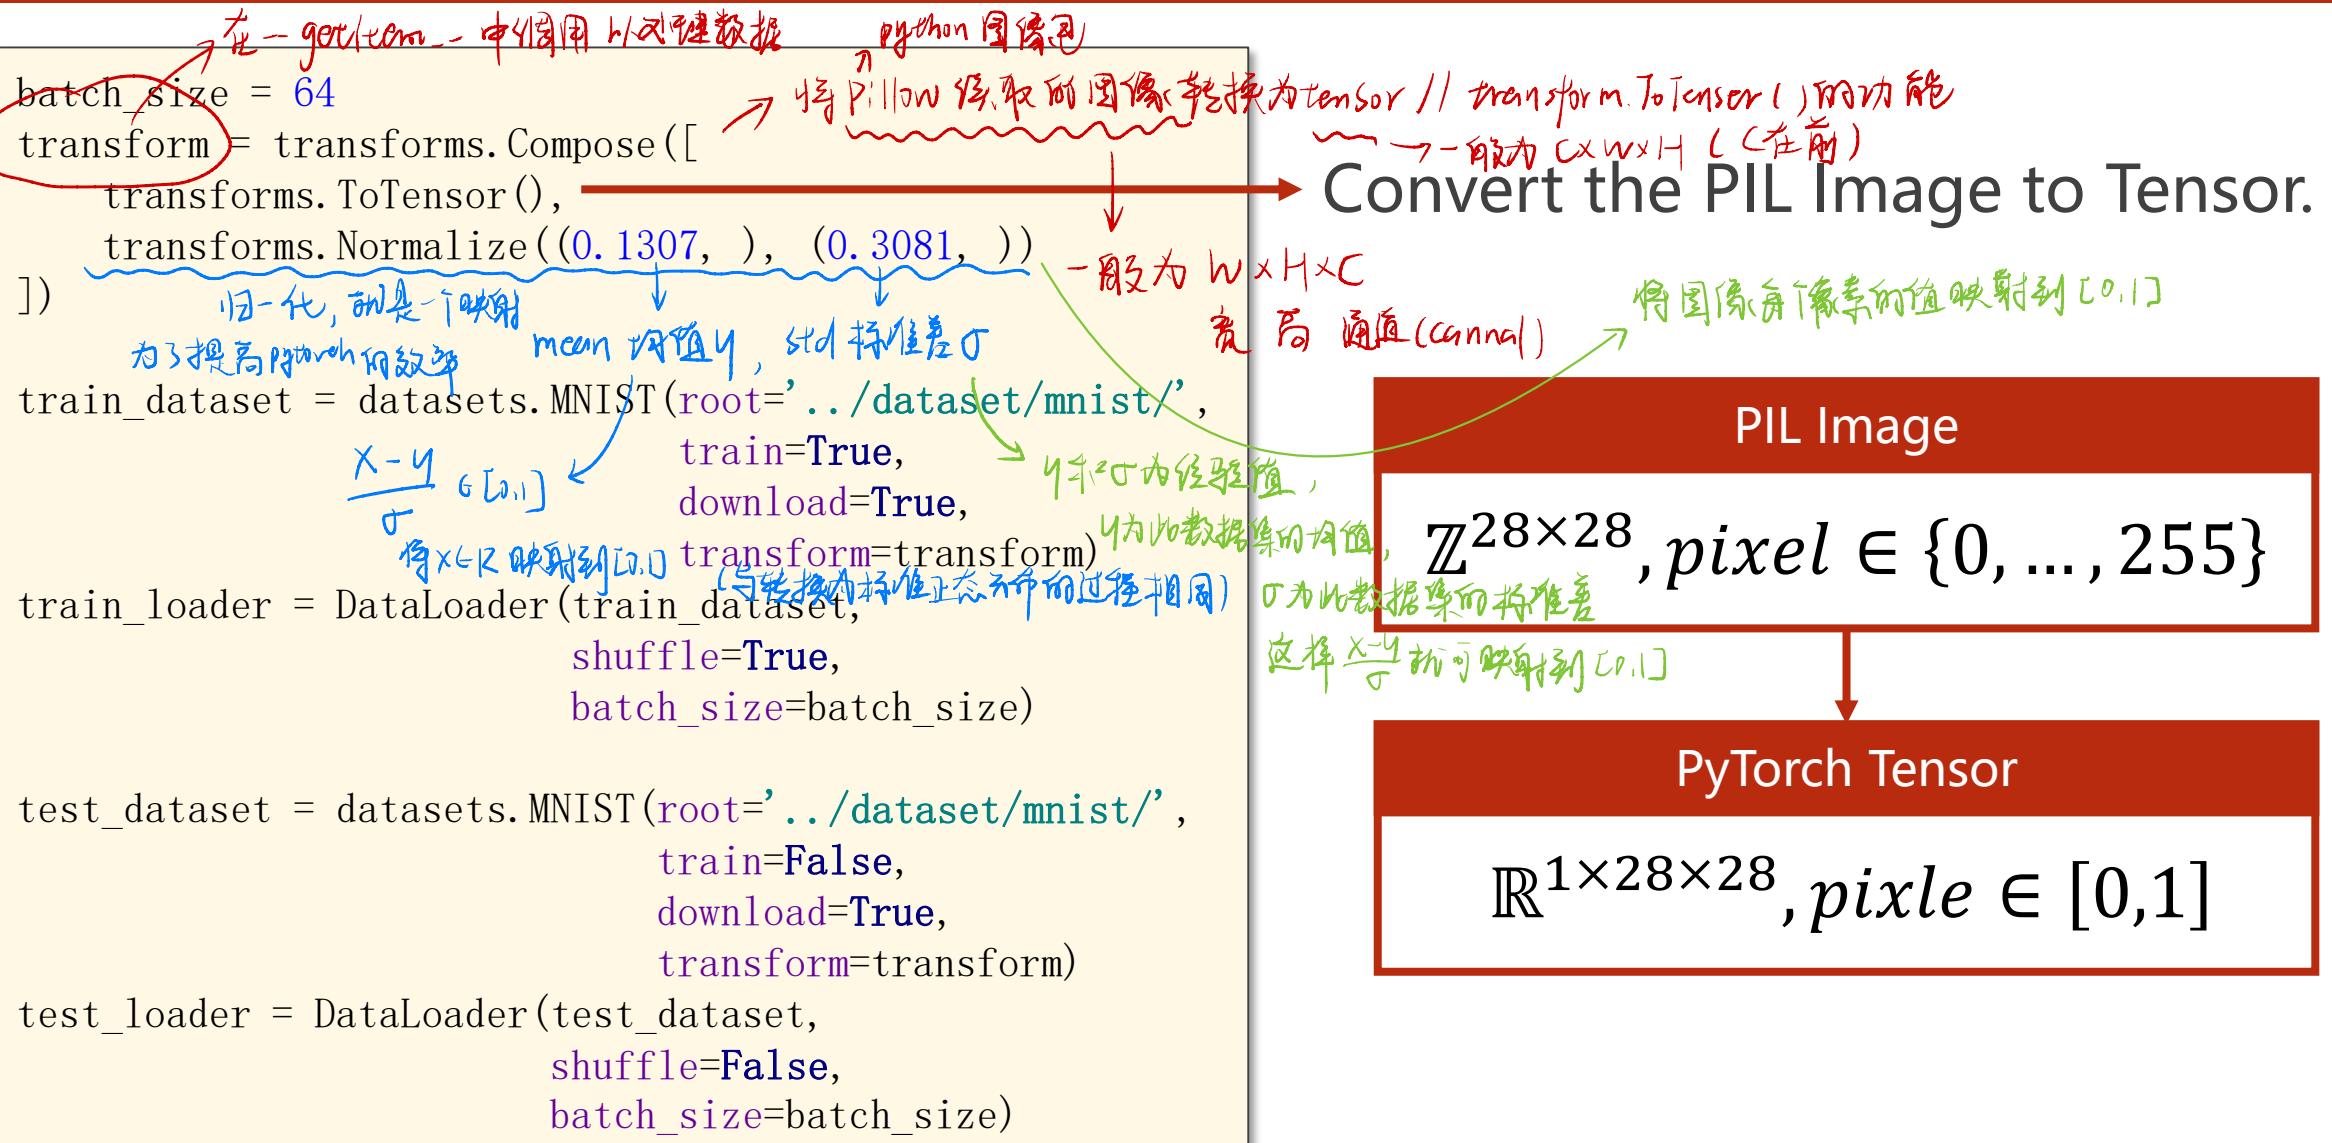
For using function relu()

Implementation – 0. Import Package

```
import torch
from torchvision import transforms
from torchvision import datasets
from torch.utils.data import DataLoader
import torch.nn.functional as F
import torch.optim as optim
```

For constructing Optimizer

Implementation – 1. Prepare Dataset



Implementation – 1. Prepare Dataset

```
batch_size = 64
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307, ), (0.3081, )))
])

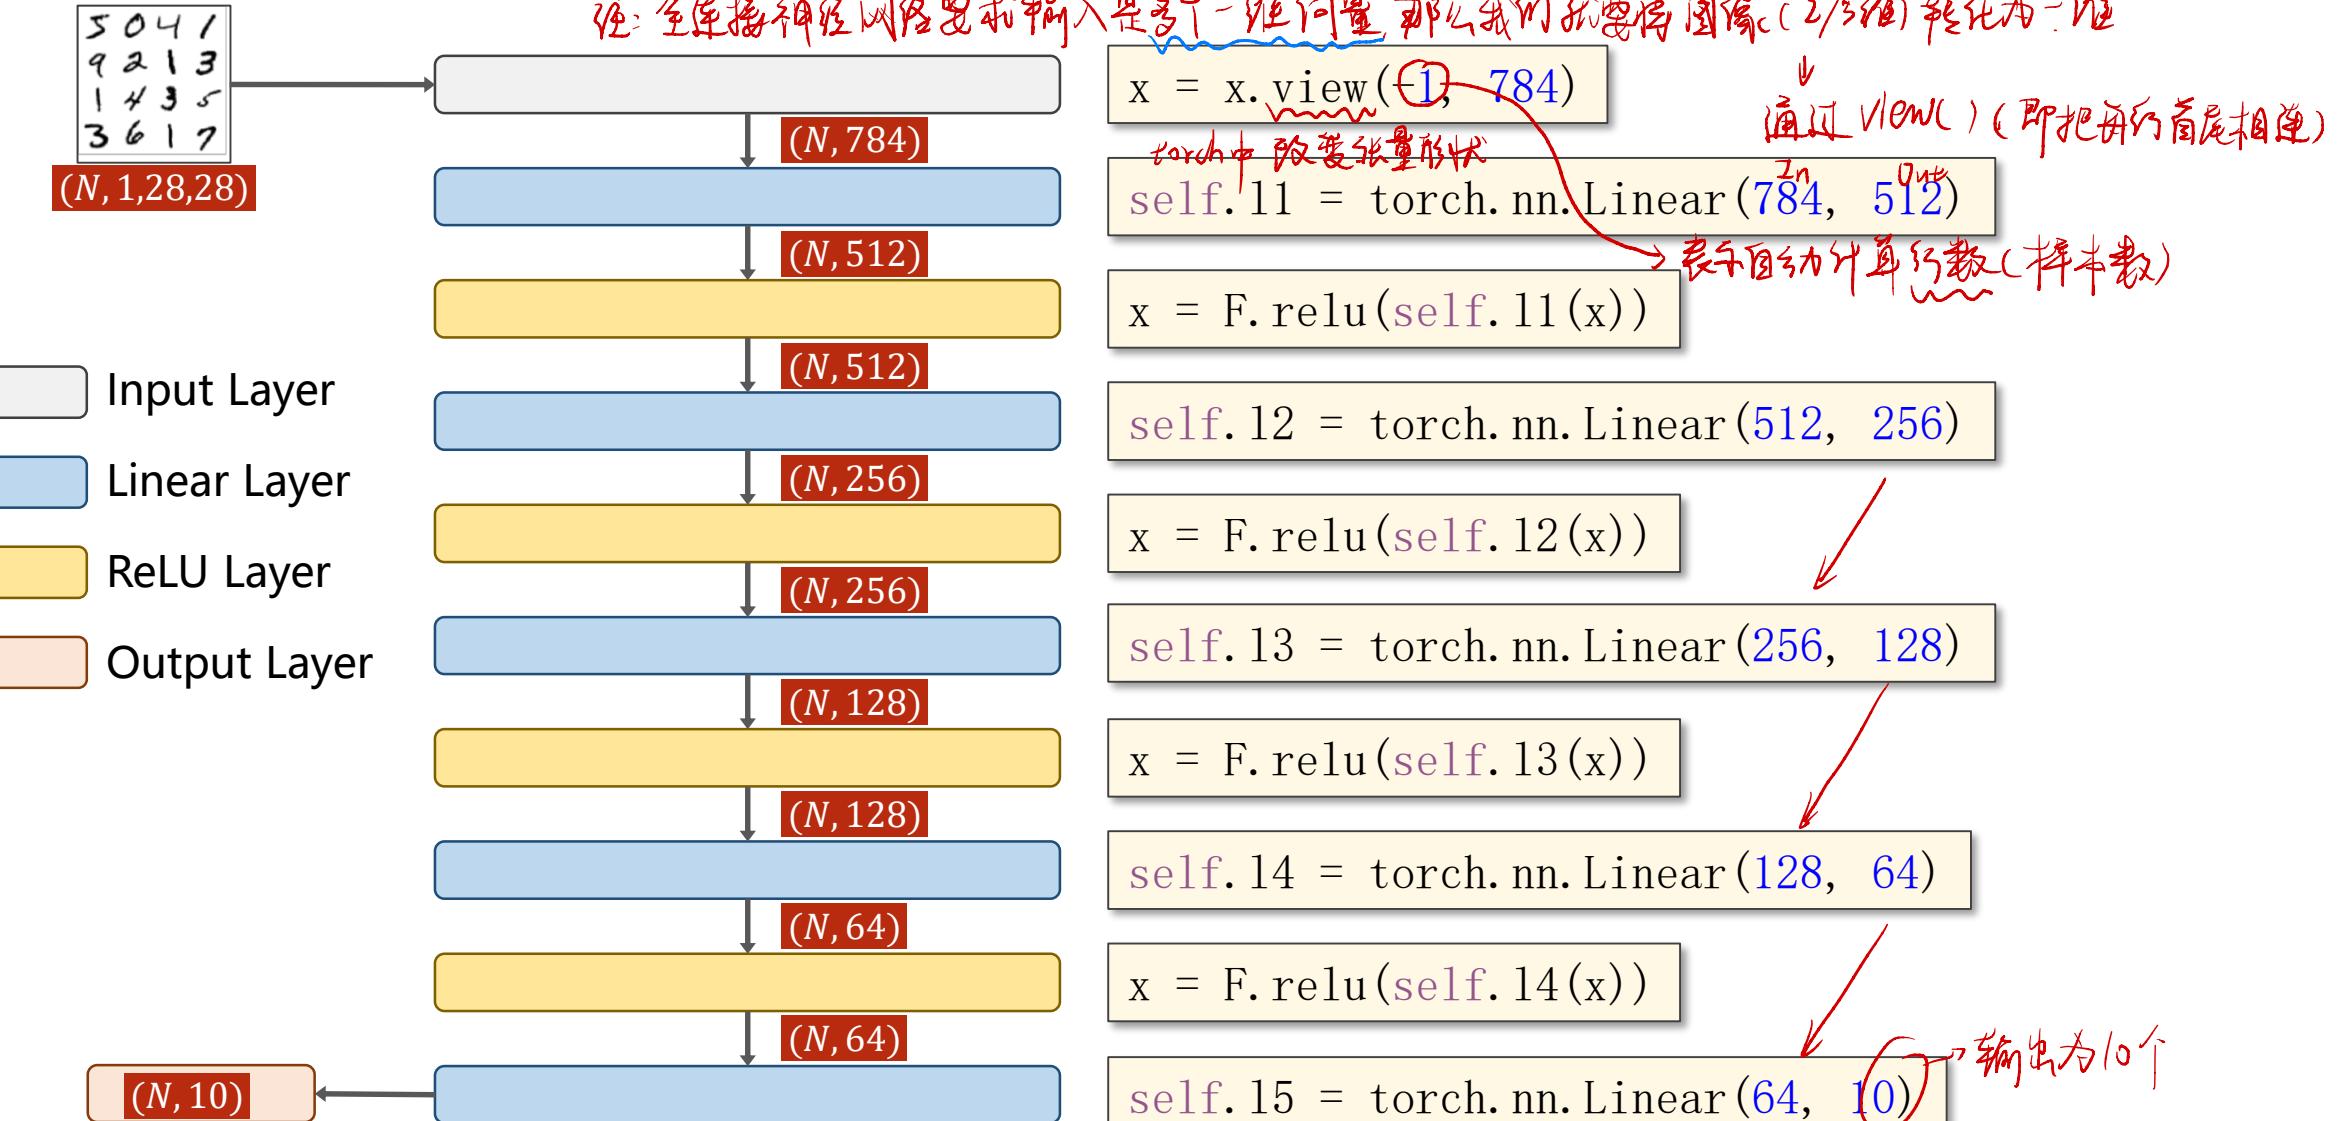
train_dataset = datasets.MNIST(root='../../dataset/mnist/',
                               train=True,
                               download=True,
                               transform=transform)
train_loader = DataLoader(train_dataset,
                          shuffle=True,
                          batch_size=batch_size)

test_dataset = datasets.MNIST(root='../../dataset/mnist/',
                              train=False,
                              download=True,
                              transform=transform)
test_loader = DataLoader(test_dataset,
                        shuffle=False,
                        batch_size=batch_size)
```

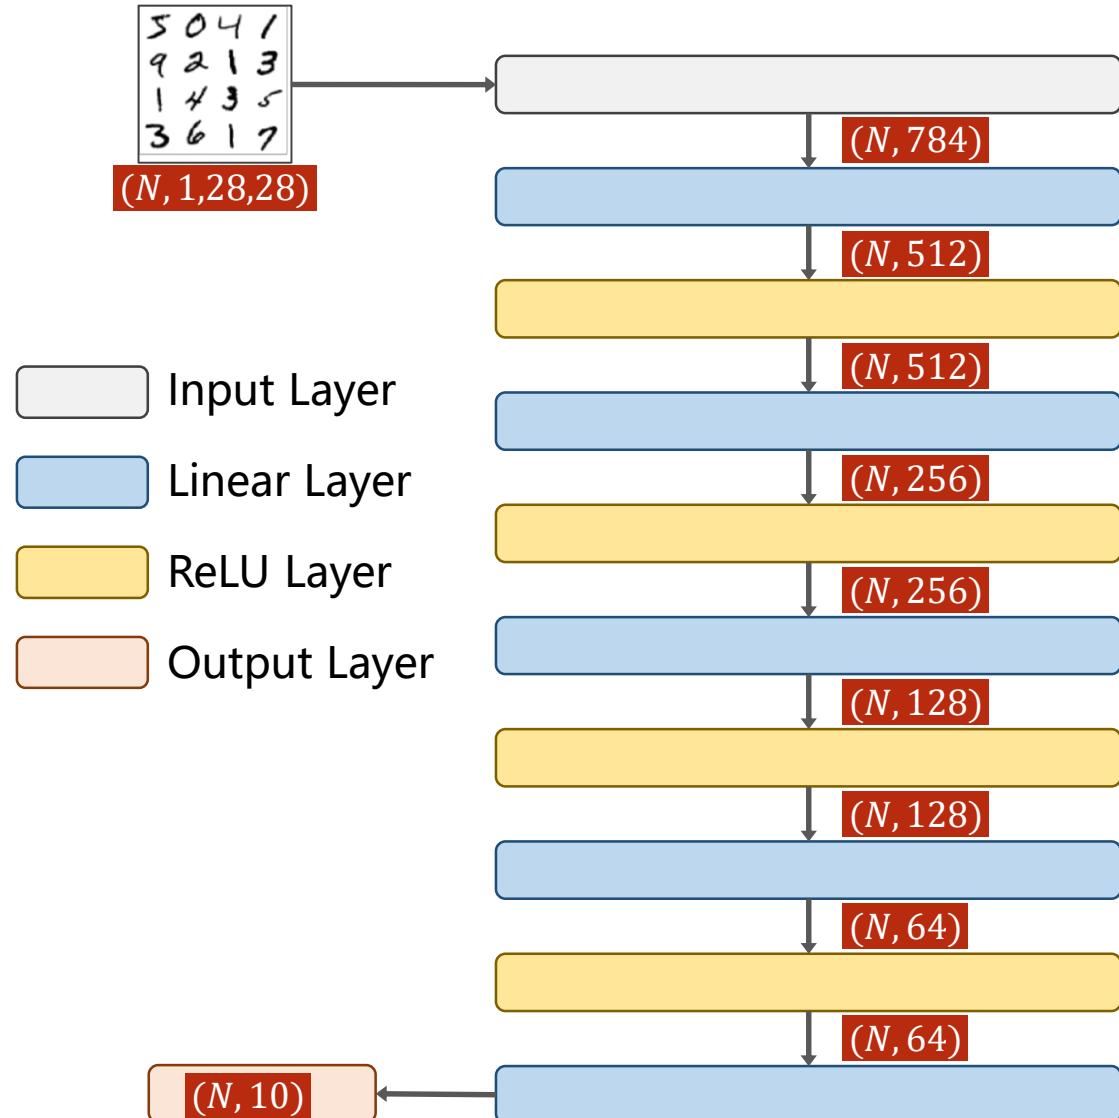
The parameters are **mean** and **std** respectively. It use formulation below:

$$Pixel_{norm} = \frac{Pixel_{origin} - mean}{std}$$

Implementation – 2. Design Model



Implementation – 2. Design Model



Input Layer

Linear Layer

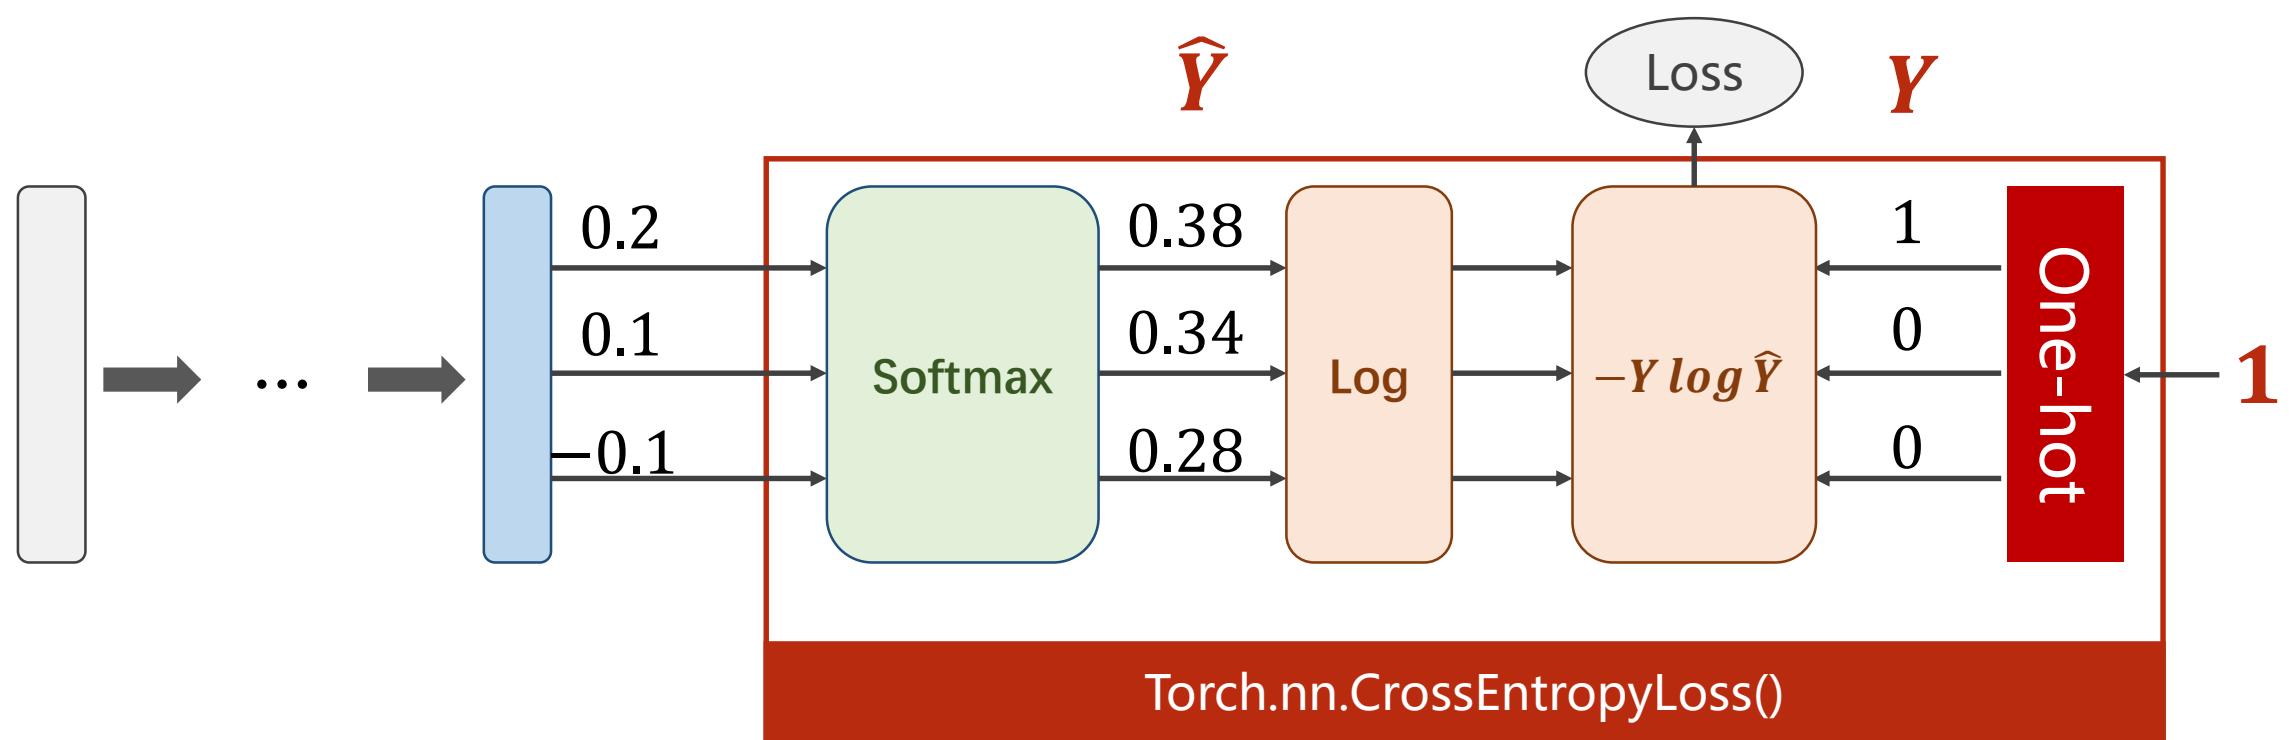
ReLU Layer

Output Layer

```
class Net(torch.nn.Module):  
    def __init__(self):  
        super(Net, self).__init__()  
        self.11 = torch.nn.Linear(784, 512)  
        self.12 = torch.nn.Linear(512, 256)  
        self.13 = torch.nn.Linear(256, 128)  
        self.14 = torch.nn.Linear(128, 64)  
        self.15 = torch.nn.Linear(64, 10)  
  
    def forward(self, x):  
        x = x.view(-1, 784) // 只有在--init--中声明了才可以在forward  
        x = F.relu(self.11(x)) // 可以调用  
        x = F.relu(self.12(x)) // 用relu对每一层的输出  
        x = F.relu(self.13(x)) // 没有  
        x = F.relu(self.14(x)) // 最后一层不作激活  
        return self.15(x)  
  
model = Net()
```

Implementation – 3. Construct Loss and Optimizer

```
criterion = torch.nn.CrossEntropyLoss()  
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.5)
```



Implementation – 4. Train and Test

```
def train(epoch): // 把一轮循环封装到函数中
    running_loss = 0.0
    for batch_idx, data in enumerate(train_loader, 0):
        inputs, target = data
        optimizer.zero_grad() # 0 grad
        # 1 forward + backward + update
        outputs = model(inputs)
        loss = criterion(outputs, target)
        2 - loss.backward()
        3 - optimizer.step()

        running_loss += loss.item()
        if batch_idx % 300 == 299:
            print(' [%d, %5d] loss: %.3f' % (epoch + 1, batch_idx + 1, running_loss / 300))
            running_loss = 0.0
```

每300次打印一次

Implementation – 4. Train and Test

```
def train(epoch):
    running_loss = 0.0
    for batch_idx, data in enumerate(train_loader, 0):
        inputs, target = data
        optimizer.zero_grad()

        # forward + backward + update
        outputs = model(inputs)
        loss = criterion(outputs, target)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        if batch_idx % 300 == 299:
            print(' [%d, %5d] loss: %.3f' % (epoch + 1, batch_idx + 1, running_loss / 300))
            running_loss = 0.0
```

Implementation – 4. Train and Test

```
def train(epoch):
    running_loss = 0.0
    for batch_idx, data in enumerate(train_loader, 0):
        inputs, target = data
        optimizer.zero_grad()

        # forward + backward + update
        outputs = model(inputs)
        loss = criterion(outputs, target)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        if batch_idx % 300 == 299:
            print(' [%d, %5d] loss: %.3f' % (epoch + 1, batch_idx + 1, running_loss / 300))
            running_loss = 0.0
```

Implementation – 4. Train and Test

```
def train(epoch):
    running_loss = 0.0
    for batch_idx, data in enumerate(train_loader, 0):
        inputs, target = data
        optimizer.zero_grad()

        # forward + backward + update
        outputs = model(inputs)
        loss = criterion(outputs, target)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() //不打印计算图
        if batch_idx % 300 == 299:
            print(' [%d, %5d] loss: %.3f' % (epoch + 1, batch_idx + 1, running_loss / 300))
            running_loss = 0.0
```

Implementation – 4. Train and Test

测试：

```
def test():
    correct = 0
    total = 0
    with torch.no_grad():
        for data in test_loader: batch
            images, labels = data
            outputs = model(images)
            _, predicted = torch.max(outputs.data, dim=1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
    print('Accuracy on test set: %d %%' % (100 * correct / total))
```

不计算梯度

Implementation – 4. Train and Test

```
def test():
    correct = 0
    total = 0
    with torch.no_grad():
        for data in test_loader:
            images, labels = data
            outputs = model(images)
            _, predicted = torch.max(outputs.data, dim=1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
    print('Accuracy on test set: %d %%' % (100 * correct / total))
```

Implementation – 4. Train and Test

```
def test():
    correct = 0
    total = 0
    with torch.no_grad():
        for data in test_loader:
            images, labels = data
            outputs = model(images)
            _, predicted = torch.max(outputs.data, dim=1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
    print('Accuracy on test set: %d %%' % (100 * correct / total))
```

[$\frac{0}{N}$] , 即 N 个样本的 label.
label, $N \times 1$ 矩阵 (2 维向量) label.size 那为 $(N, 1)$, label.size[0] = N
输出是一矩阵 [三], 每一行是一个结果
每行结果有 10 个数
(10 个类别预测)
样本数
正确数
打印
不是总量
每行
返回最大值的下标
即输出最大值.

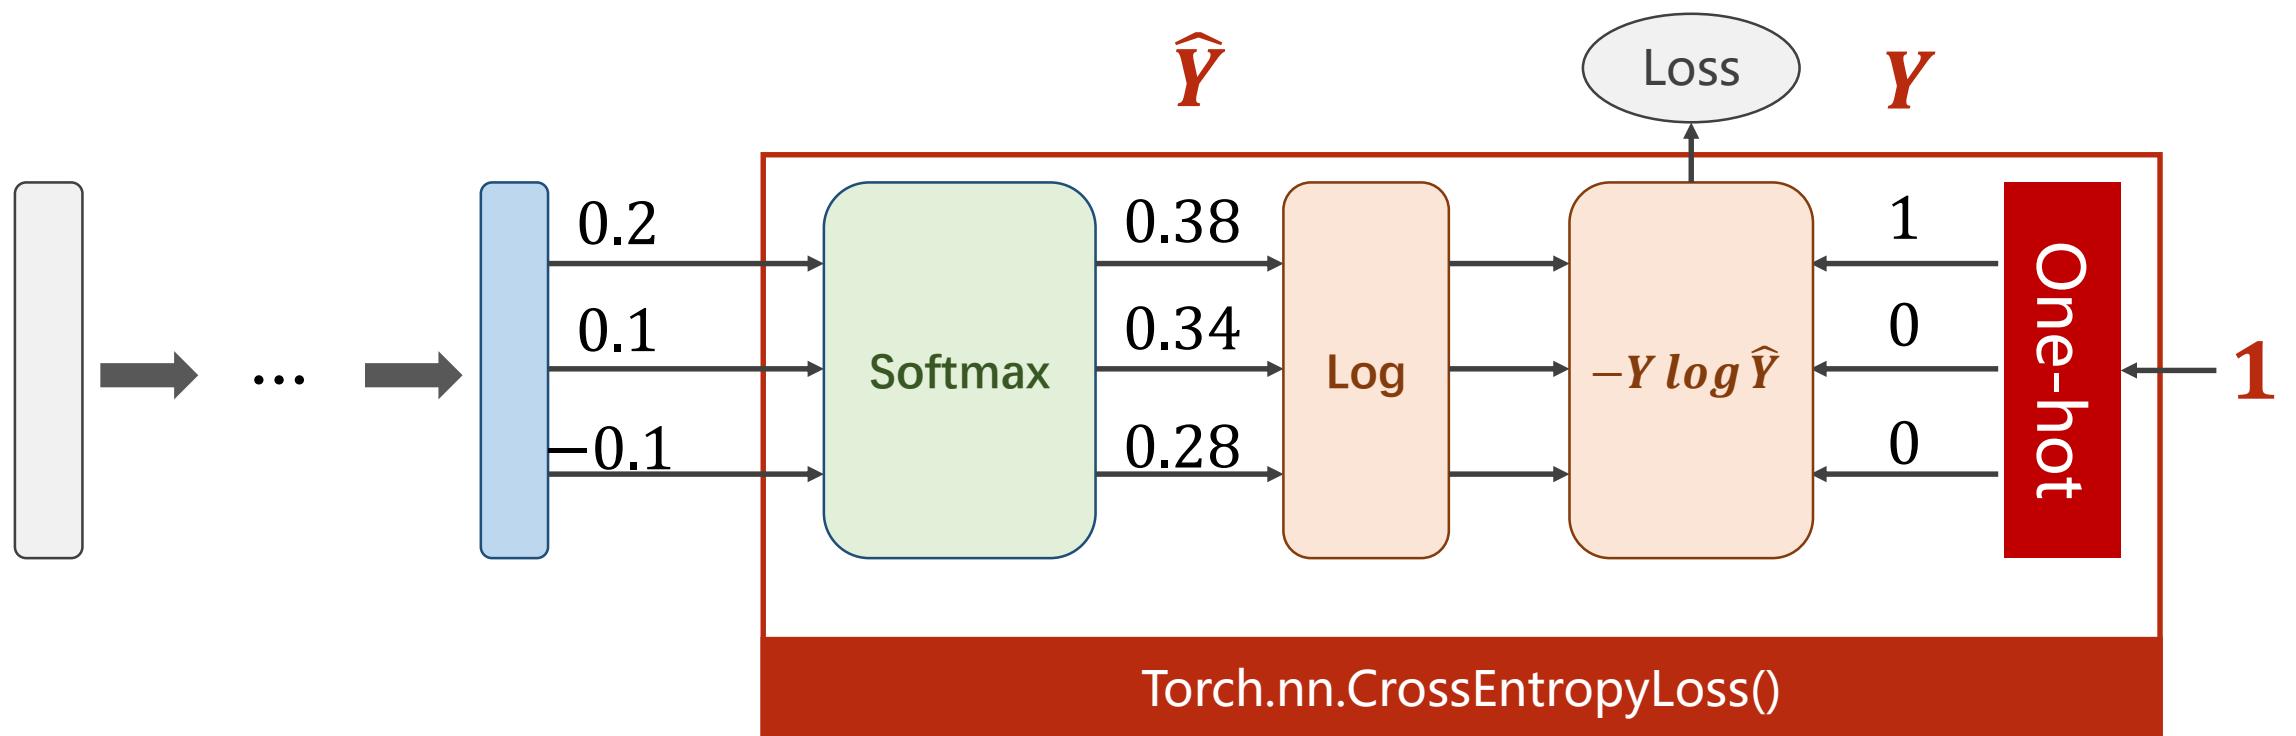
Implementation – 4. Train and Test

```
if __name__ == '__main__':
    for epoch in range(10):
        train(epoch)  # 需训练的数据
        test()
```

全连接：把所有数据（有用 and 不用）都用神经网络
特征提取，只用特征训练 CNN
~~~~~ {  
    AI: FFT, 小图 ~  
    运动: CNN

```
[1, 300] loss: 0.335
[1, 600] loss: 0.154
[1, 900] loss: 0.067
Accuracy on test set: 90 %
[2, 300] loss: 0.048
[2, 600] loss: 0.040
[2, 900] loss: 0.035
Accuracy on test set: 93 %
.....
[9, 300] loss: 0.005
[9, 600] loss: 0.006
[9, 900] loss: 0.007
Accuracy on test set: 97 %
[10, 300] loss: 0.005
[10, 600] loss: 0.005
[10, 900] loss: 0.005
Accuracy on test set: 97 %
```

# Softmax and CrossEntropyLoss



# Exercise 9-2: Classifier Implementation

- Try to implement a classifier for:
  - Otto Group Product Classification Challenge
  - Dataset: <https://www.kaggle.com/c/otto-group-product-classification-challenge/data>

The screenshot shows the 'Data' tab of a Kaggle competition page. At the top, there are tabs for Overview, Data, Kernels, Discussion, Leaderboard, Rules, and a blue 'Late Submission' button. Below the tabs, the word 'Data' is centered. To the right are buttons for API, ?, Download All, and a close icon. The main area is divided into three columns: 'Data Sources', 'About this file', and 'Columns'. The 'Data Sources' column lists 'sampleSubmission.... 144k x 10', 'test.csv 144k x 94', and 'train.csv 61.9k x 95', with 'train.csv' highlighted by a red box. The 'About this file' column says 'No description yet'. The 'Columns' column lists the features: # id, # feat\_1, # feat\_2, # feat\_3, # feat\_4, # feat\_5, and # feat\_6.

| Data Sources                   | About this file    | Columns  |
|--------------------------------|--------------------|----------|
| sampleSubmission.... 144k x 10 | No description yet | # id     |
| test.csv 144k x 94             |                    | # feat_1 |
| train.csv 61.9k x 95           |                    | # feat_2 |
|                                |                    | # feat_3 |
|                                |                    | # feat_4 |
|                                |                    | # feat_5 |
|                                |                    | # feat_6 |



# PyTorch Tutorial

## 09. Softmax Classifier



# PyTorch Tutorial

## 09. Softmax Classifier



# PyTorch Tutorial

## 09. Softmax Classifier