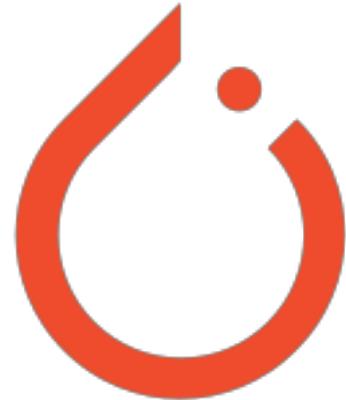


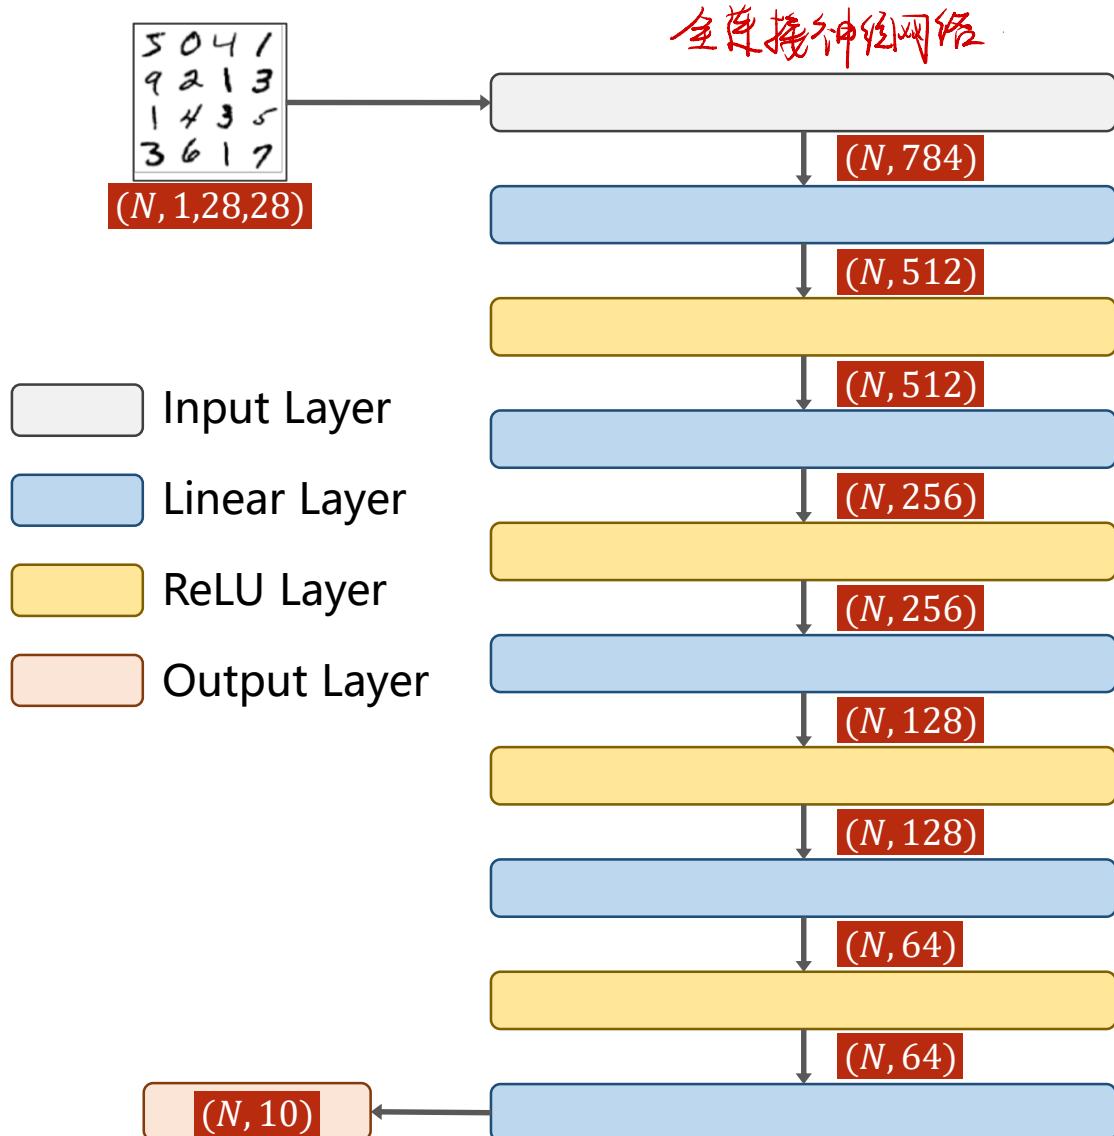
神经网络 {  
全连接神经网络  
卷积神经网络  
循环神经网络



# PyTorch Tutorial

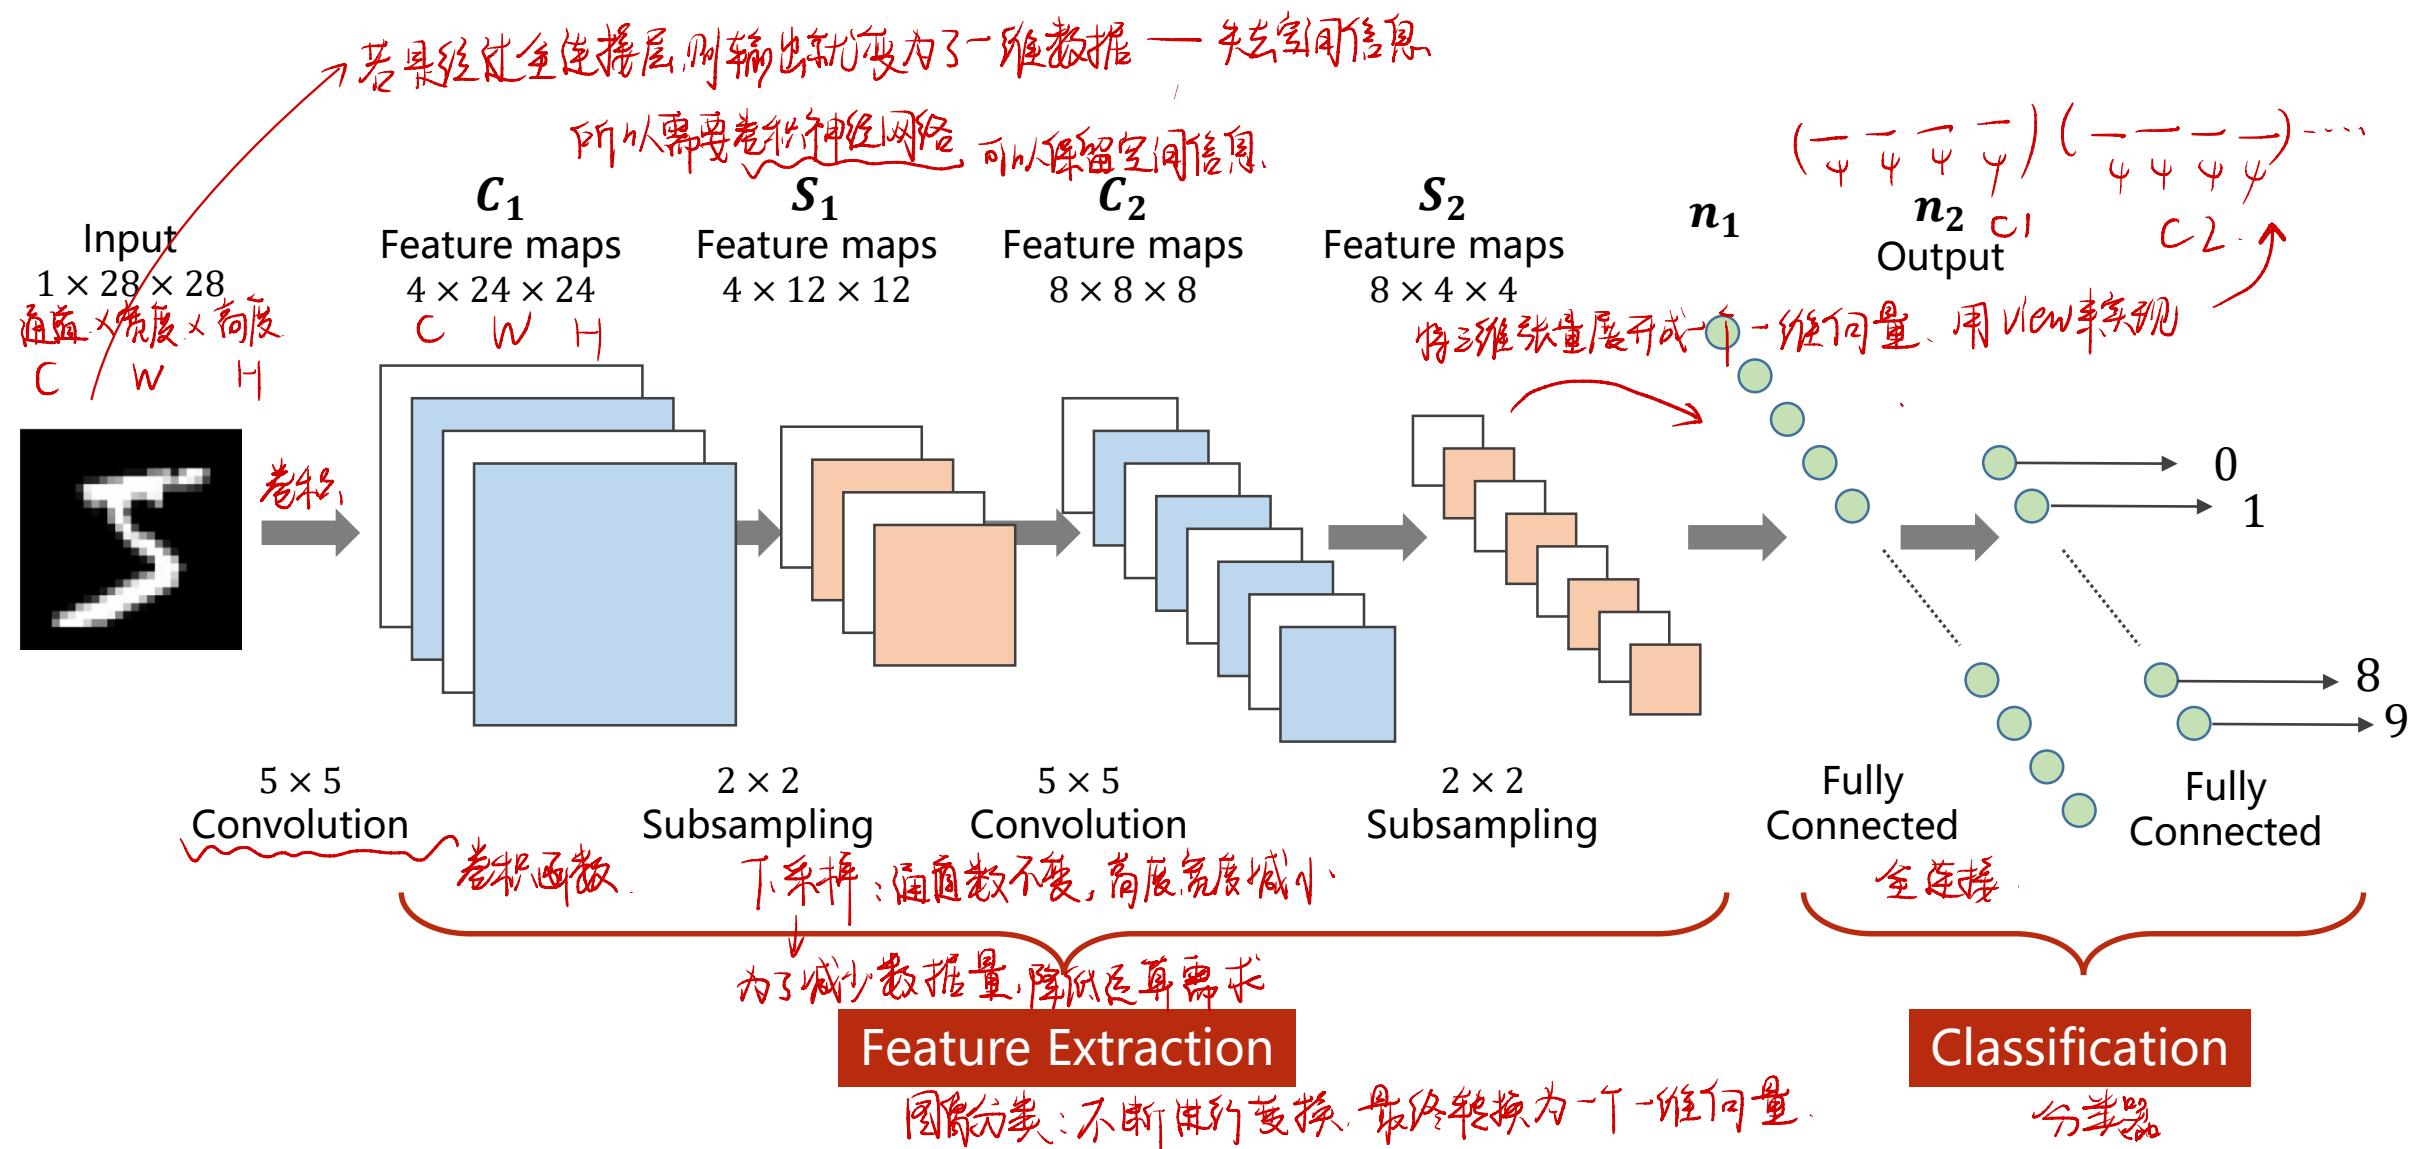
## 10. Basic CNN

# Revision: Fully Connected Neural Network

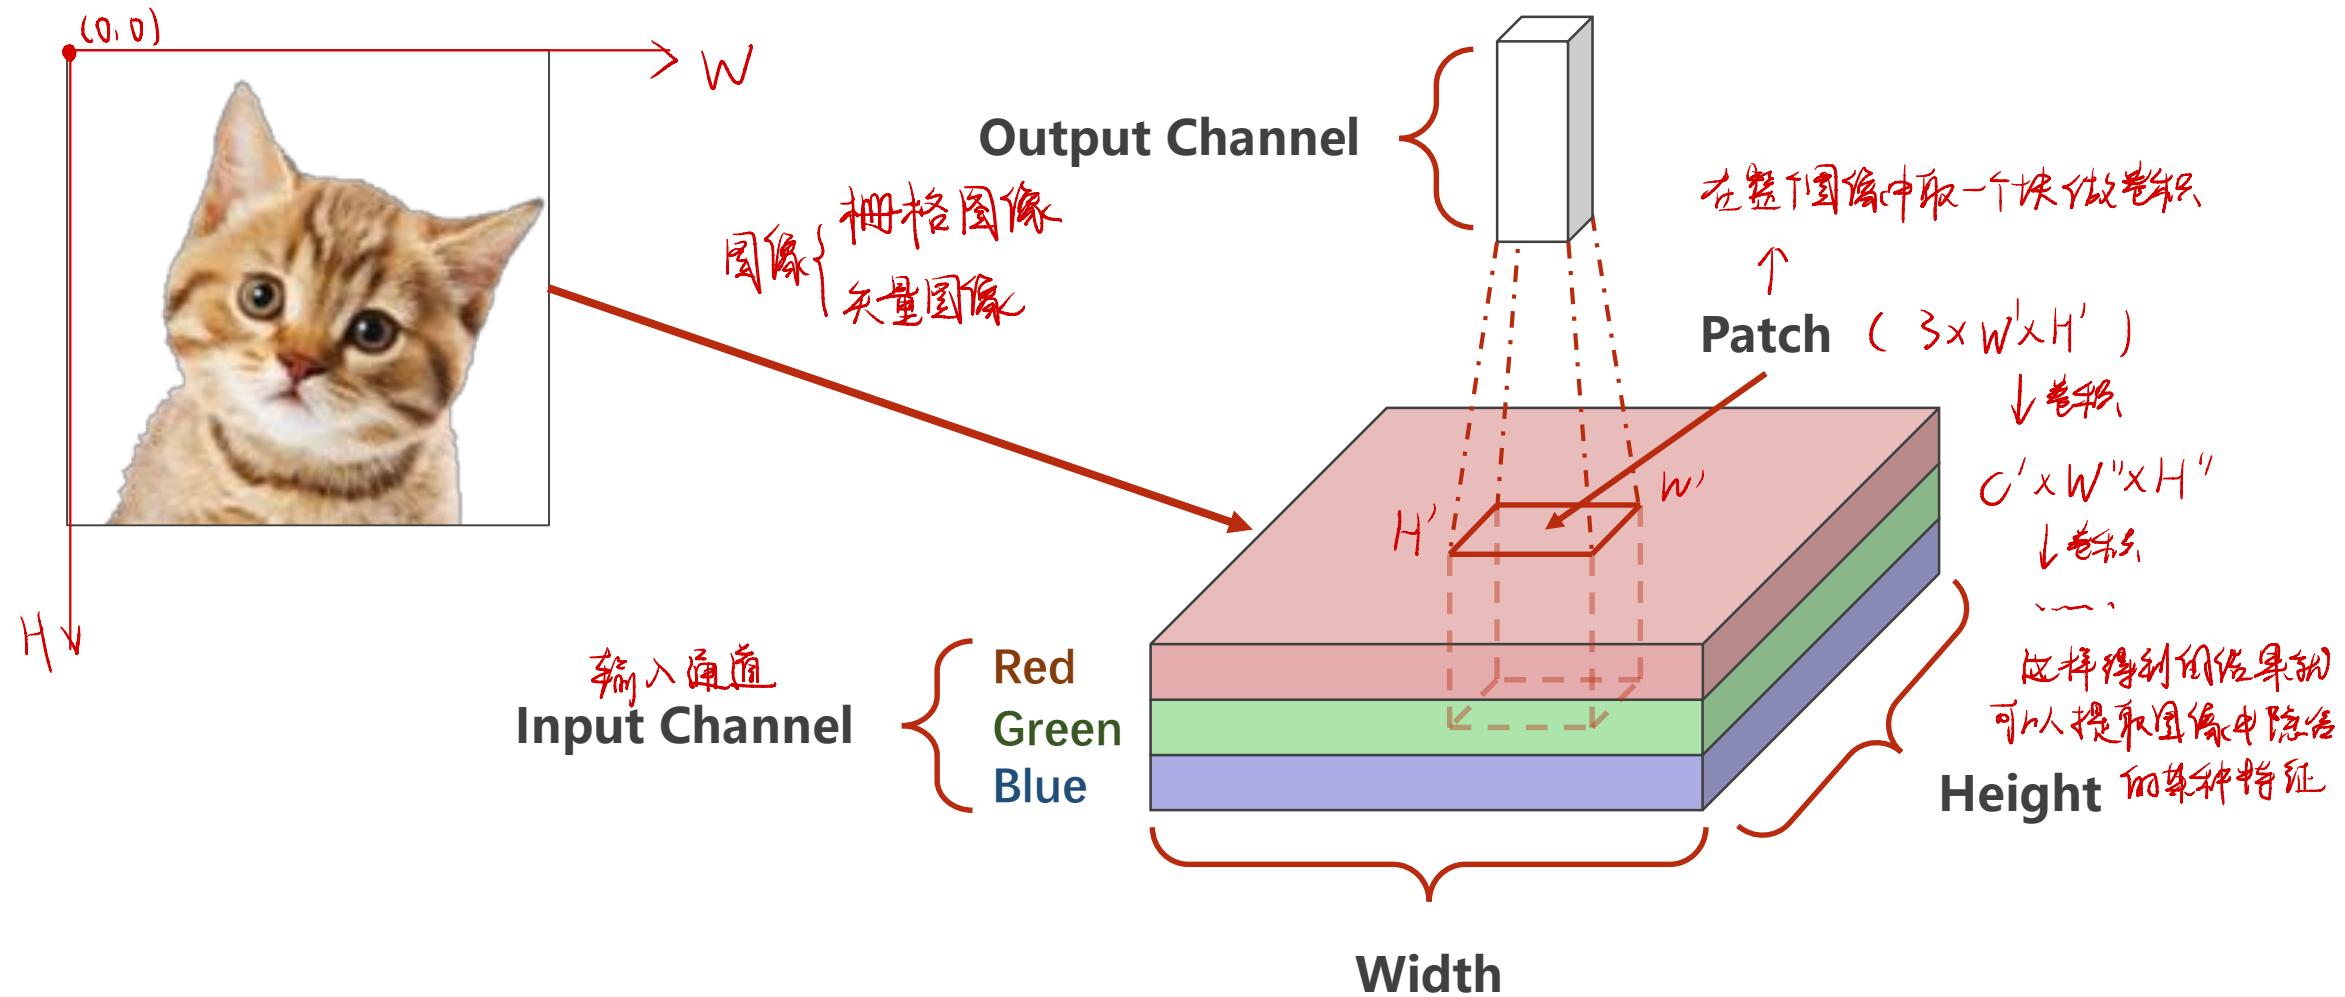


```
class Net(torch.nn.Module):  
    def __init__(self):  
        super(Net, self).__init__()  
        self.l1 = torch.nn.Linear(784, 512)  
        self.l2 = torch.nn.Linear(512, 256)  
        self.l3 = torch.nn.Linear(256, 128)  
        self.l4 = torch.nn.Linear(128, 64)  
        self.l5 = torch.nn.Linear(64, 10)  
  
    def forward(self, x):  
        x = x.view(-1, 784)  
        x = F.relu(self.l1(x))  
        x = F.relu(self.l2(x))  
        x = F.relu(self.l3(x))  
        x = F.relu(self.l4(x))  
        return self.l5(x)  
  
model = Net()
```

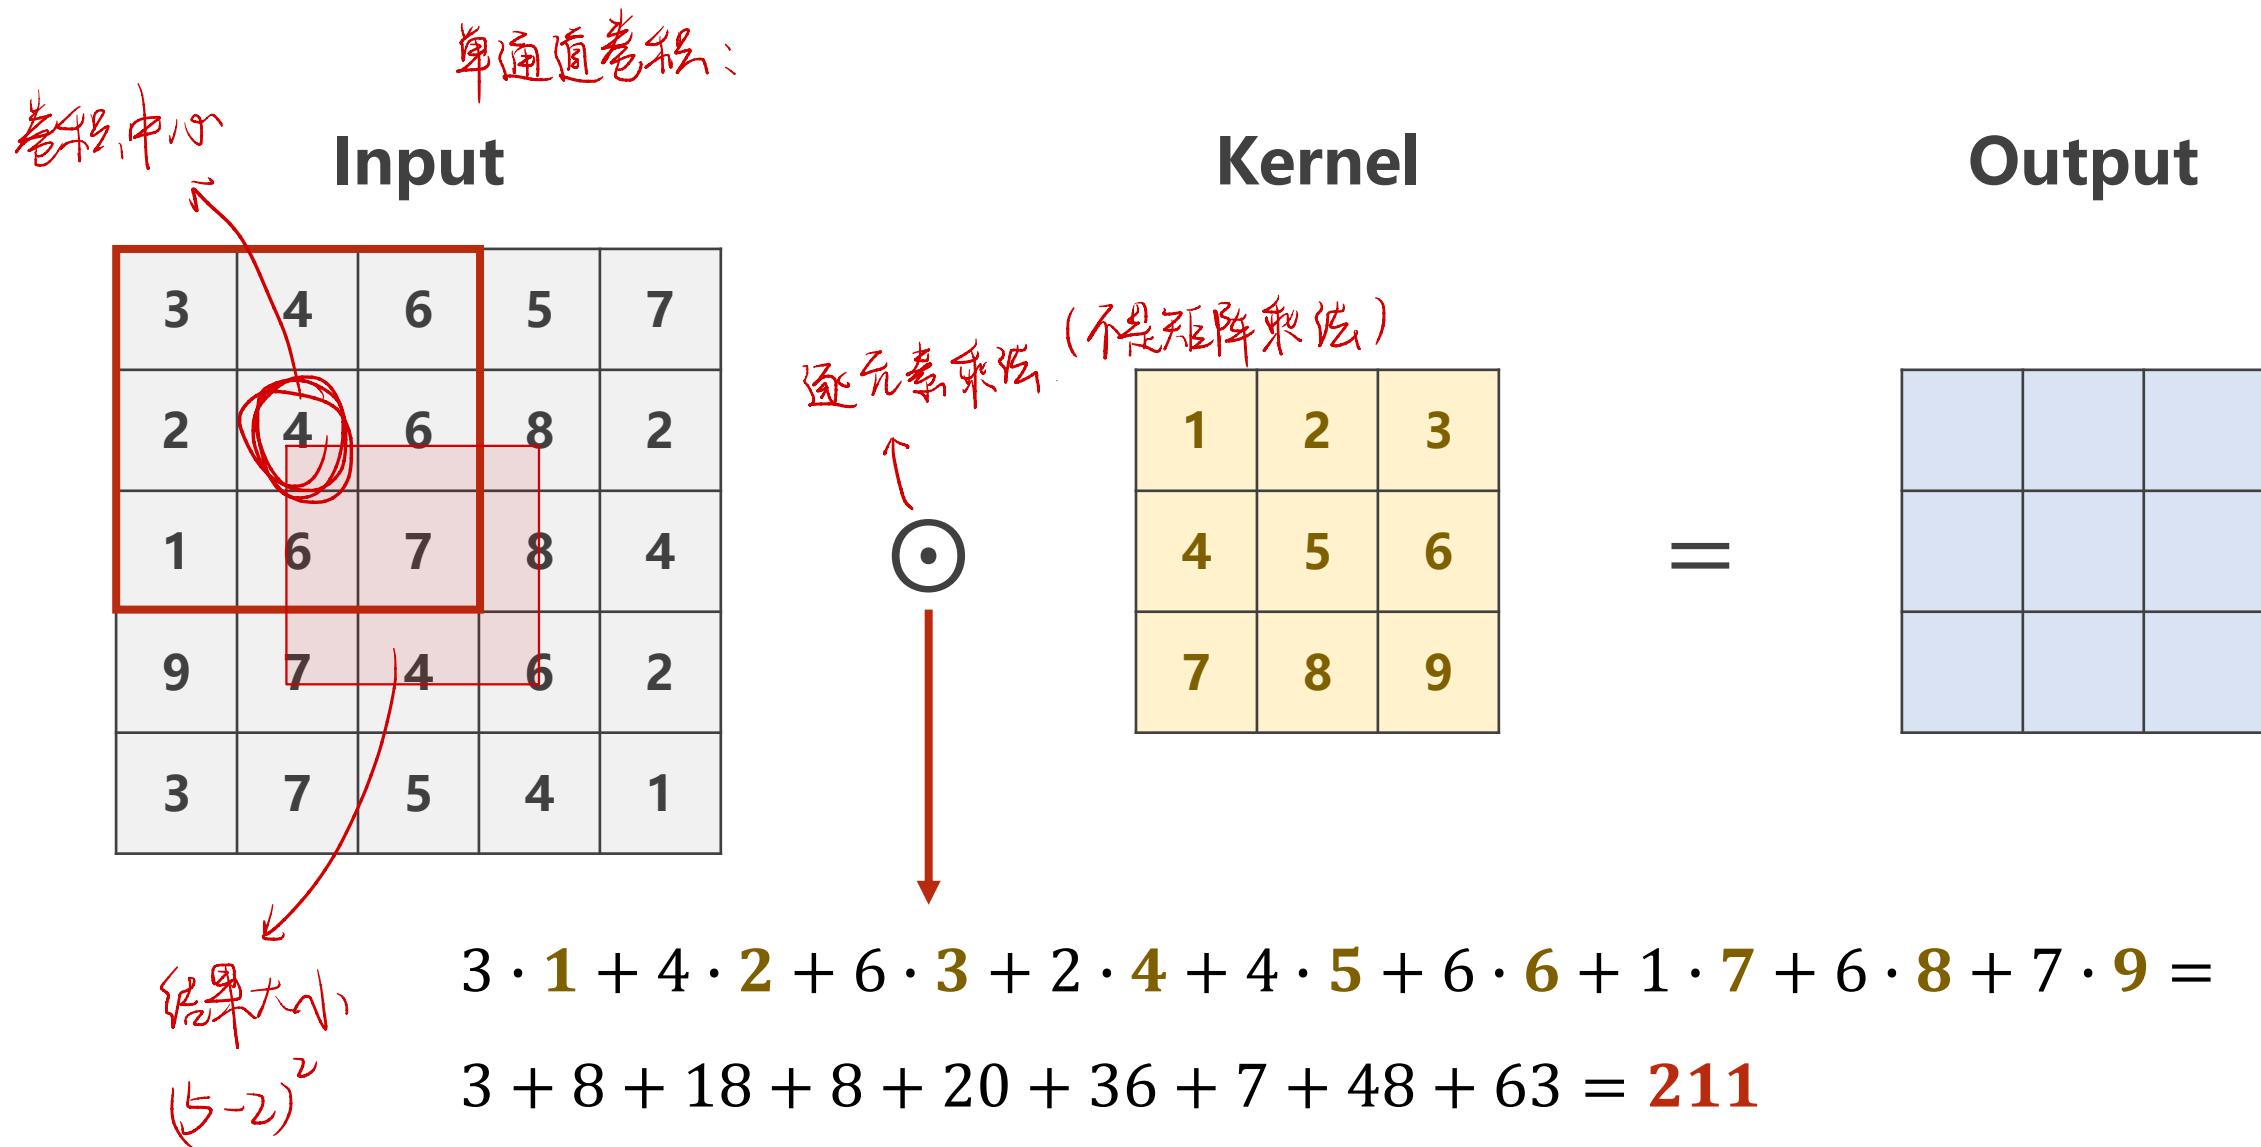
# Convolutional Neural Network



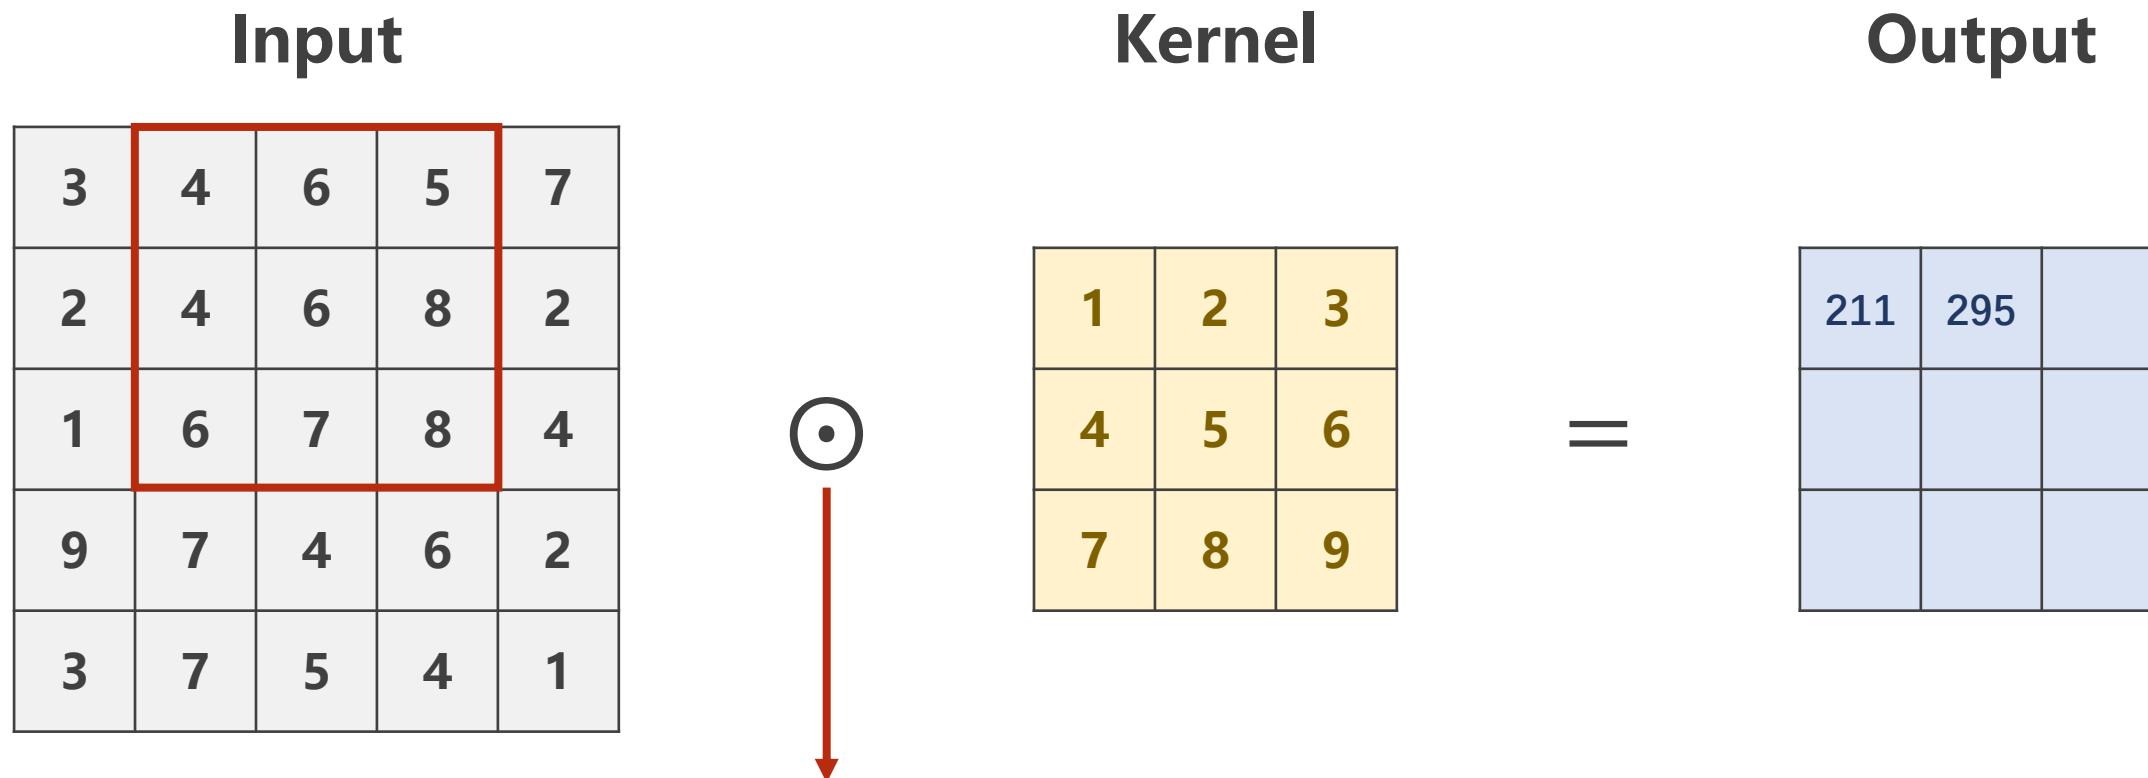
# Convolution



# Convolution – Single Input Channel



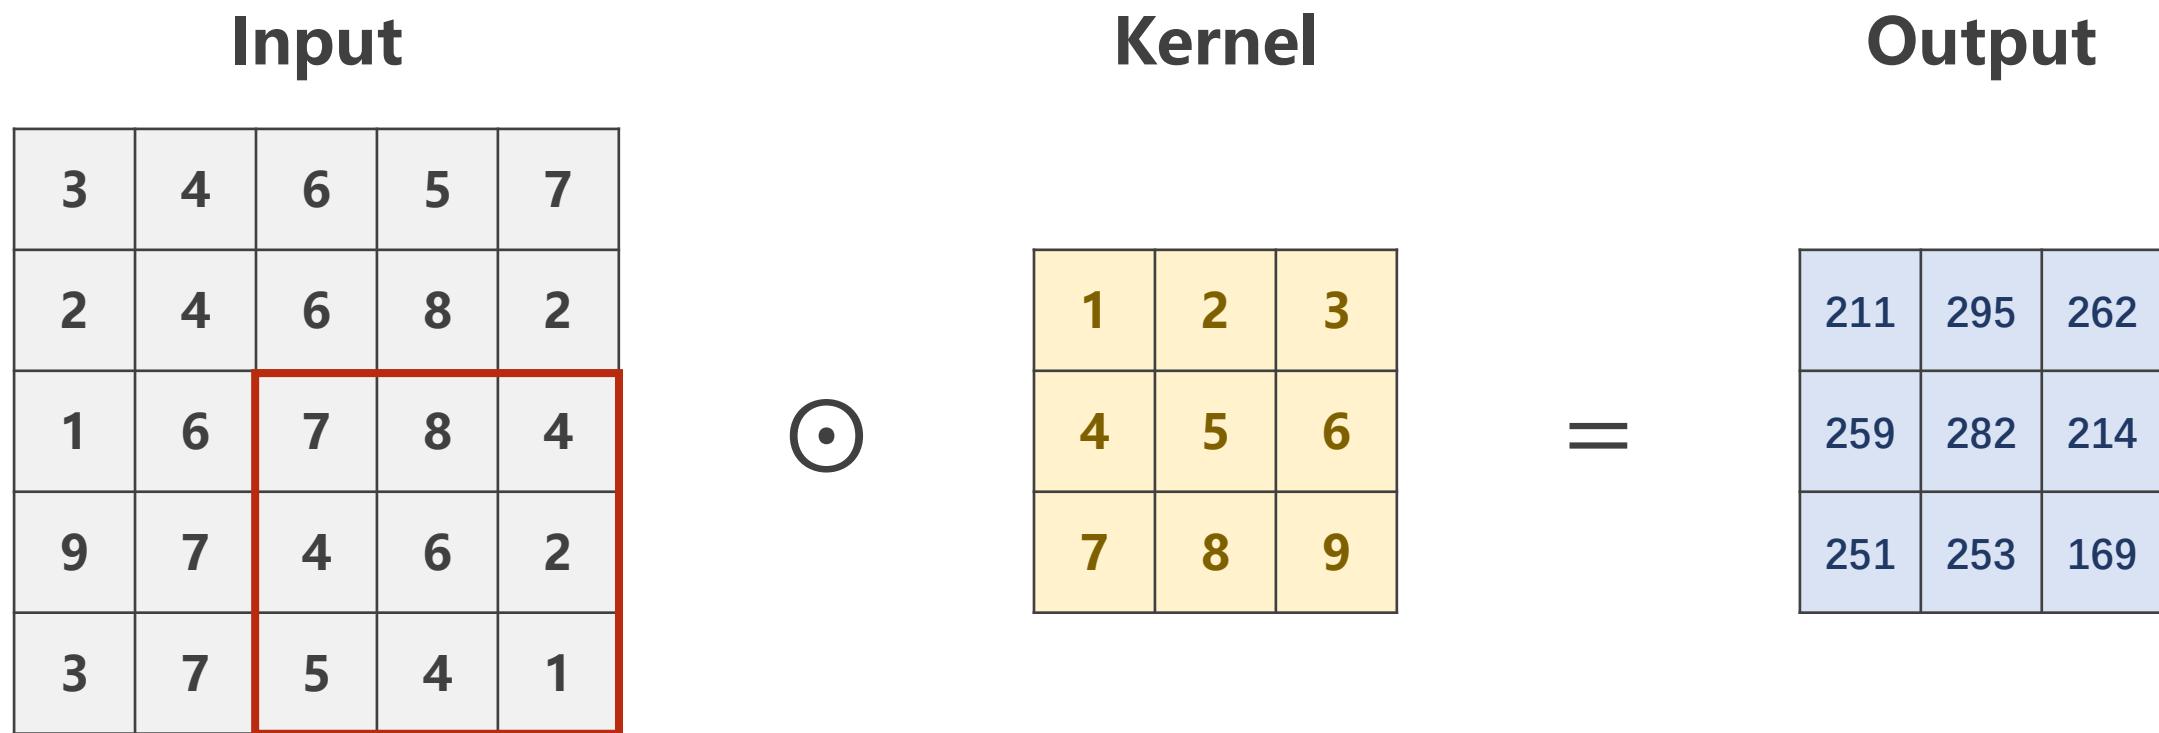
# Convolution – Single Input Channel



$$4 \cdot 1 + 6 \cdot 2 + 5 \cdot 3 + 4 \cdot 4 + 6 \cdot 5 + 8 \cdot 6 + 6 \cdot 7 + 7 \cdot 8 + 8 \cdot 9 =$$

$$4 + 12 + 15 + 16 + 30 + 48 + 42 + 56 + 72 = 295$$

# Convolution – Single Input Channel



# Convolution – 3 Input Channels

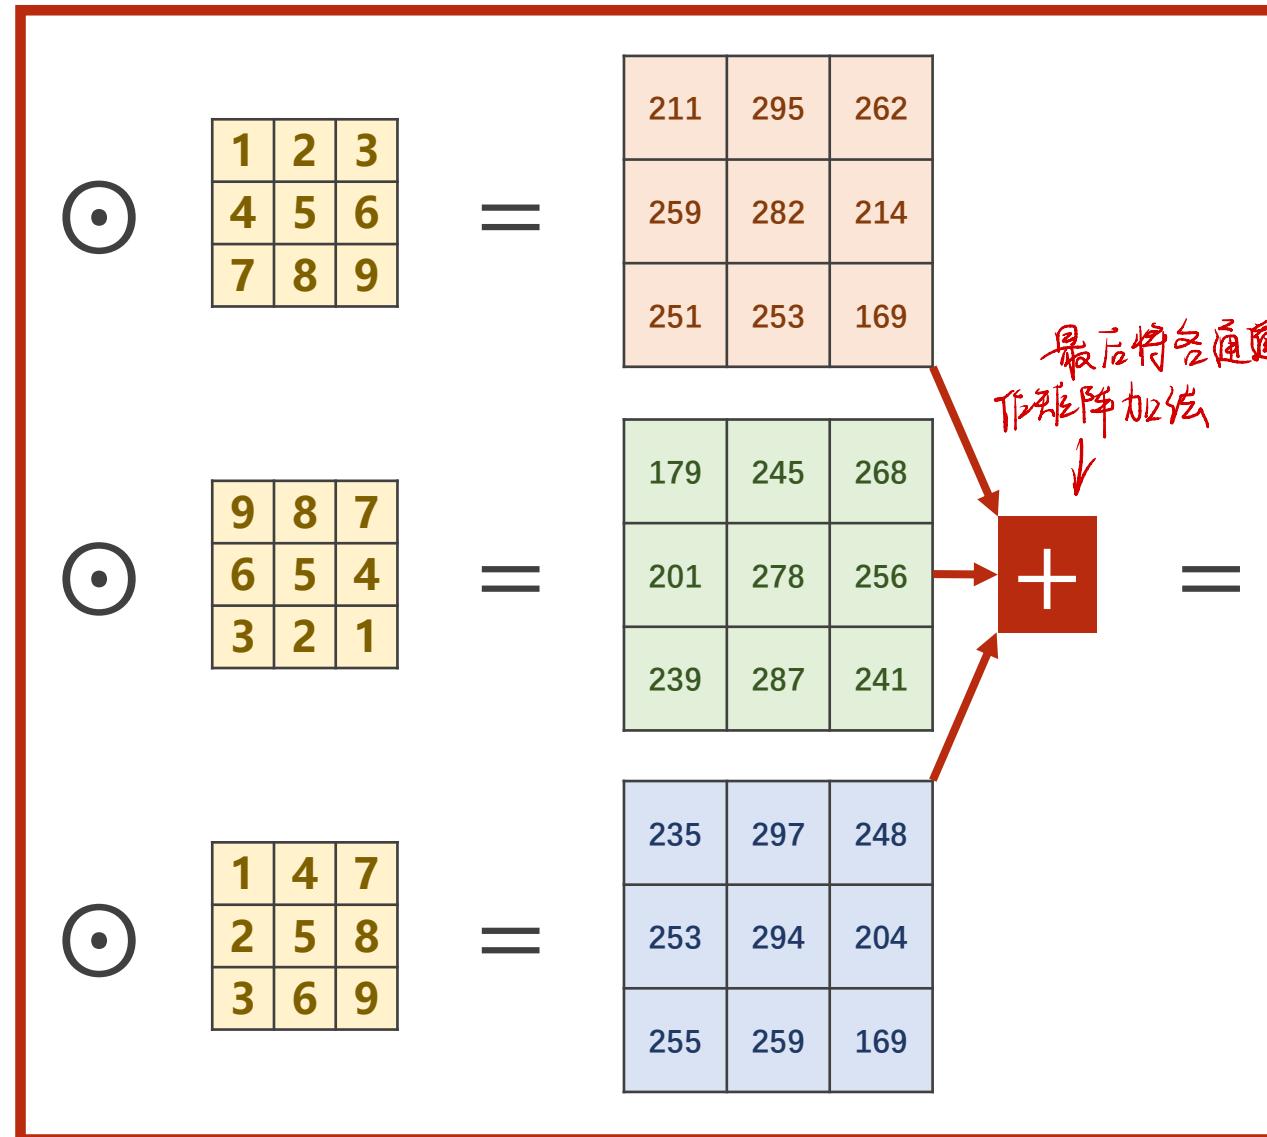
多通道卷积：

每一个通道都有  
单独的卷积核。

3	4	6	5	7
2	4	6	8	2
1	6	7	8	4
9	7	4	6	2
3	7	5	4	1

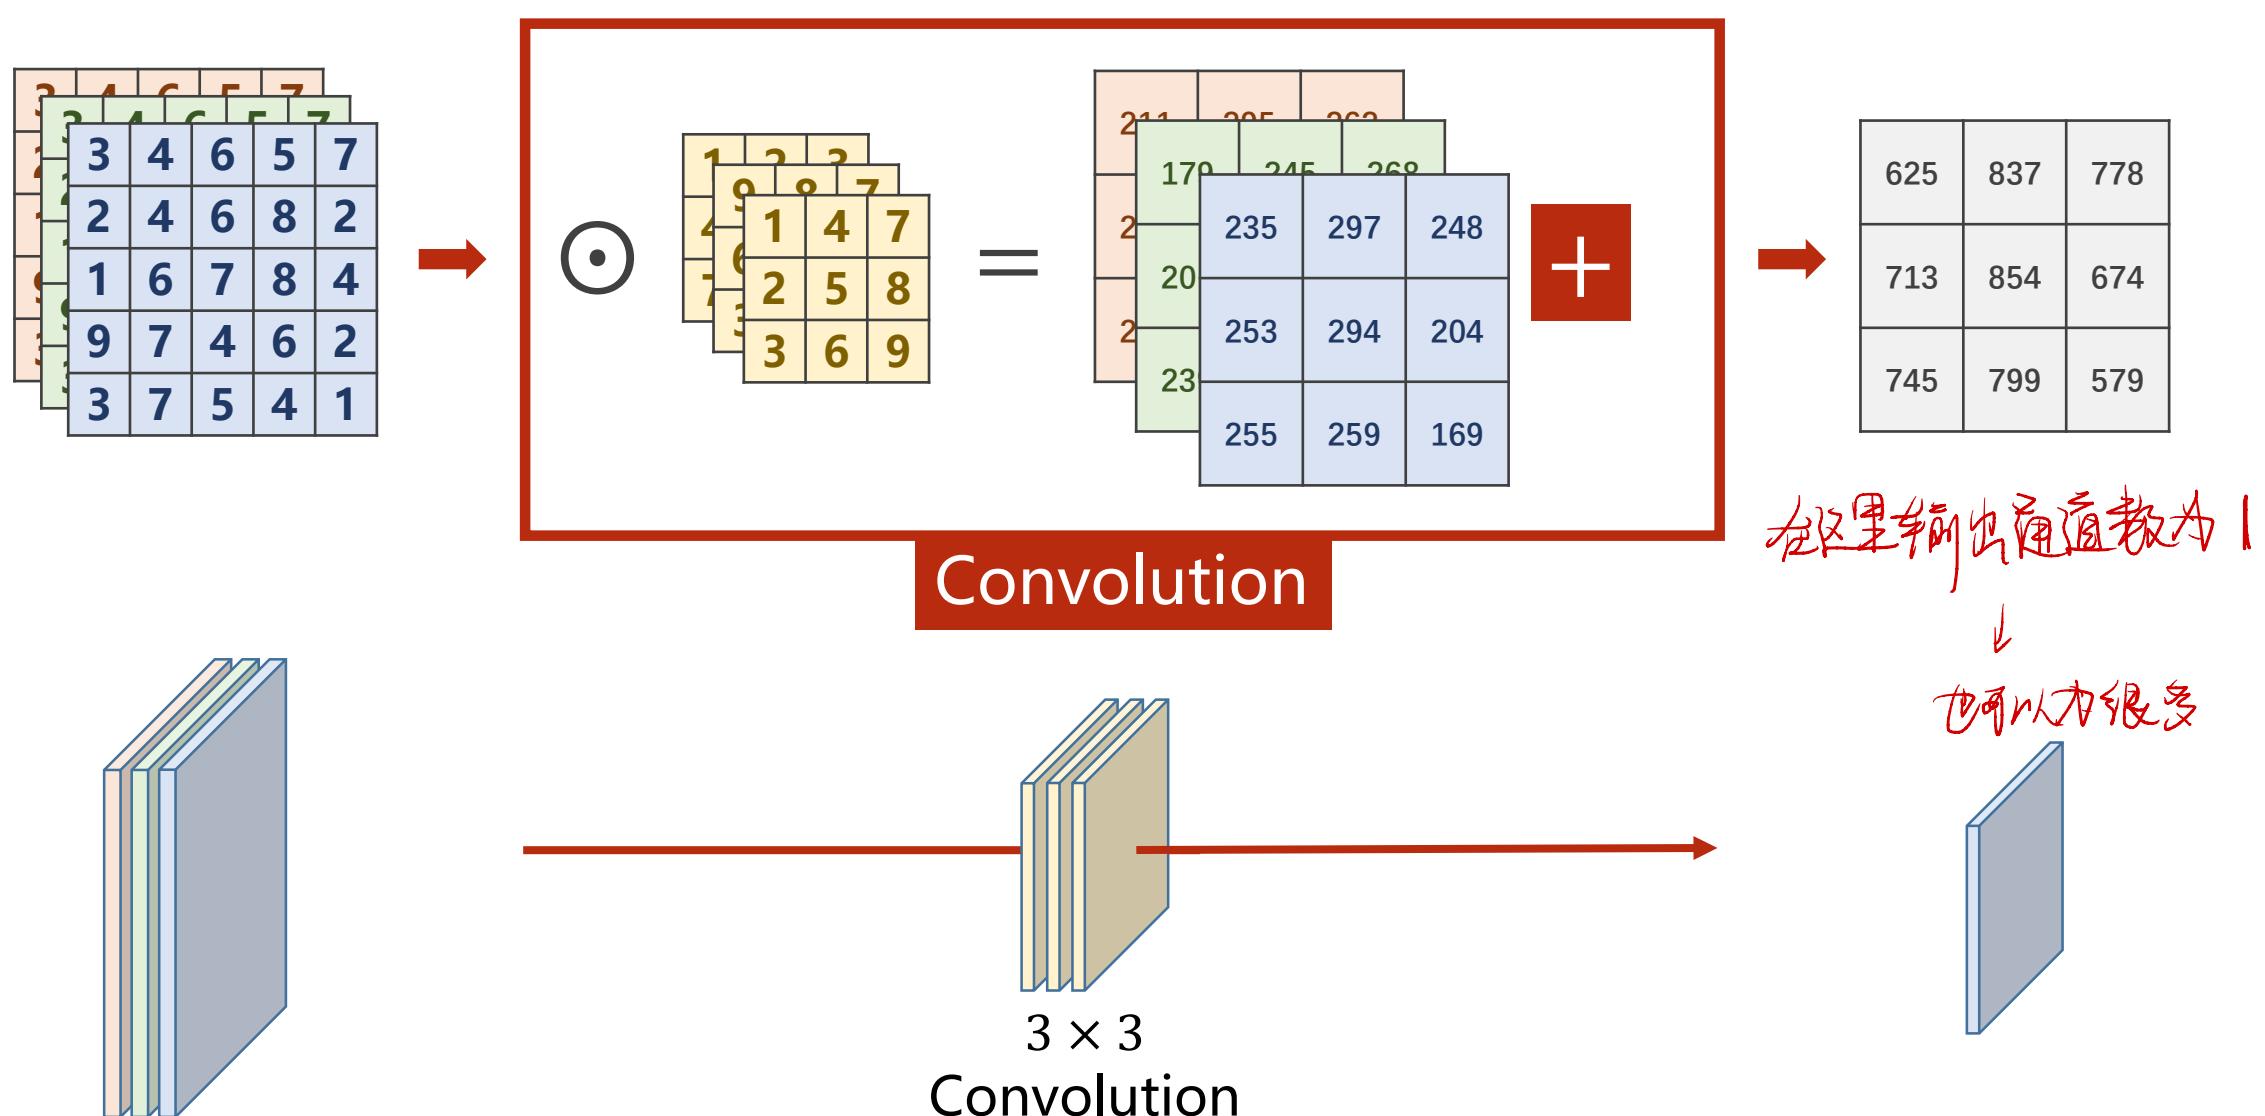
3	4	6	5	7
2	4	6	8	2
1	6	7	8	4
9	7	4	6	2
3	7	5	4	1

3	4	6	5	7
2	4	6	8	2
1	6	7	8	4
9	7	4	6	2
3	7	5	4	1

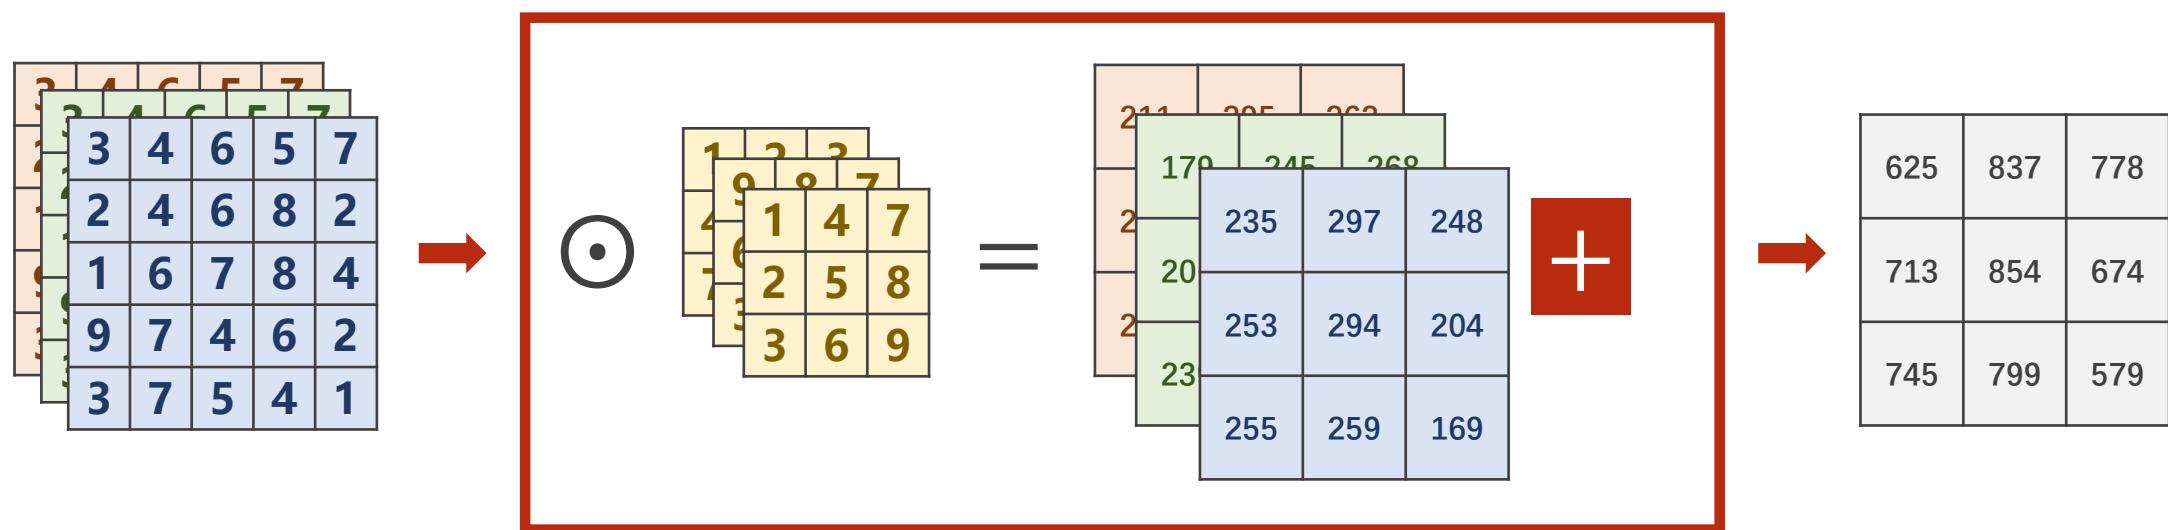


Convolution

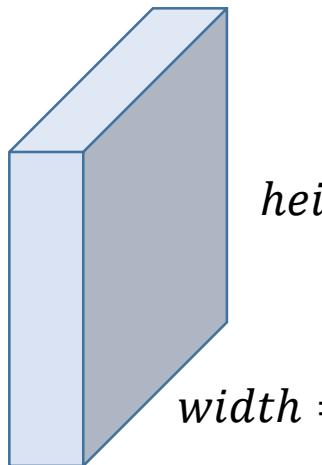
# Convolution – 3 Input Channels



# Convolution – 3 Input Channels

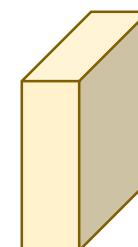


channel = 3



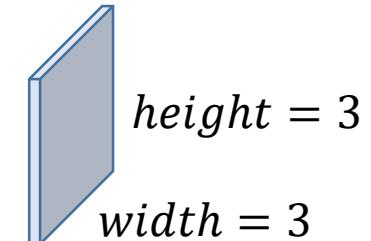
Convolution

channel = 3



$3 \times 3$   
Convolution

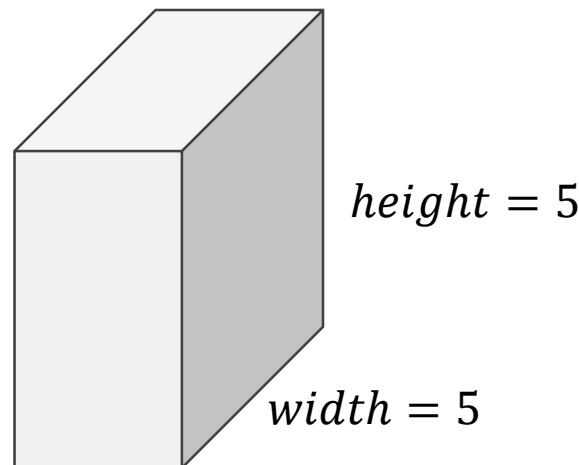
channel = 1



# Convolution – N Input Channels

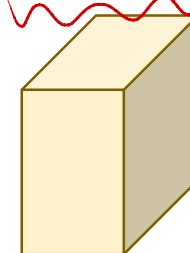
若输入通道数为  $n$

$channel = n$

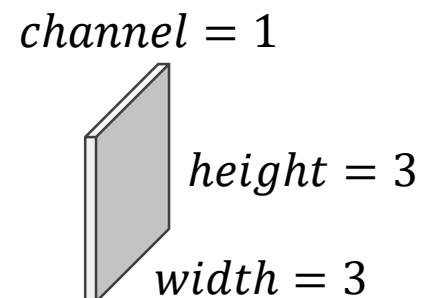


The shape of kernel tensor is  $(n, 3, 3)$

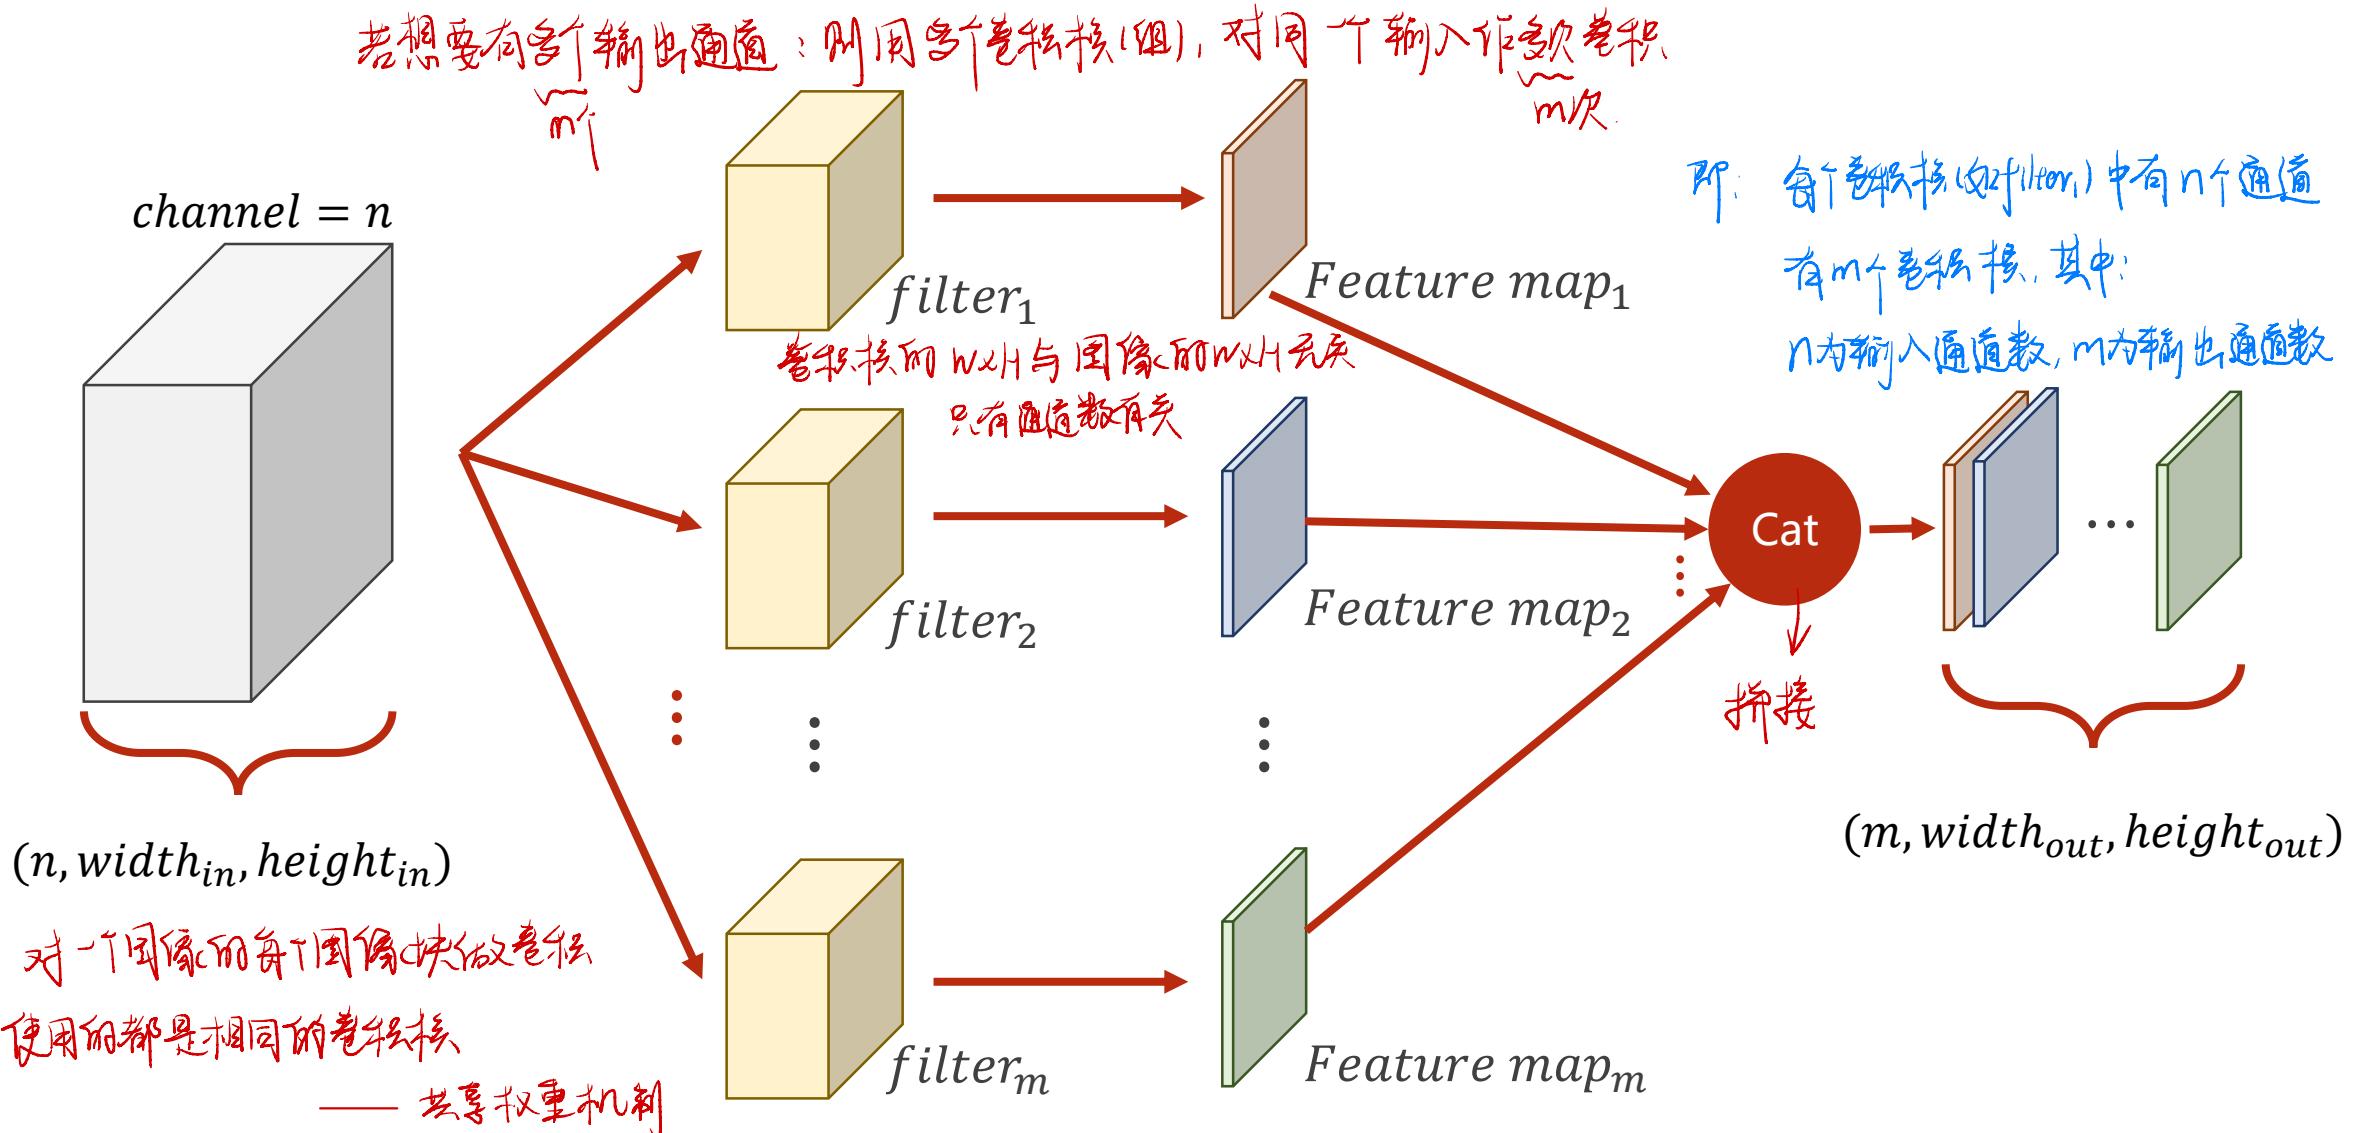
channel = n  
卷积核也要变为  $n$  个



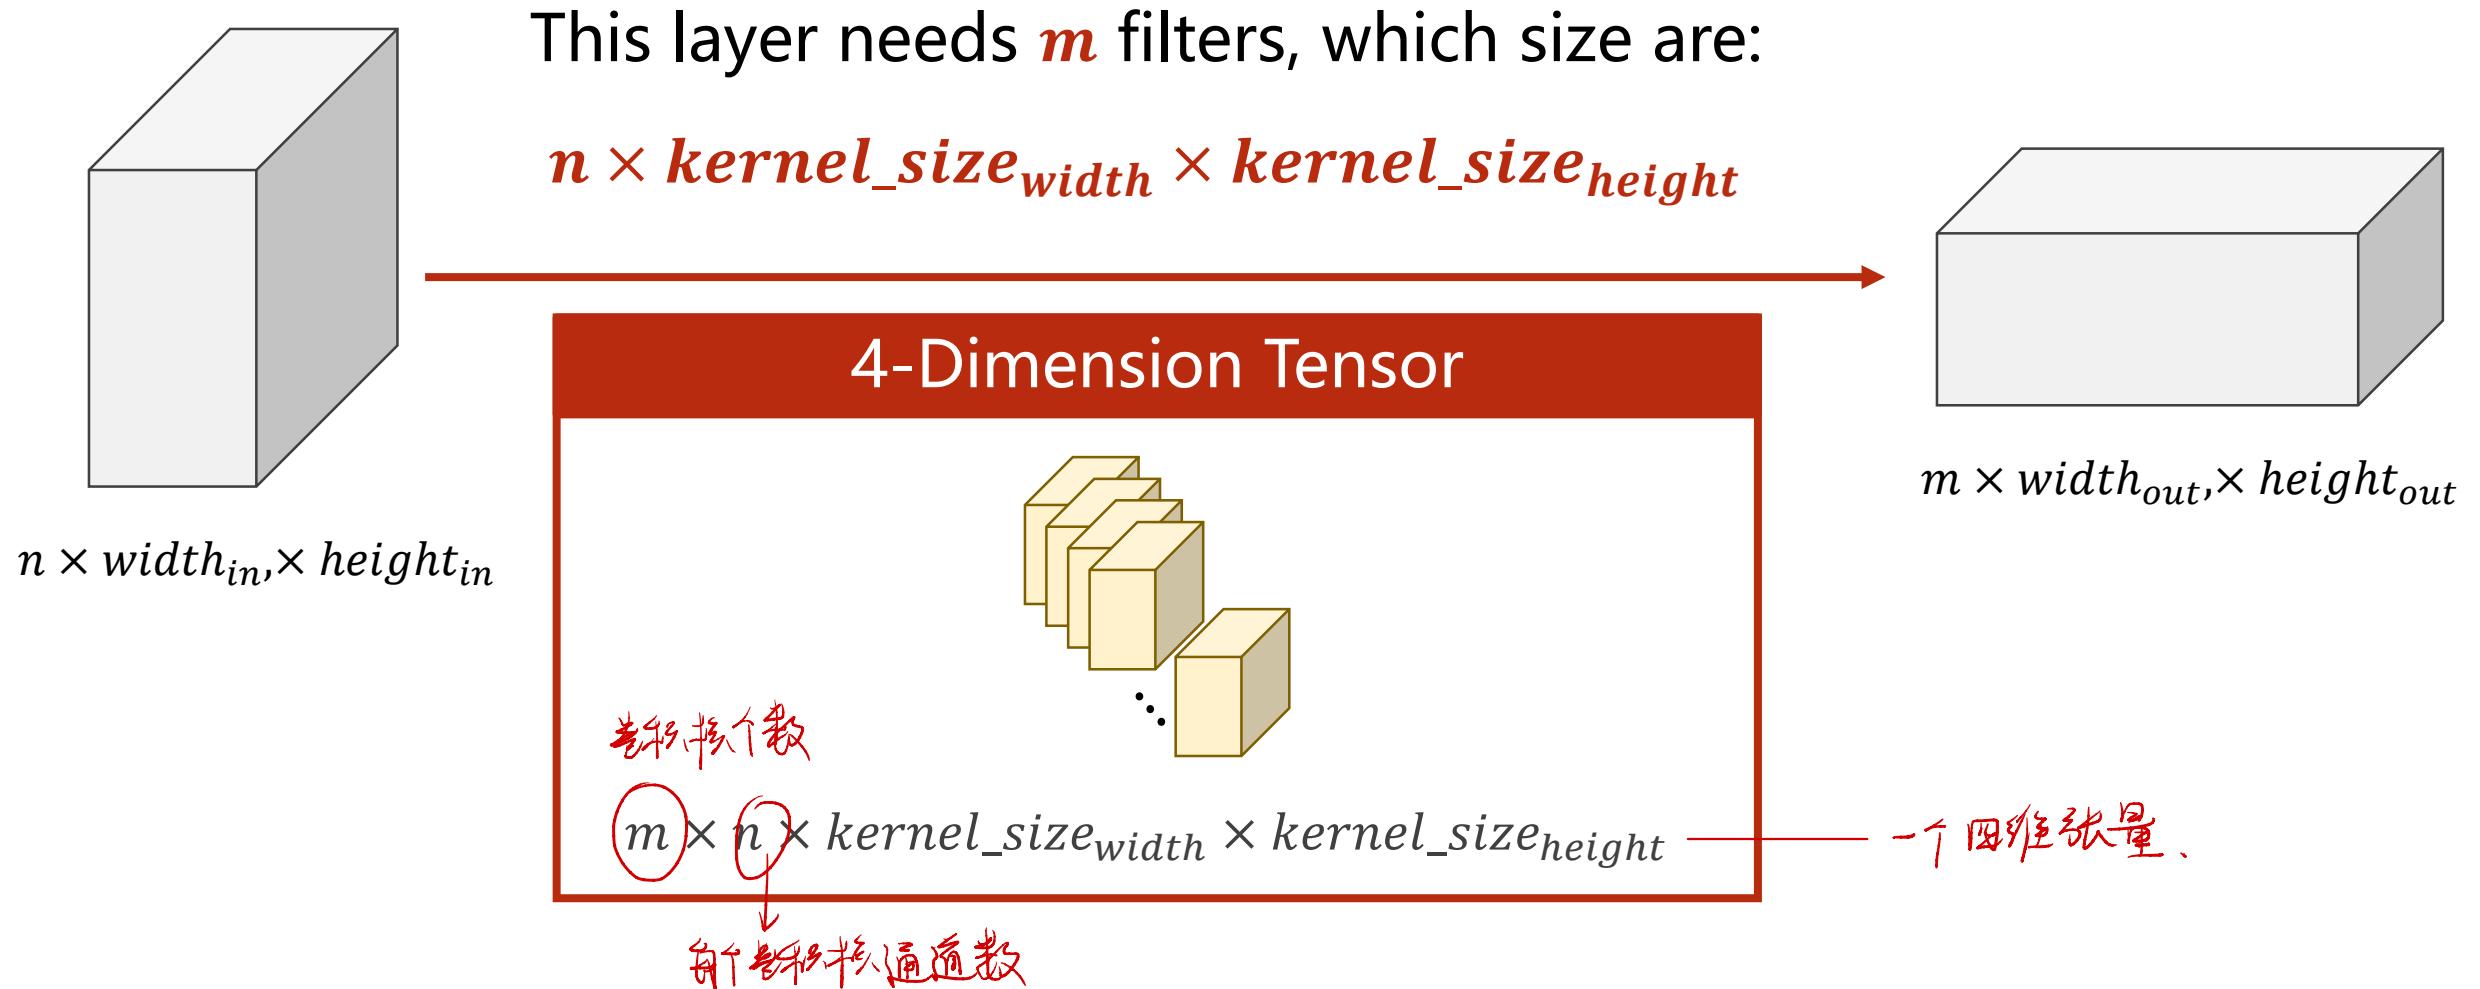
$3 \times 3$   
Convolution



# Convolution – N Input Channels and M Output Channels



# Convolutional Layer



# Convolutional Layer

计算过程：

```
import torch
in_channels, out_channels = n, m
width, height = 100, 100 // 图像大小
kernel_size = 3 // 卷积核大小
batch_size = 1 // pytorch 中所有输入必须是小批量输入
                所以要把一个图像分为多个 Batch
input = torch.randn(batch_size,
                    in_channels,
                    width,
                    height) → 先随机一个权重，后面训练(?)
```

卷积层      生成卷积层的模块      可以是正方形  
out\_channels, // m      kernel\_size=kernel\_size// 3x3      也可以长方形，一般正方形，常见大小为奇数  
conv\_layer = torch.nn.Conv2d(in\_channels, // n  
output = conv\_layer(input)  
print(input.shape)  
print(output.shape)  
print(conv\_layer.weight.shape)

# Convolutional Layer

```
import torch
in_channels, out_channels= 5, 10
width, height = 100, 100
kernel_size = 3
batch_size = 1

input = torch.randn(batch_size,
                   in_channels,
                   width,
                   height)

conv_layer = torch.nn.Conv2d(in_channels,
                           out_channels,
                           kernel_size=kernel_size)

output = conv_layer(input)

print(input.shape)
print(output.shape)
print(conv_layer.weight.shape)
```

注意：卷积层并不在乎图像 or batch 的大小，它只是遍历一遍  
从头到底

torch.Size([1, 5, 100, 100])  
torch.Size([1, 10, 98, 98])  
torch.Size([10, 5, 3, 3])

↓  
输入通道 → 输出通道

# Convolutional Layer – padding=1

卷积层的常见参数

**Input**

0	0	0	0	0	0	0	0
0	3	4	6	5	7	0	0
0	2	4	6	8	2	0	0
0	1	6	7	8	4	0	0
0	9	7	4	6	2	0	0
0	3	7	5	4	1	0	0
0	0	0	0	0	0	0	0

**Kernel**

① padding：为使输入大小

和输出大小一致，我们可以扩大图像（常见方法是在周围添0）

添加圈数与卷积核大小有关，若卷积核大小  $n \times n$

则圈数为  $\frac{n}{2}$  取整



1	2	3
4	5	6
7	8	9

=

**Output**

91				
	211	295	262	
	259	282	214	
	251	253	169	

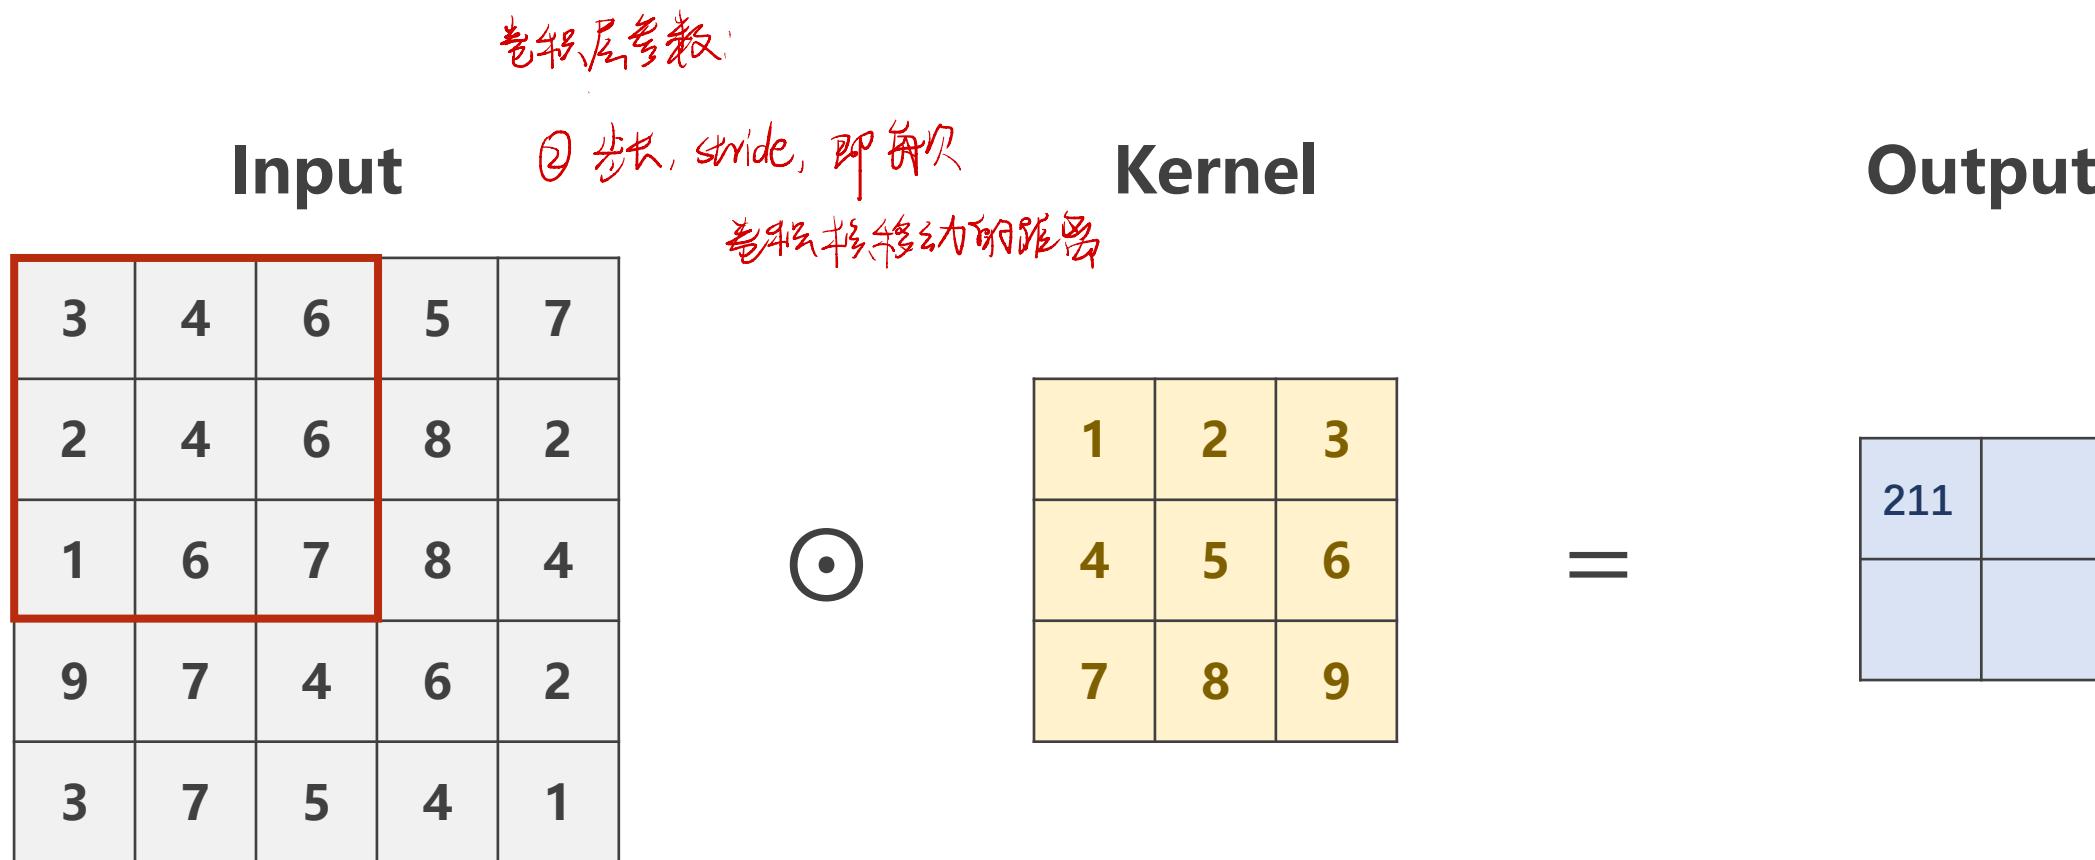
# Convolutional Layer – padding=1

```
import torch  
  
input = [3, 4, 6, 5, 7,  
        2, 4, 6, 8, 2,  
        1, 6, 7, 8, 4,  
        9, 7, 4, 6, 2,  
        3, 7, 5, 4, 1]  
input = torch.Tensor(input).view(1, 1, 5, 5)  
  
conv_layer = torch.nn.Conv2d(1, 1, kernel_size=3, padding=1, bias=False) // 构建卷积层  
  
kernel = torch.Tensor([1, 2, 3, 4, 5, 6, 7, 8, 9]).view(1, 1, 3, 3)  
conv_layer.weight.data = kernel.data // 初始化卷积权重  
  
output = conv_layer(input)  
print(output)
```

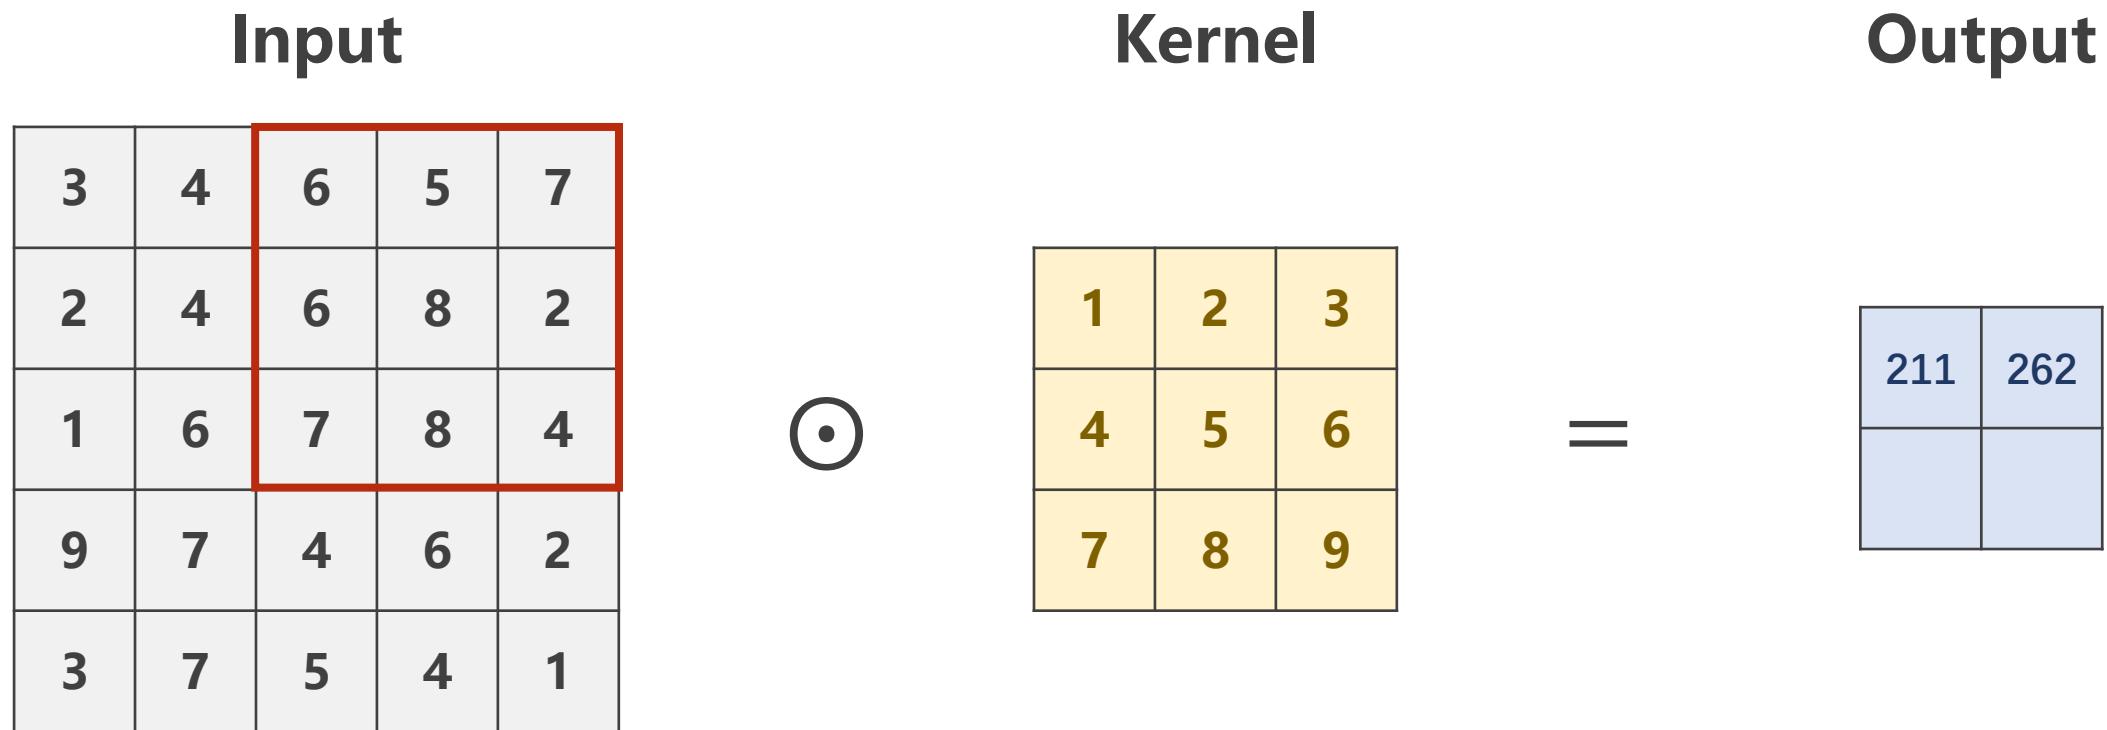
手写注释：

- “本质上依然是线性计算”
- “填充的圈数：1圈”
- “padding=1, bias=False”
- “是否加偏置量”

# Convolutional Layer – stride=2



# Convolutional Layer – stride=2



# Convolutional Layer – stride=2

**Input**

3	4	6	5	7
2	4	6	8	2
1	6	7	8	4
9	7	4	6	2
3	7	5	4	1

**Kernel**



1	2	3
4	5	6
7	8	9

=

**Output**

211	262
251	

# Convolutional Layer – stride=2

```
import torch

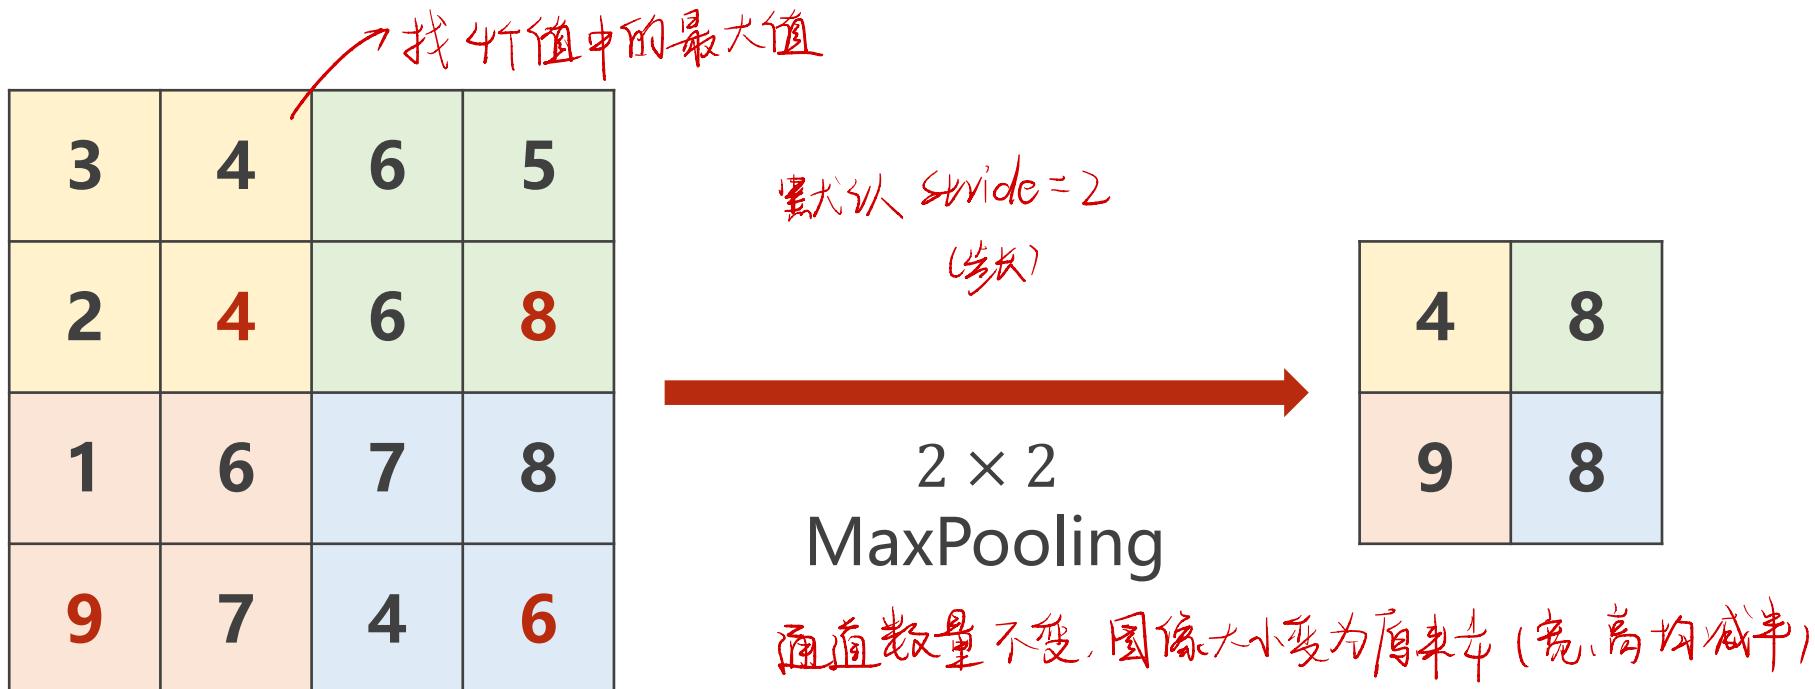
input = [3, 4, 6, 5, 7,
         2, 4, 6, 8, 2,
         1, 6, 7, 8, 4,
         9, 7, 4, 6, 2,
         3, 7, 5, 4, 1]
input = torch.Tensor(input).view(1, 1, 5, 5)

conv_layer = torch.nn.Conv2d(1, 1, kernel_size=3, stride=2, bias=False)
卷积步长
kernel = torch.Tensor([1, 2, 3, 4, 5, 6, 7, 8, 9]).view(1, 1, 3, 3)
conv_layer.weight.data = kernel.data

output = conv_layer(input)
print(output)
```

# Max Pooling Layer

下采样：最大池化层



# Max Pooling Layer

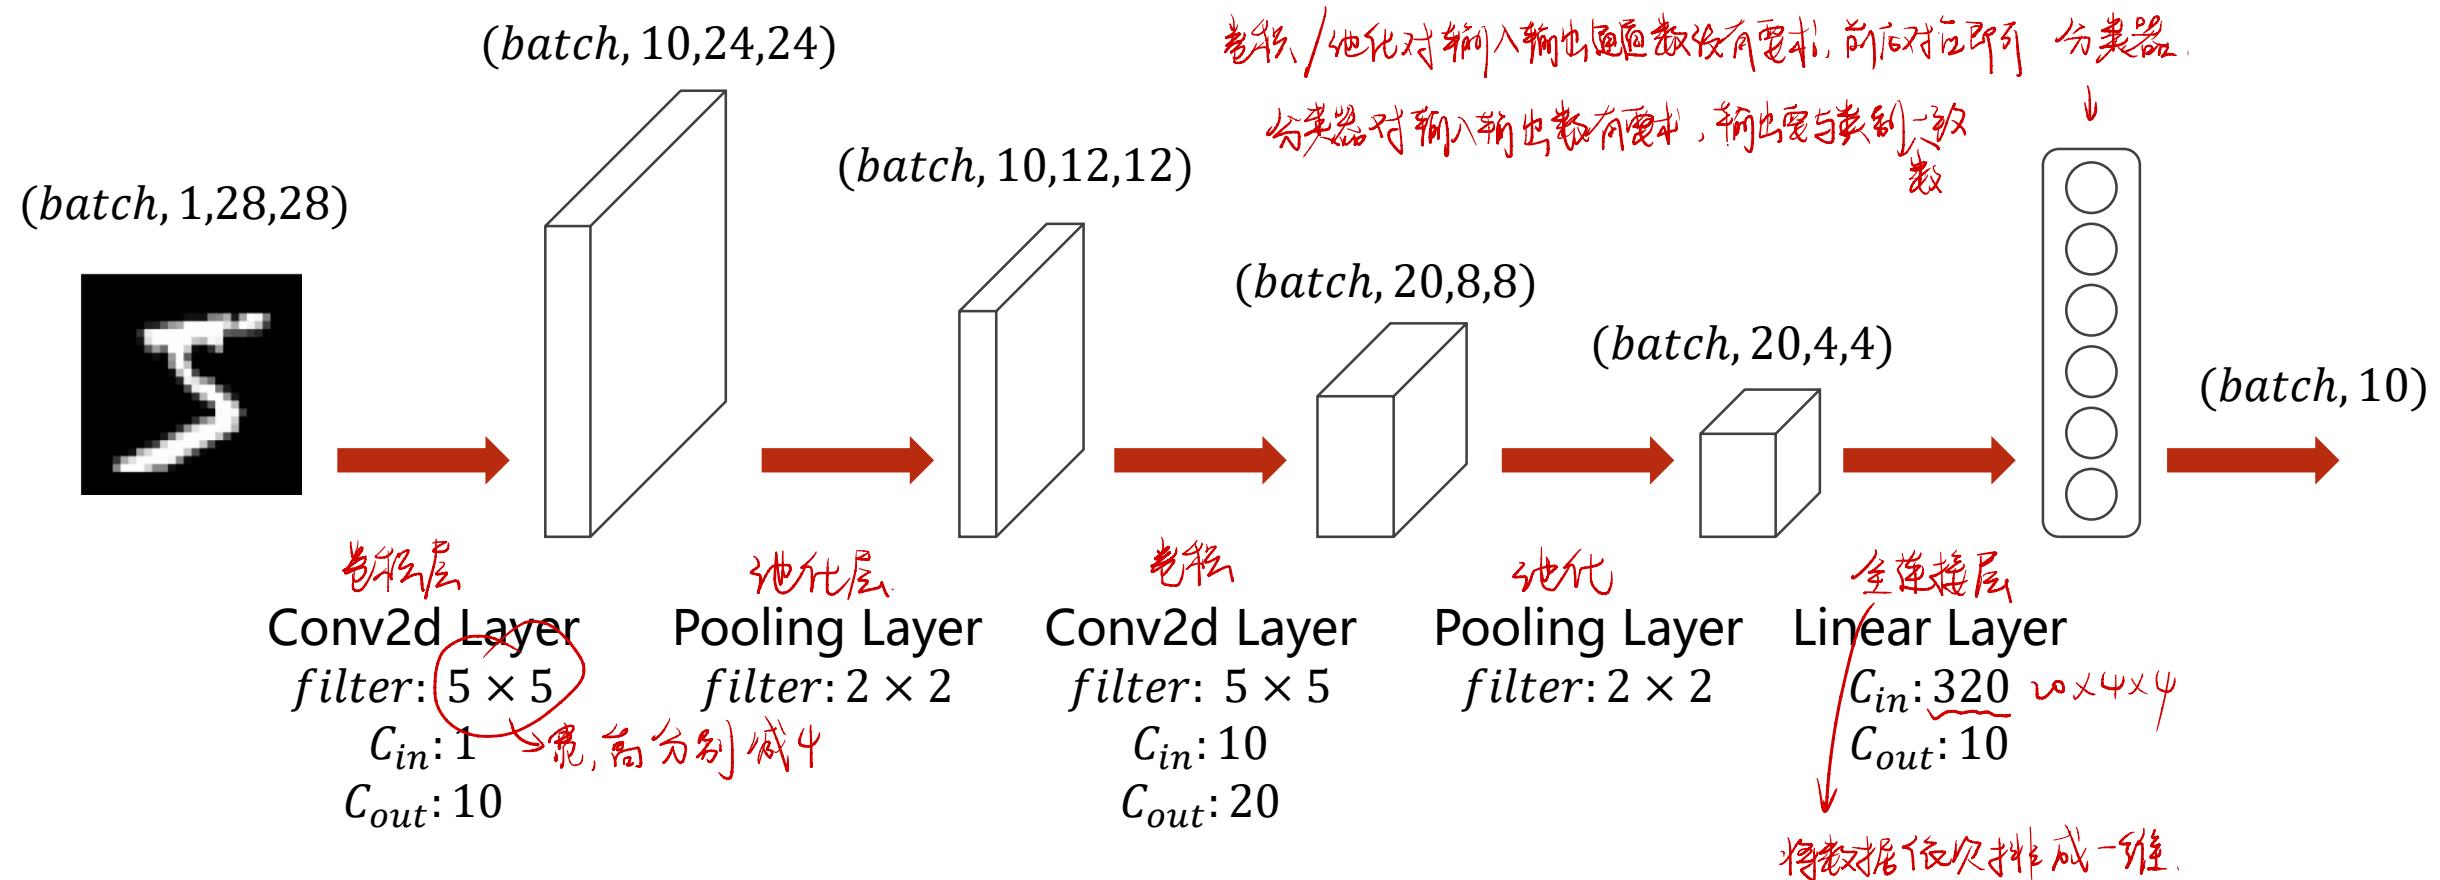
```
import torch

input = [3, 4, 6, 5,
         2, 4, 6, 8,
         1, 6, 7, 8,
         9, 7, 4, 6,
        ]
input = torch.Tensor(input).view(1, 1, 4, 4)

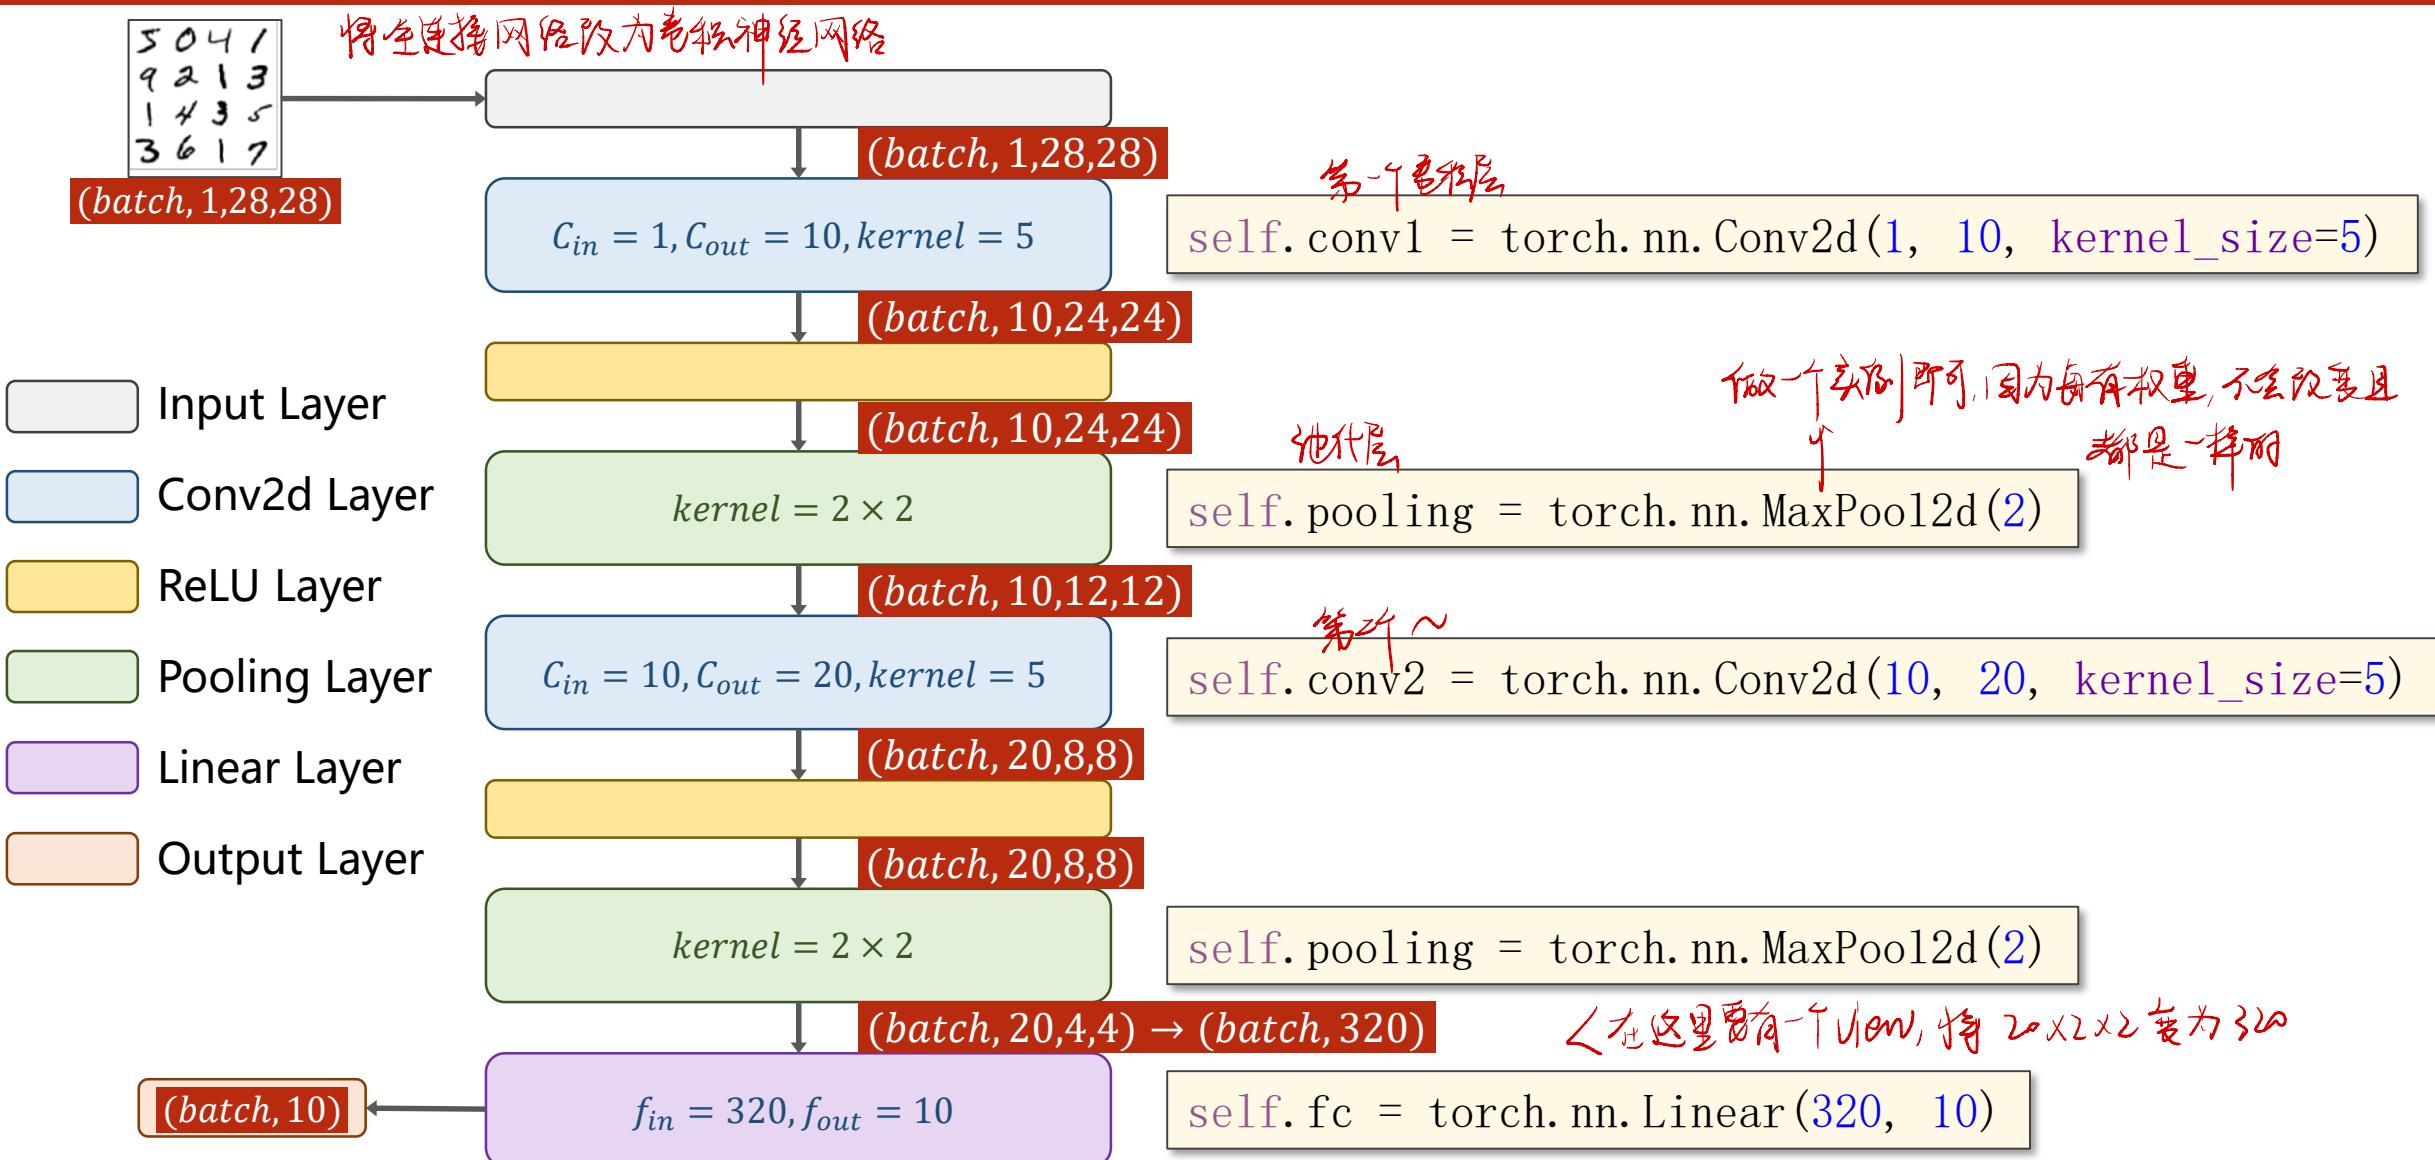
maxpooling_layer = torch.nn.MaxPool2d(kernel_size=2)

output = maxpooling_layer(input)
print(output)
```

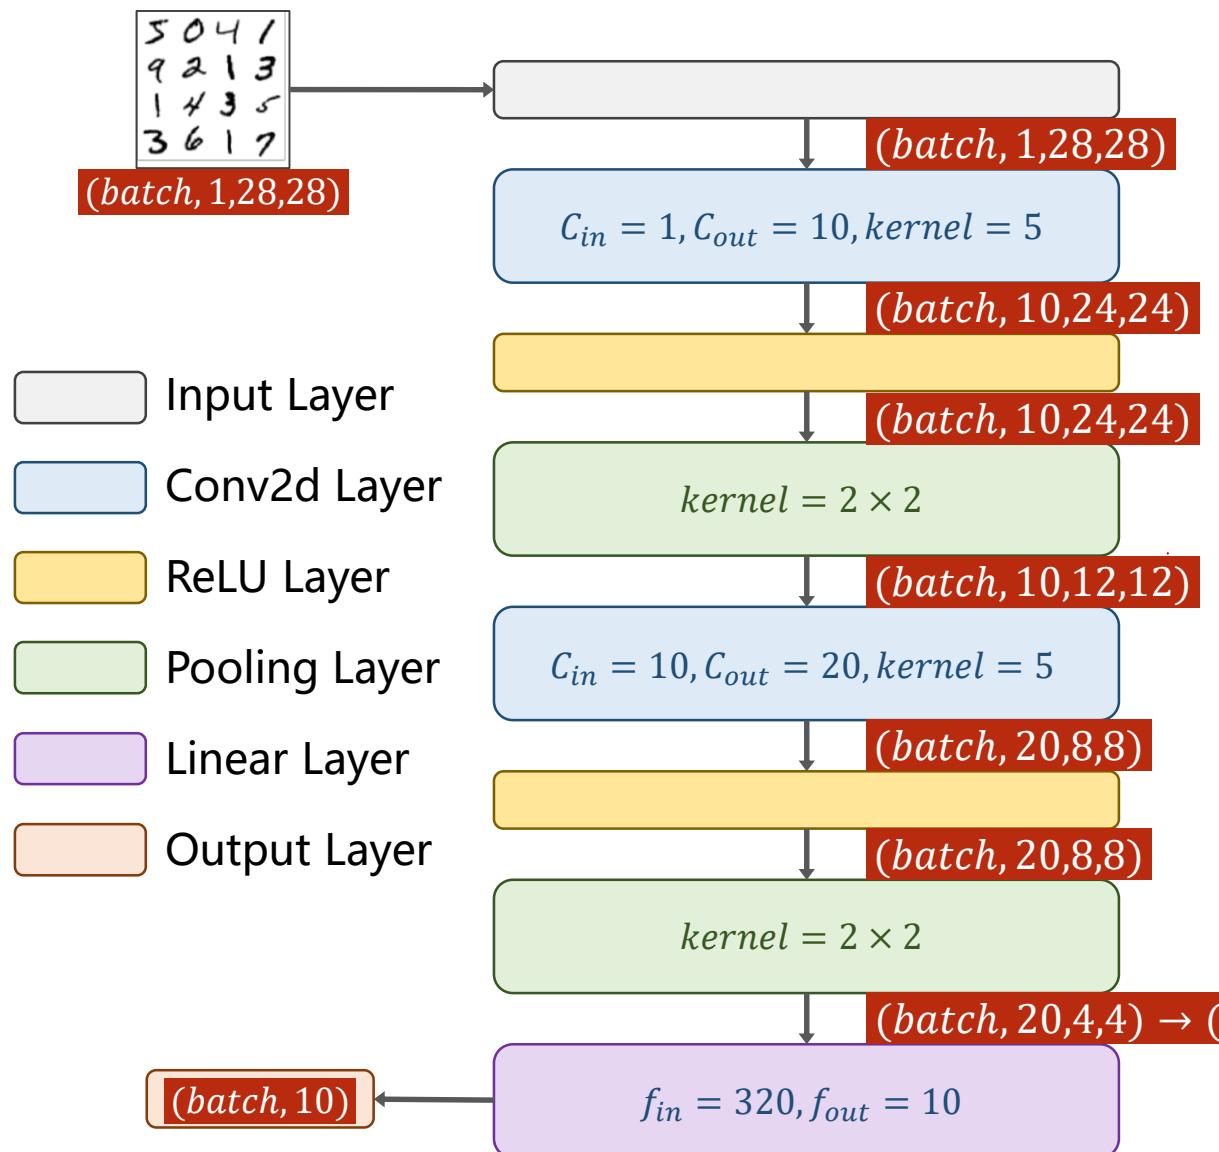
# A Simple Convolutional Neural Network



# Revision: Fully Connected Neural Network



# Revision: Fully Connected Neural Network



```
class Net(torch.nn.Module):  
    def __init__(self):  
        super(Net, self).__init__()  
        self.conv1 = torch.nn.Conv2d(1, 10, kernel_size=5)  
        self.conv2 = torch.nn.Conv2d(10, 20, kernel_size=5)  
        self.pooling = torch.nn.MaxPool2d(2)  
        self.fc = torch.nn.Linear(320, 10) // 全连接层  
  
    def forward(self, x):  
        # Flatten data from (n, 1, 28, 28) to (n, 784)  
        batch_size = x.size(0) // 先求batch大小 固定写法  
        x = F.relu(self.pooling(self.conv1(x)))  
        x = F.relu(self.pooling(self.conv2(x)))  
        x = x.view(batch_size, -1) # flatten // view  
        x = self.fc(x) // 全连接一下很方便，因为要用交叉熵的那个方法  
        return x
```

model = Net()

先卷积，再池化，再ReLU  
非线性化

# How to use GPU – 1. Move Model to GPU

如何使用GPU

```
class Net(torch.nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = torch.nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = torch.nn.Conv2d(10, 20, kernel_size=5)
        self.pooling = torch.nn.MaxPool2d(2)
        self.fc = torch.nn.Linear(320, 10)

    def forward(self, x):
        # Flatten data from (n, 1, 28, 28) to (n, 784)
        batch_size = x.size(0)
        x = F.relu(self.pooling(self.conv1(x)))
        x = F.relu(self.pooling(self.conv2(x)))
        x = x.view(batch_size, -1)
        x = self.fc(x)
        return x
```

Define device as the first visible cuda device if we have CUDA available.

```
model = Net()
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

第一块显卡

# How to use GPU – 1. Move Model to GPU

标记模型时

```
class Net(torch.nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = torch.nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = torch.nn.Conv2d(10, 20, kernel_size=5)
        self.pooling = torch.nn.MaxPool2d(2)
        self.fc = torch.nn.Linear(320, 10)

    def forward(self, x):
        # Flatten data from (n, 1, 28, 28) to (n, 784)
        batch_size = x.size(0)
        x = F.relu(self.pooling(self.conv1(x)))
        x = F.relu(self.pooling(self.conv2(x)))
        x = x.view(batch_size, -1) # flatten
        x = self.fc(x)
        return x
```

model = Net()  
device = torch.device("cuda:0" if torch.cuda.is\_available() else "cpu")  
**model.to(device)** // 将所有内容放入显卡

Convert parameters and buffers of all modules to CUDA Tensor.

# How to use GPU – 2. Move Tensors to GPU

训练时

```
def train(epoch):
    running_loss = 0.0
    for batch_idx, data in enumerate(train_loader, 0):
        inputs, target = data
        inputs, target = inputs.to(device), target.to(device) ← 添加这一行，将数据放入GPU
        optimizer.zero_grad()

        # forward + backward + update
        outputs = model(inputs)
        loss = criterion(outputs, target)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        if batch_idx % 300 == 299:
            print(' [%d, %5d] loss: %.3f' % (epoch + 1, batch_idx + 1, running_loss / 2000))
            running_loss = 0.0
```

在这里使用了 mini-batch

**inputs, target = inputs.to(device), target.to(device)** ← 添加这一行，将数据放入GPU

Send the inputs and targets at every  
step to the GPU.

# How to use GPU – 2. Move Tensors to GPU

GPU 成对

```
def test():
    correct = 0
    total = 0
    with torch.no_grad():
        for data in test_loader:
            inputs, target = data
            inputs, target = inputs.to(device), target.to(device) ← 添加这一行
            outputs = model(inputs)
            _, predicted = torch.max(outputs.data, dim=1)
            total += target.size(0)
            correct += (predicted == target).sum().item()
    print('Accuracy on test set: %d %% [%d/%d]' % (100 * correct / total, correct, total))
```

Send the inputs and targets at every step to the GPU.

# Results

**Accuracy on test set: 6 % [637/10000]**

[1, 300] loss: 0.098

[1, 600] loss: 0.035

[1, 900] loss: 0.025

**Accuracy on test set: 96 % [9605/10000]**

[2, 300] loss: 0.021

[2, 600] loss: 0.017

[2, 900] loss: 0.015

**Accuracy on test set: 97 % [9709/10000]**

.....

[9, 300] loss: 0.006

[9, 600] loss: 0.006

[9, 900] loss: 0.007

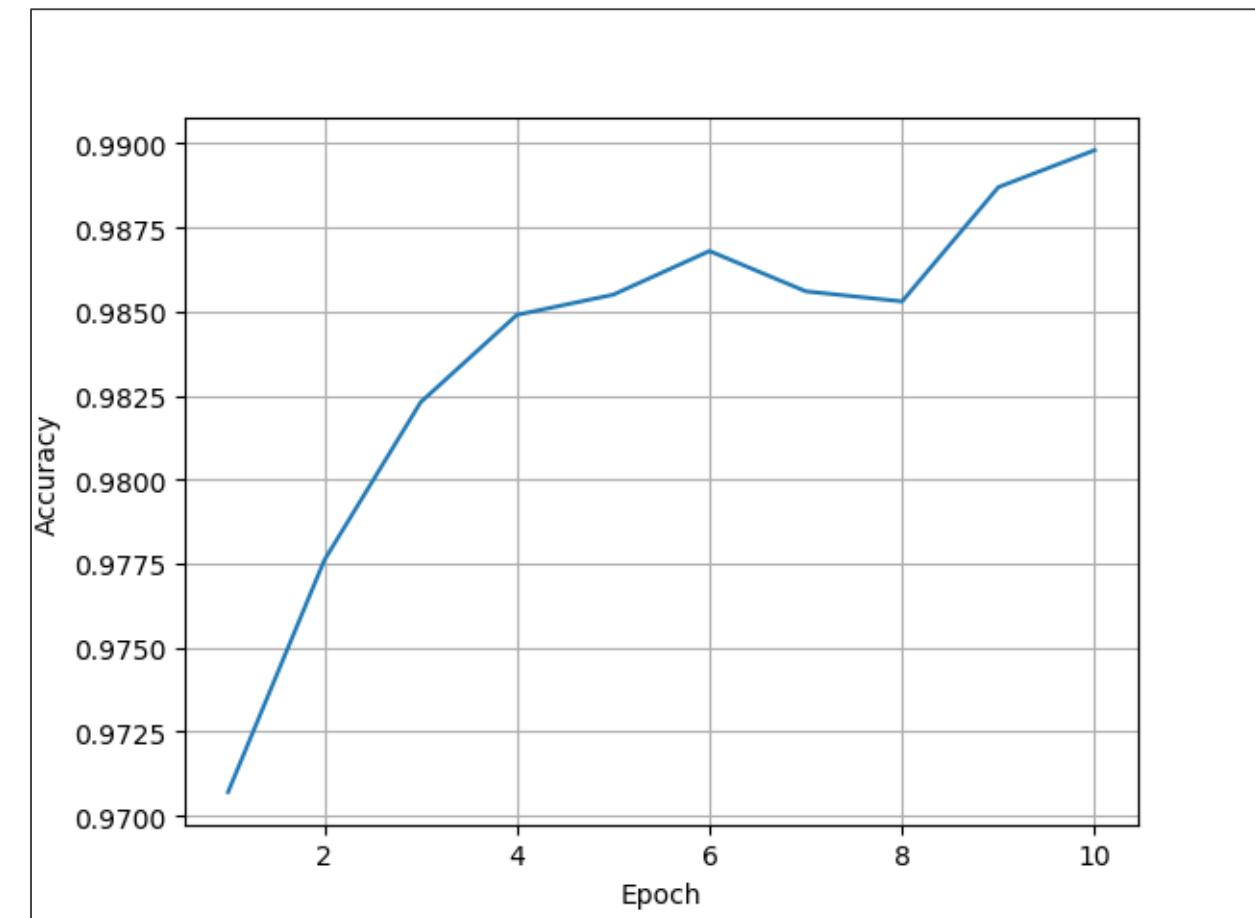
**Accuracy on test set: 98 % [9857/10000]**

[10, 300] loss: 0.006

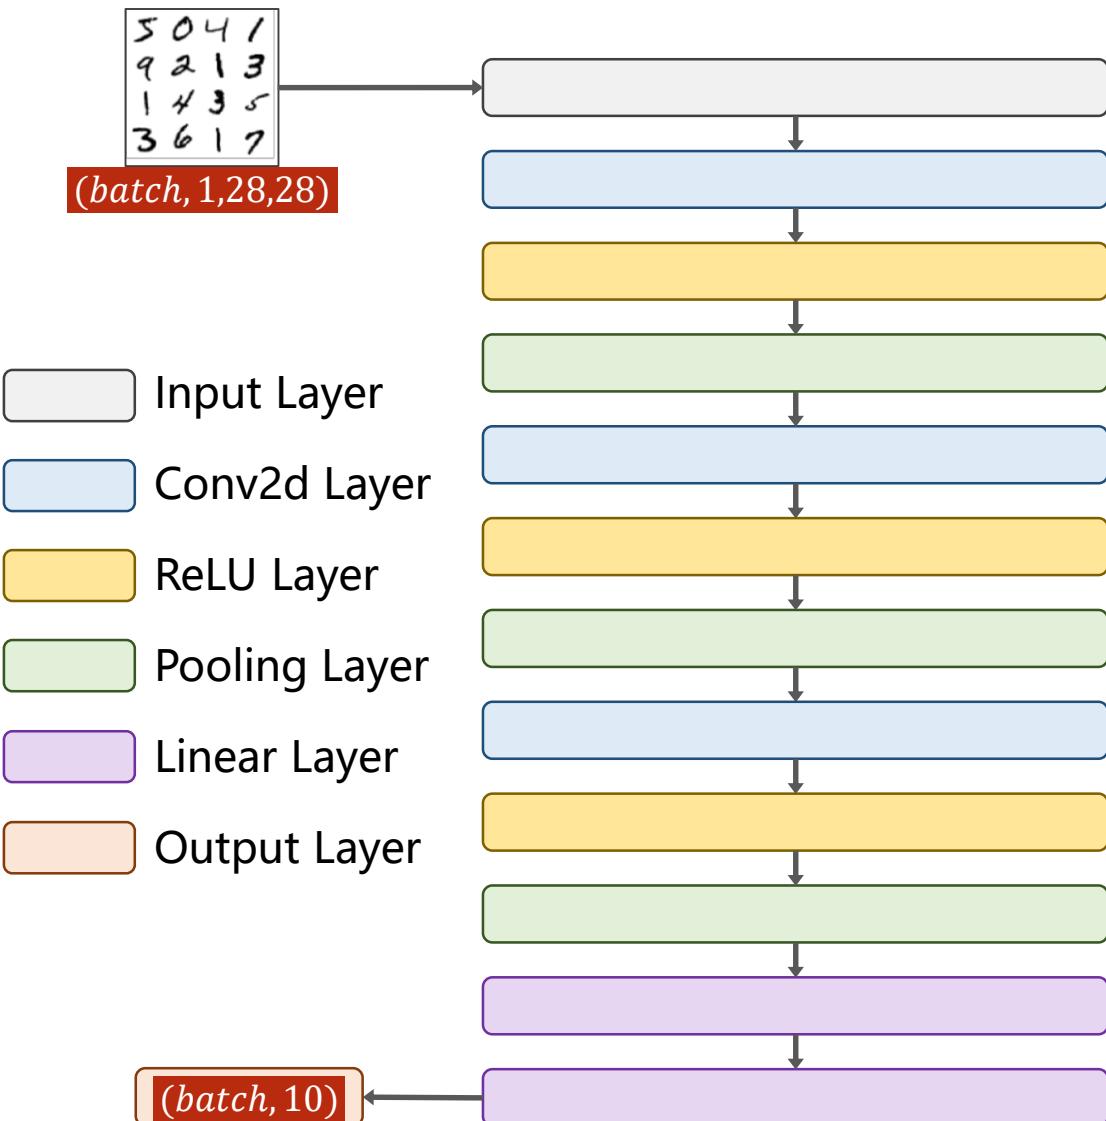
[10, 600] loss: 0.006

[10, 900] loss: 0.006

**Accuracy on test set: 98 % [9869/10000]**



# Exercise 10-1



- Try a more complex CNN:
  - Conv2d Layer \*3
  - ReLU Layer \* 3
  - MaxPooling Layer \* 3
  - Linear Layer \* 3
- Try different configuration of this CNN:
  - Compare their performance.



# PyTorch Tutorial

## 10. Basic CNN